



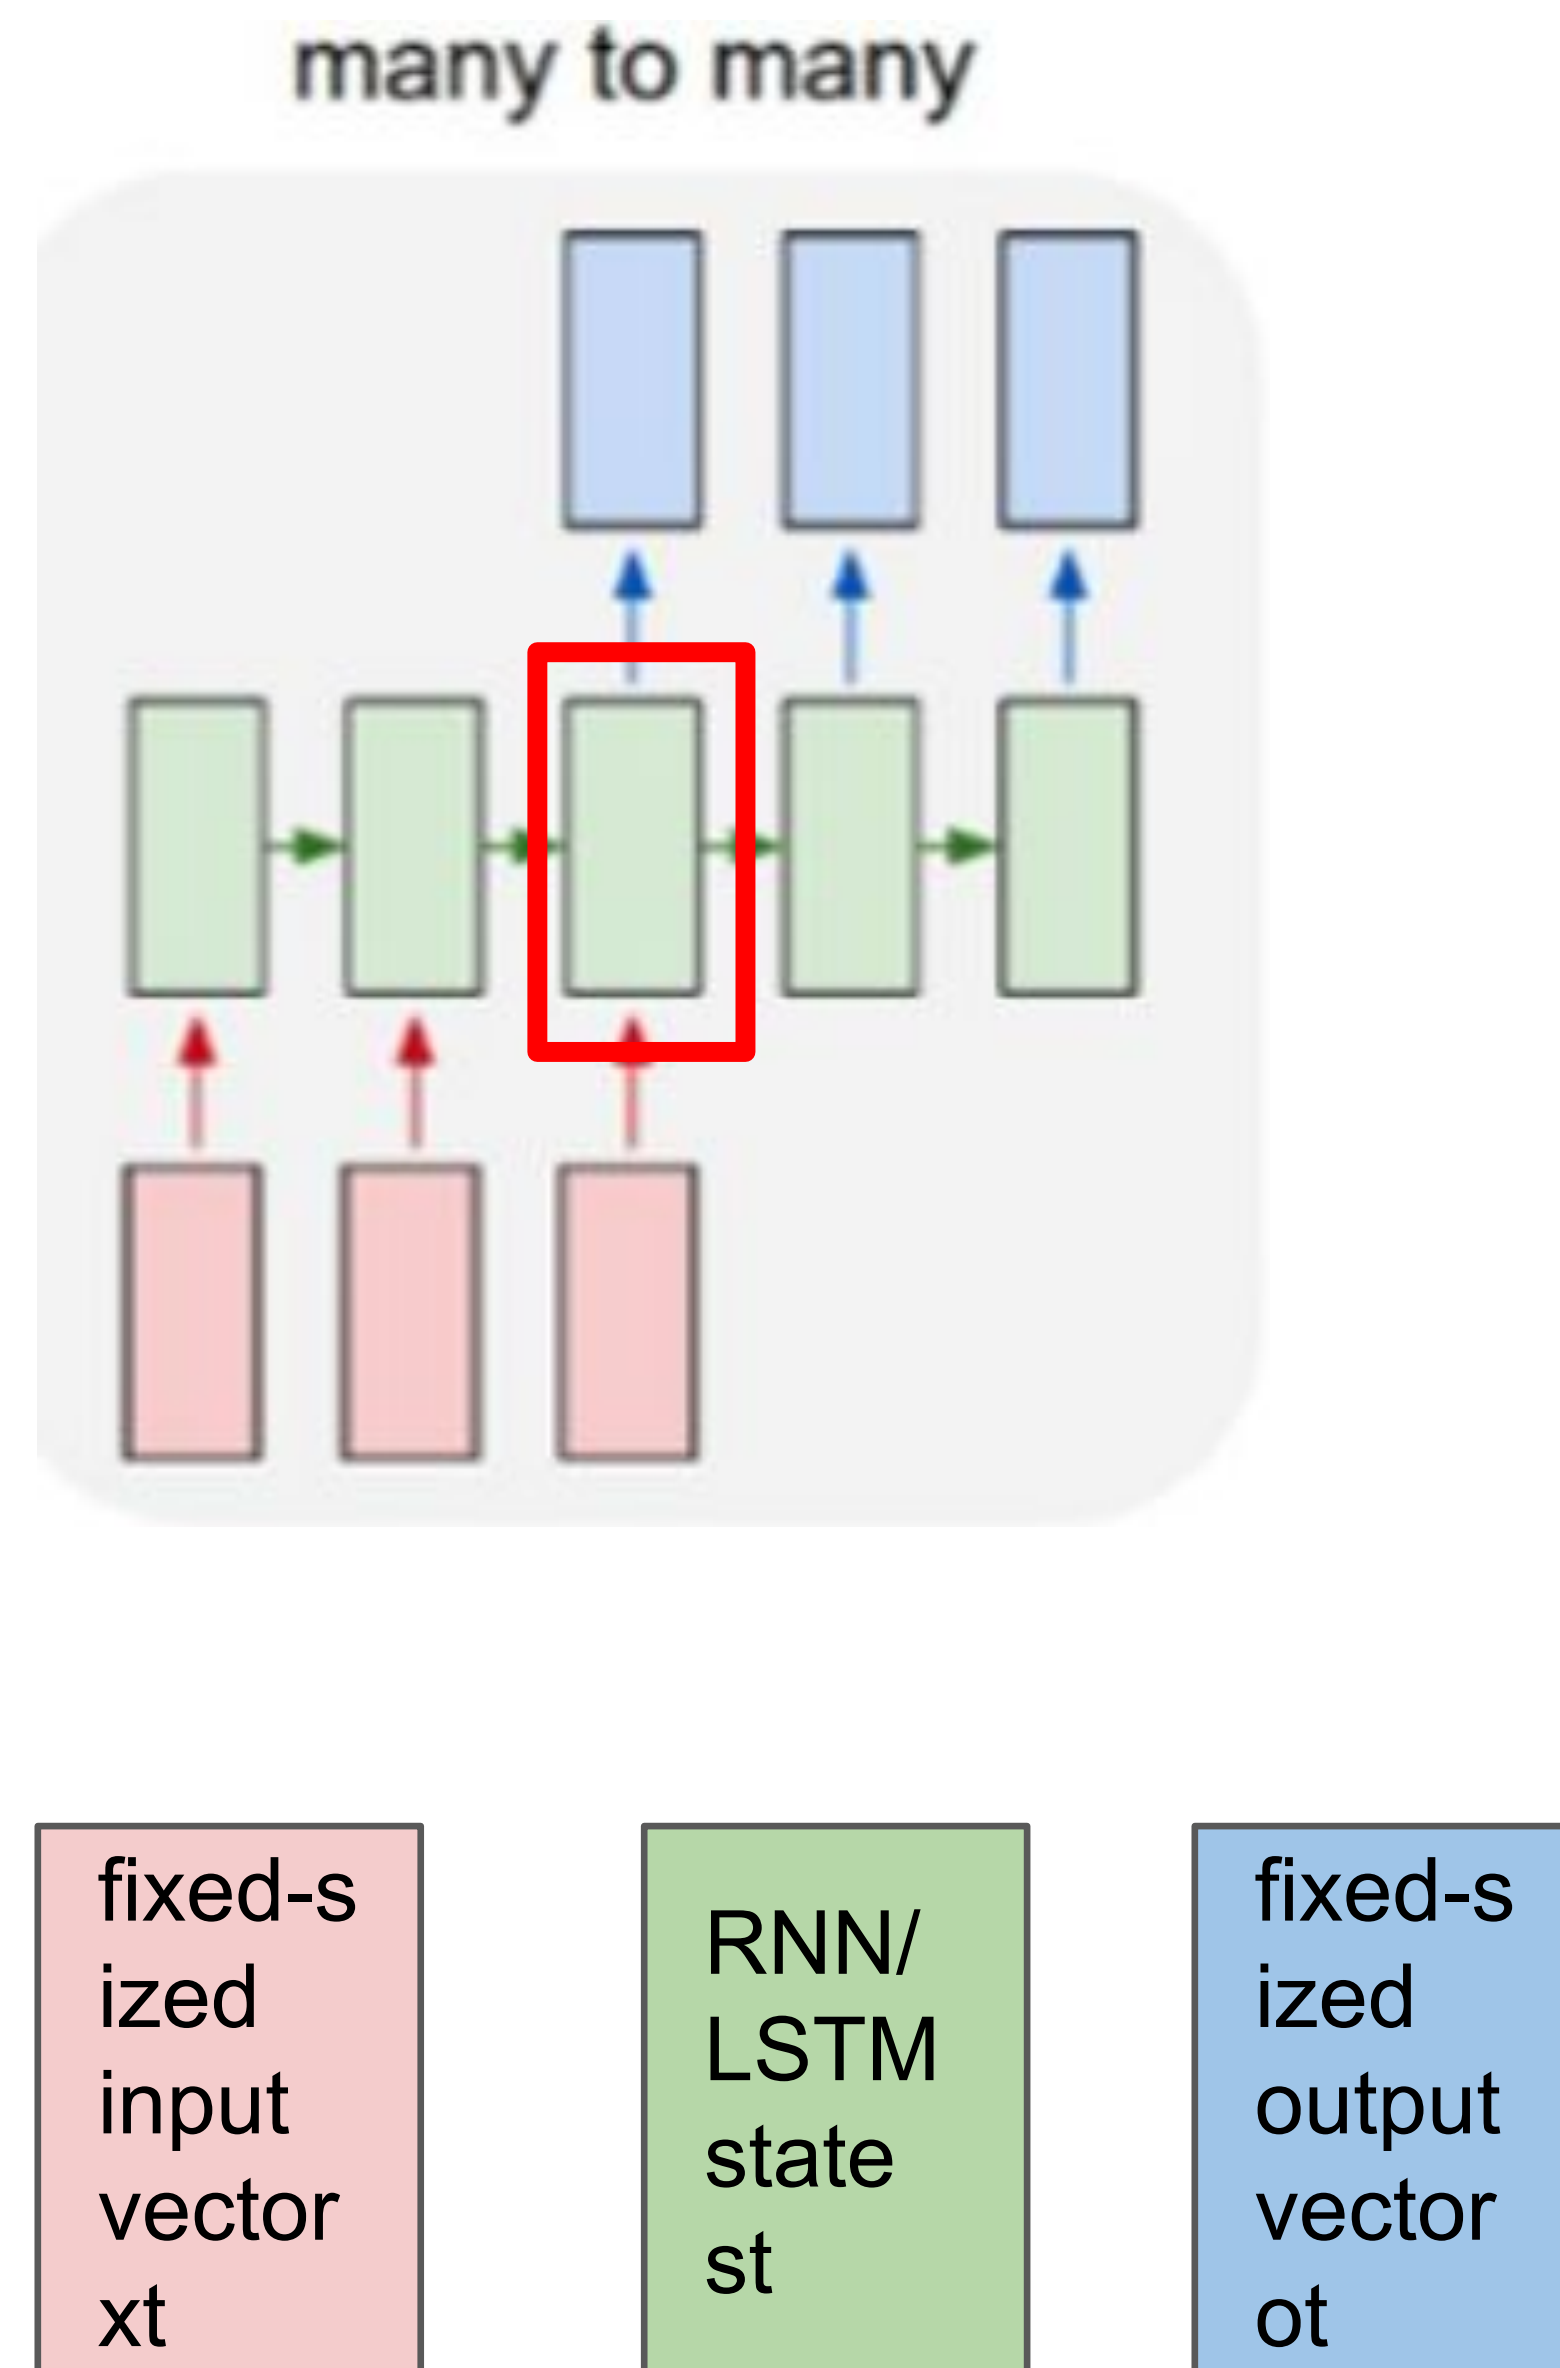
IF3270 Pembelajaran Mesin

Attention & Transformer

Tim Pengajar IF3270



Review: Sequence Model Architecture



- Generate **a hidden vector** as representation of input sequence, then use it to generate output.
- For encoder-decoder (Seq2Seq: Sequence to Sequence) architecture.

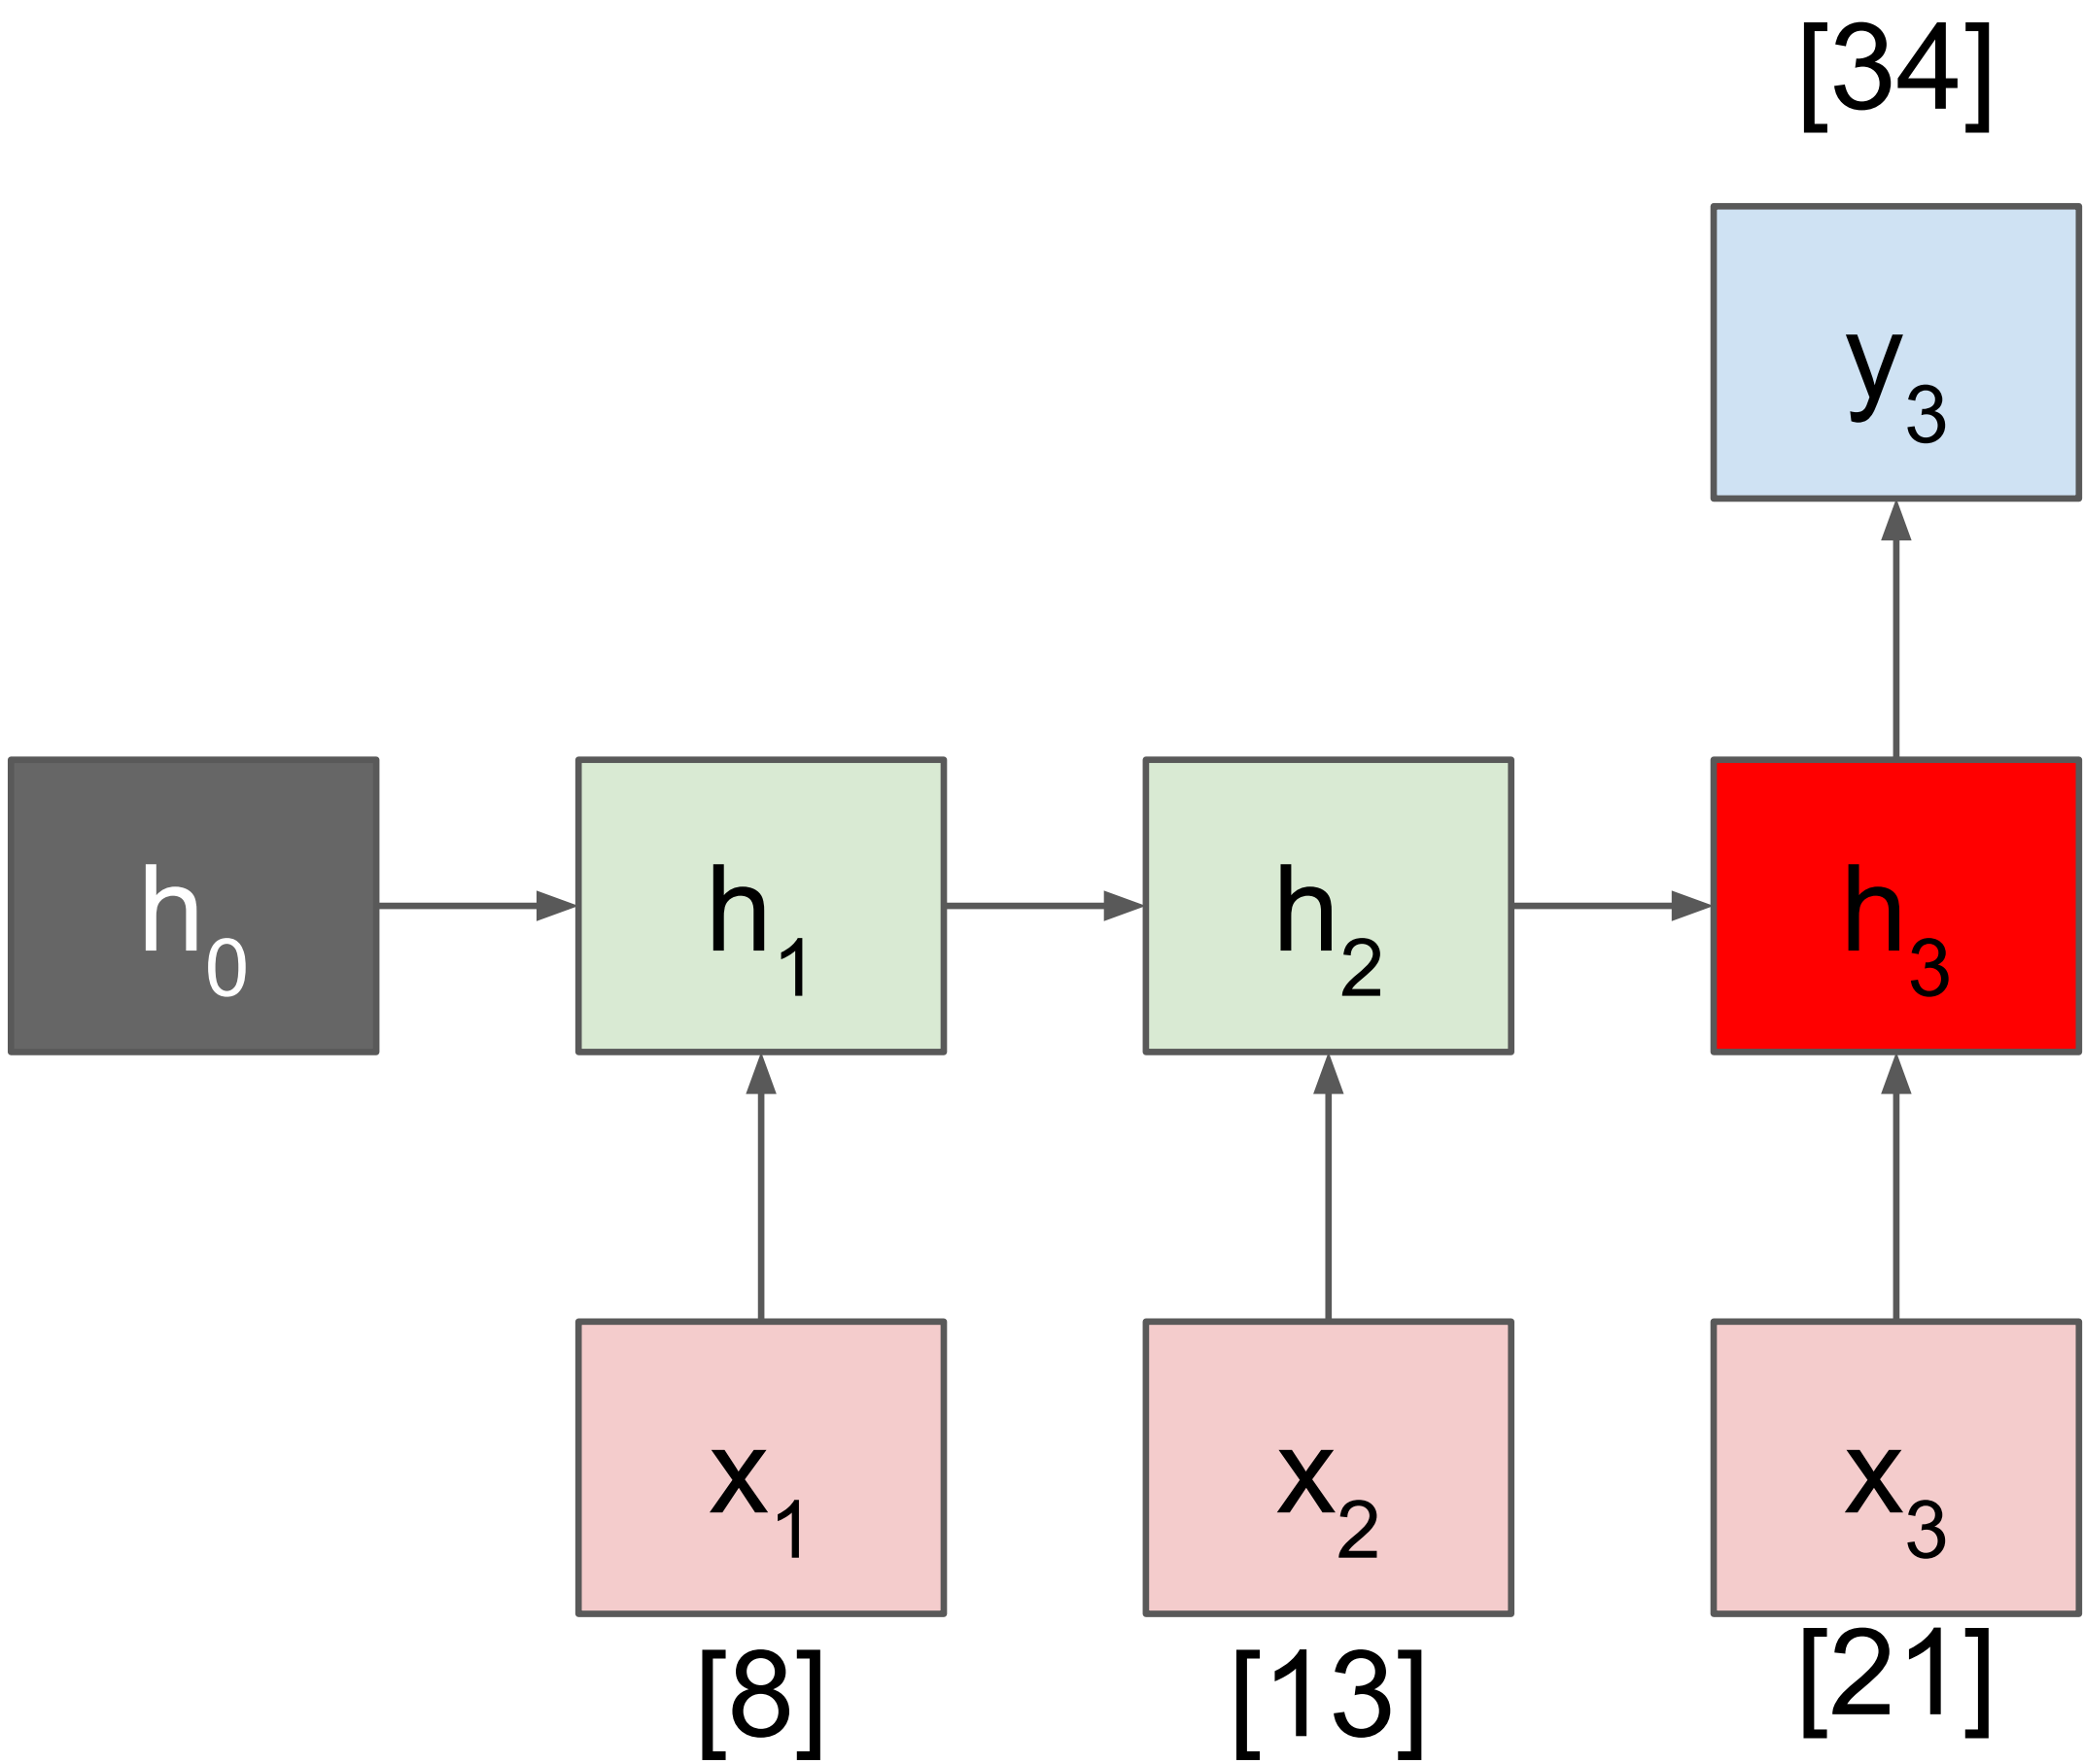
Traditional Seq2Seq model **discard all the intermediate states of the encoder** and use **only its final states** (vector) to initialize the decoder

Illustration: Many to One (Simple Model)

Dataset Fibonacci: 1,2,3, f_{1200}

Target f: $x_1 \ x_2 \ \dots \ x_{ts} \rightarrow x_{ts+1}$

```
modelRNN = Sequential()  
modelRNN.add(SimpleRNN(hidden_units, input_shape=(ts,1),  
activation='tanh'))  
modelRNN.add(Dense(units=dense_units, activation='linear'))  
modelRNN.compile(loss='mse', optimizer='adam')  
modelRNN.fit(trainX, trainY, epochs=epochs)
```



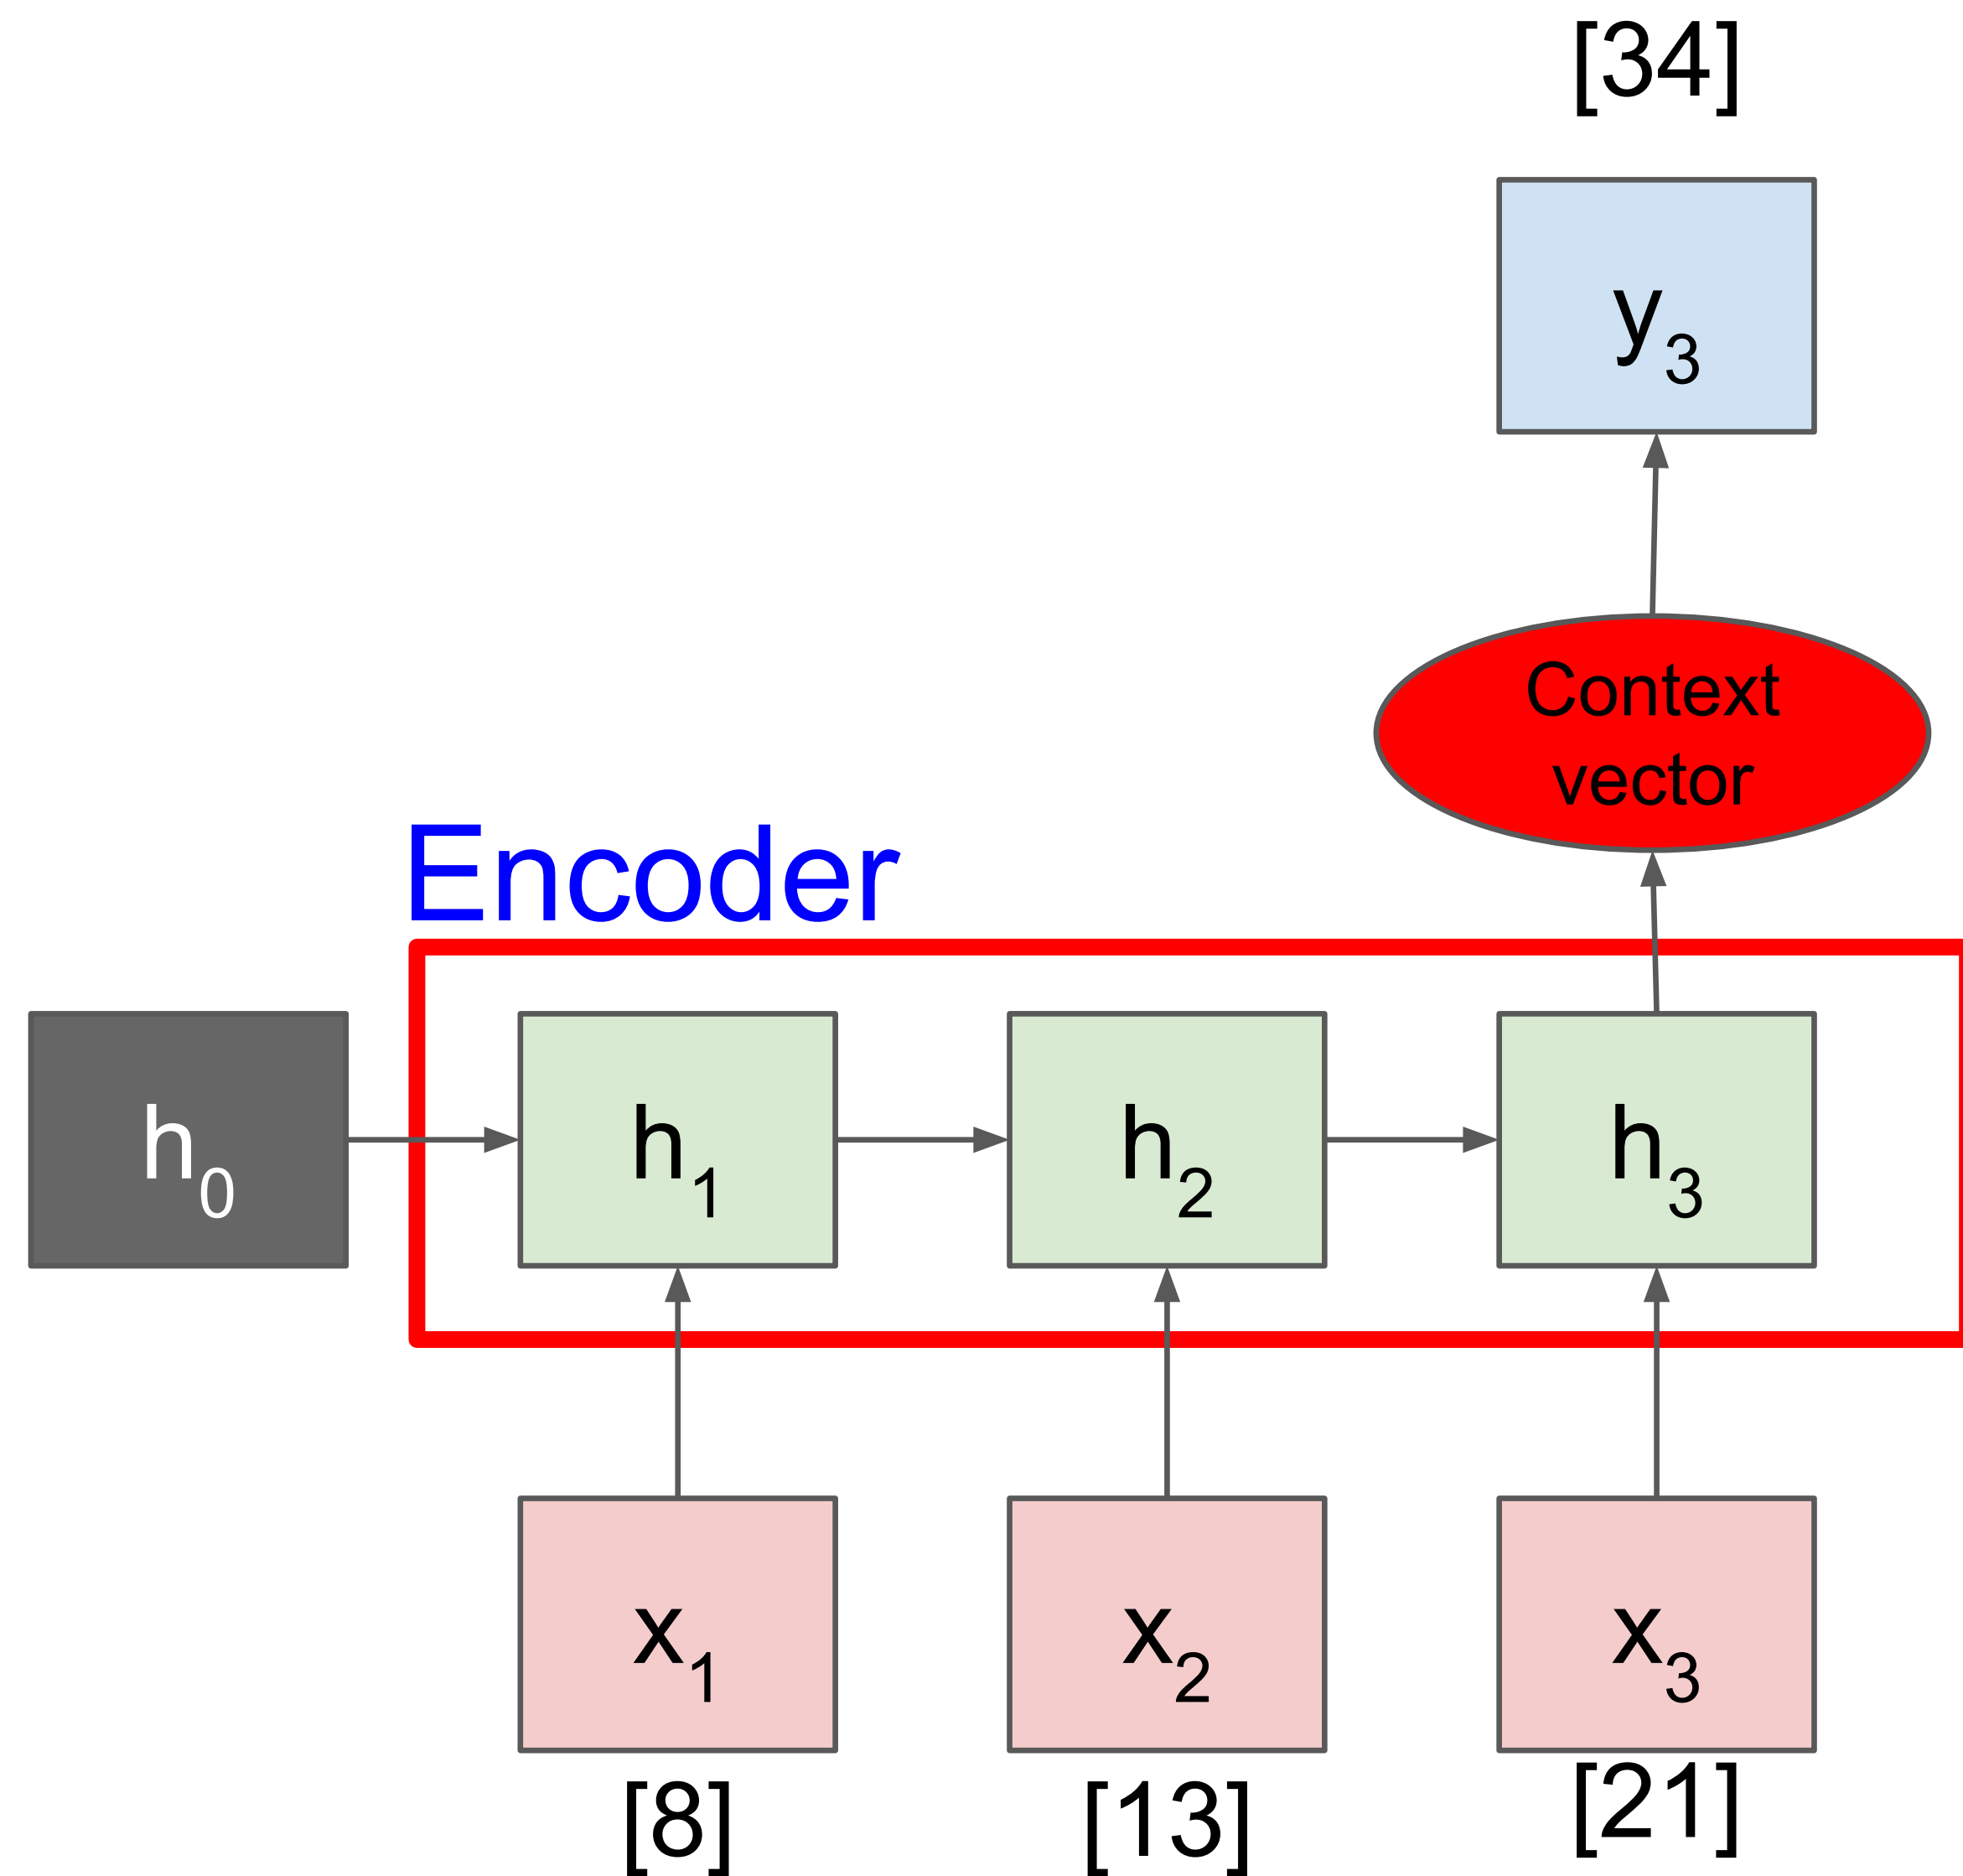
Model: "sequential_49"

Layer (type)	Output Shape	Param #
=====	=====	=====
simple_rnn_52 (SimpleRNN)	(None, 2)	8
dense_62 (Dense)	(None, 1)	3
=====	=====	=====
Total params: 11 (44.00 Byte)		
Trainable params: 11 (44.00 Byte)		
Non-trainable params: 0 (0.00 Byte)		

Hidden_units =2
Dense_units = 1

#Par = (1+2+1)*2+(2+1)*1=11

Predict Next Element in Series



Input: a sequence of vectors

$$\mathbf{x} = (x_1, \dots, x_{T_x})$$

Hidden state at time t:

$$h_t = f(x_t, h_{t-1})$$

Context vector: a vector generated from the sequence of the hidden states.


$$c = q(\{h_1, \dots, h_{T_x}\})$$

$$q(\{h_1, \dots, h_T\}) = h_T$$

Encoder - Decoder: Machine Translation

Dataset: parallel corpus

Target f: $x_1 x_2 \dots x_m \rightarrow y_1 y_2 \dots y_n$



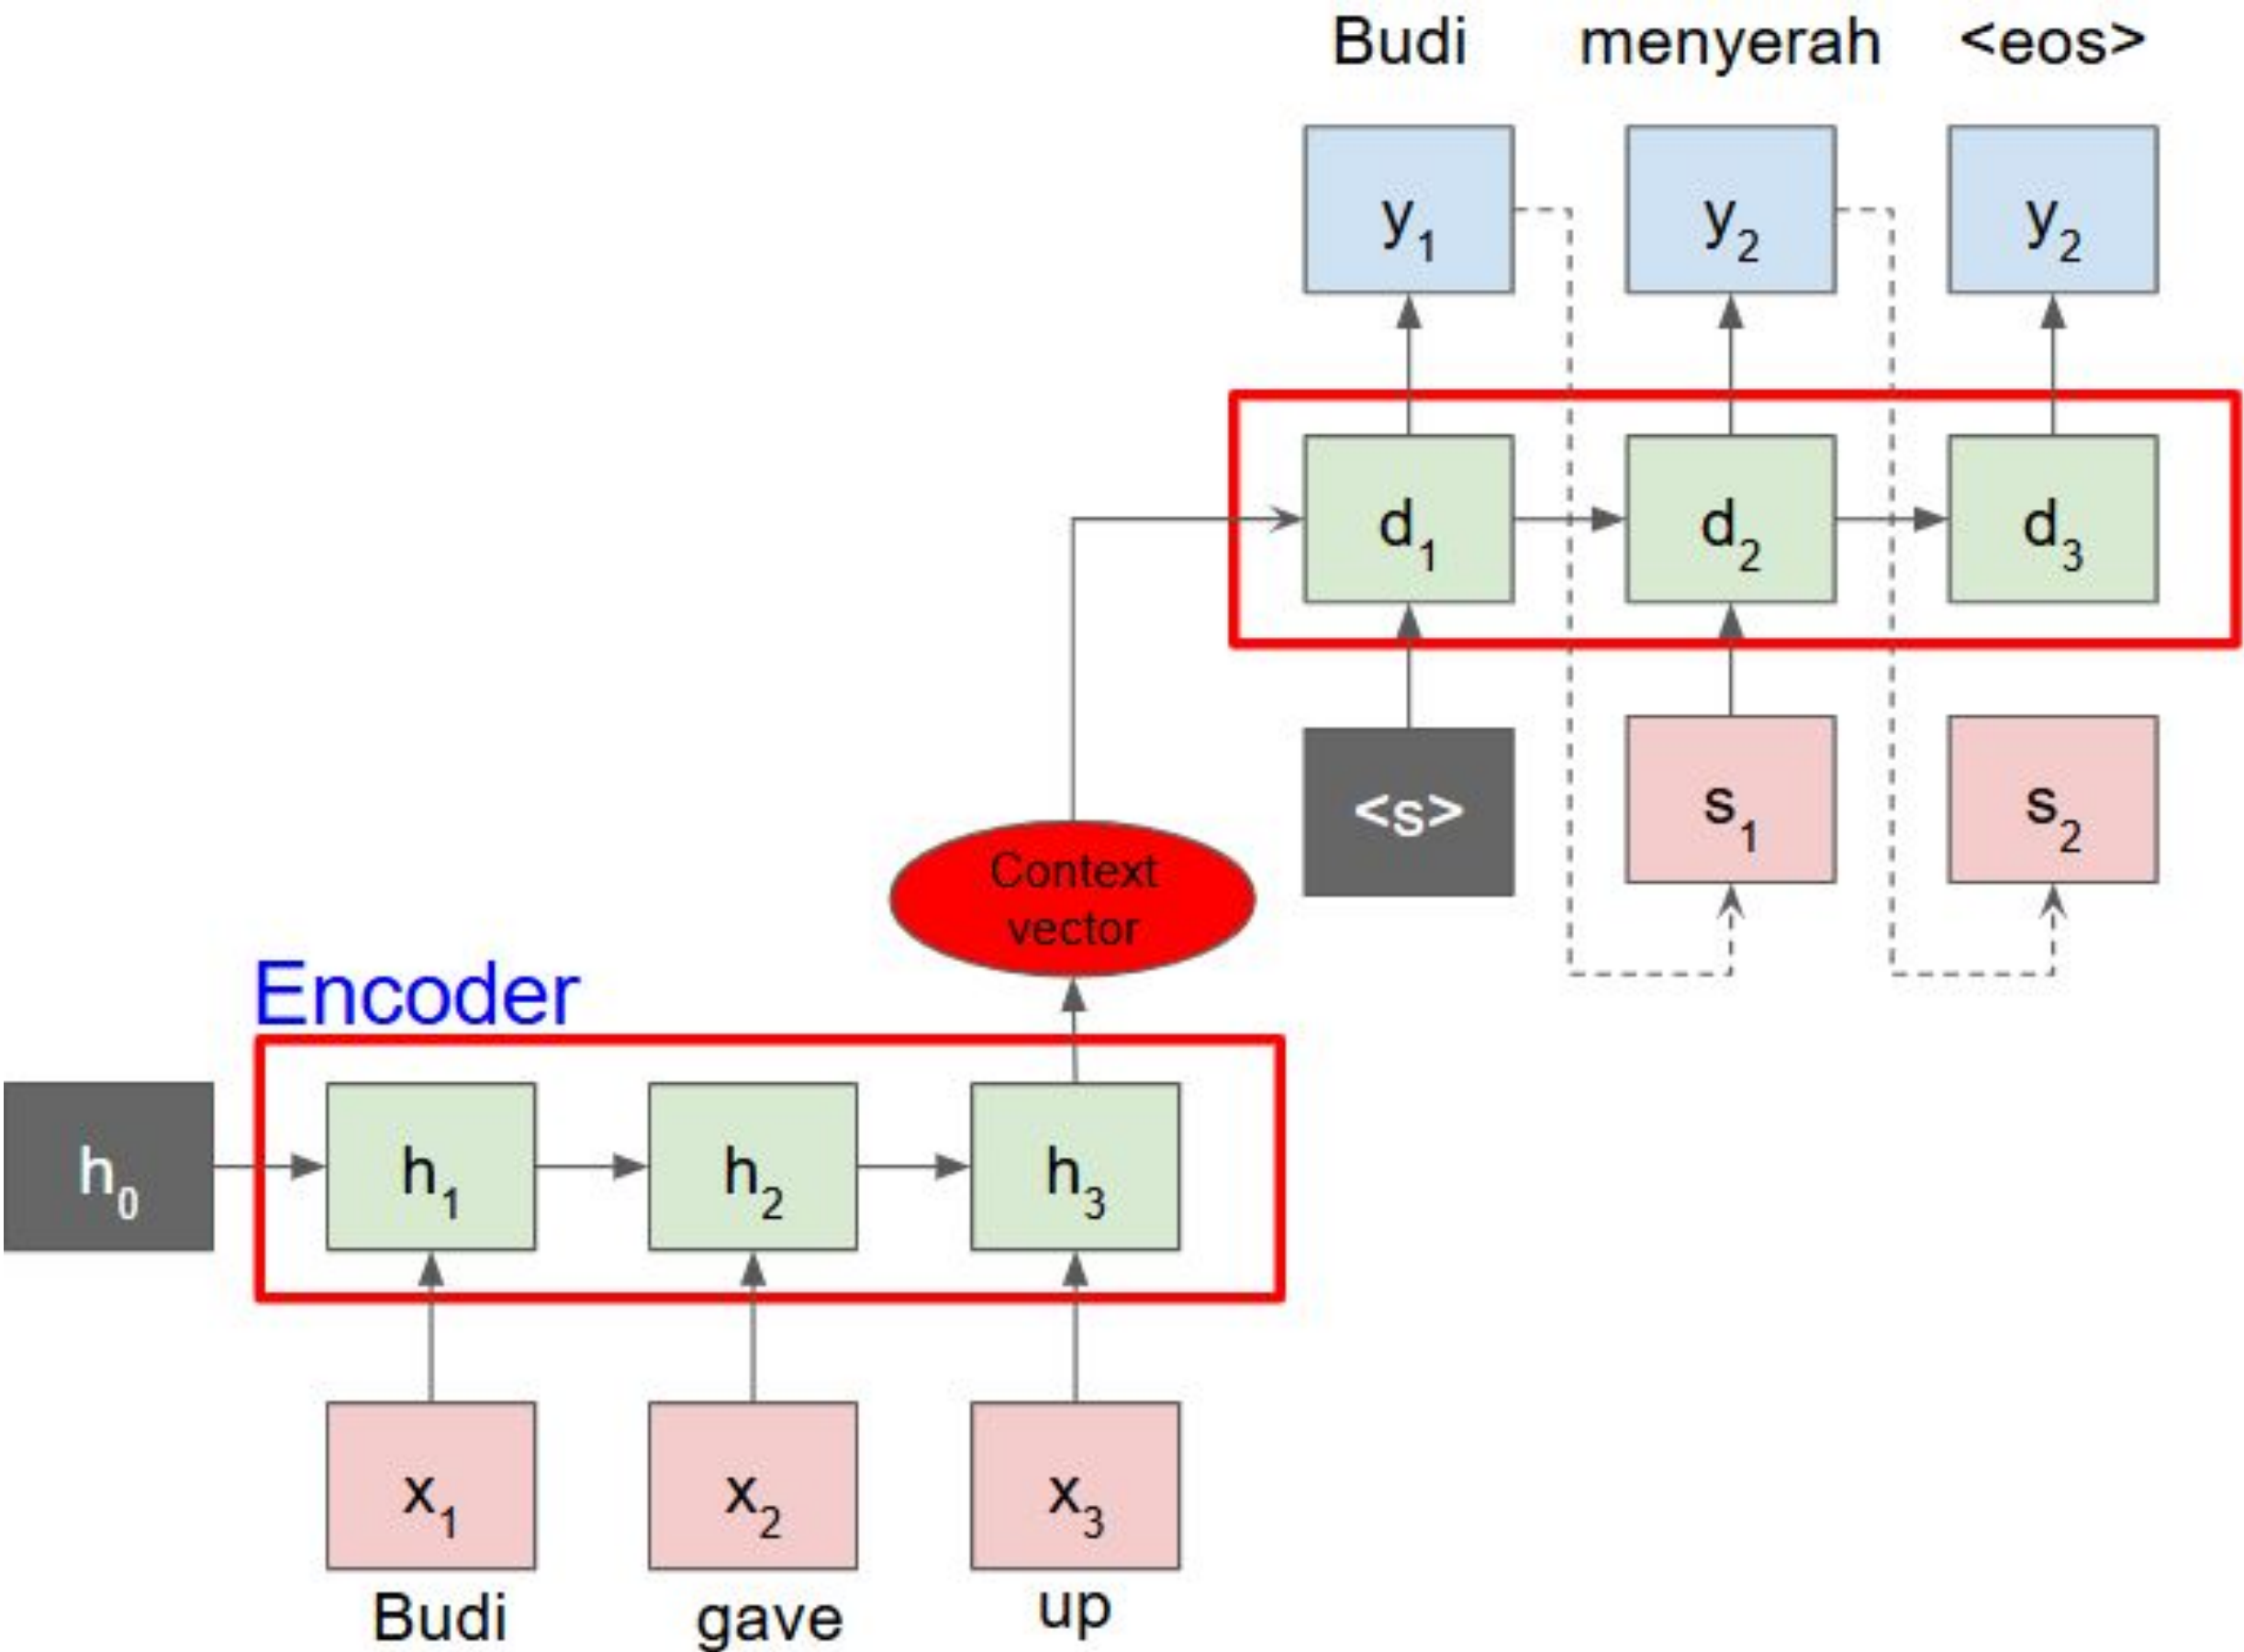
Decoder is trained to predict the **next word** y_t given the **context vector** c and all the previously predicted words $\{y_1, \dots, y_{t-1}\}$.

$$p(y_t | \{y_1, \dots, y_{t-1}\}, c) = g(y_{t-1}, s_t, c)$$

The decoder defines a probability over the translation y by decomposing the joint probability into the ordered conditionals:

$$p(\mathbf{y}) = \prod_{t=1}^T p(y_t | \{y_1, \dots, y_{t-1}\}, c)$$

Encoder - Decoder Architecture Example



Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
input_3 (InputLayer)	[(None, None, 74)]	0	[]
input_4 (InputLayer)	[(None, None, 78)]	0	[]
lstm_2 (LSTM)	[(None, 256), (None, 256), (None, 256)]	338944	['input_3[0][0]']
lstm_3 (LSTM)	[(None, None, 256), (None, 256), (None, 256)]	343040	['input_4[0][0]', 'lstm_2[0][1]', 'lstm_2[0][2]']
dense_1 (Dense)	(None, None, 78)	20046	['lstm_3[0][0]']

=====

Total params: 702030 (2.68 MB)
Trainable params: 702030 (2.68 MB)
Non-trainable params: 0 (0.00 Byte)

=====

Encoder - Decoder: Weakness & Solution

- The **final hidden state** of the **encoder** creates an **information bottleneck**. It has to capture the meaning of the whole input sequence because this is all the **decoder** has access to when **generating the output**.
 - especially challenging for long sequences
- Alternative solution: allowing the **decoder** to have access to **all** of the **encoder's hidden states**.
 - The general mechanism for this is called **attention** and is a key component in many modern neural network architectures.

ATTENTION!

Central idea behind Attention is **not to throw away those intermediate encoder states** but to utilize all the states in order to **construct the context vectors** required by the decoder to generate the output sequence

Bahdanau, D., Cho, K. H., & Bengio, Y. (2015, January). Neural machine translation by jointly learning to align and translate. In *3rd International Conference on Learning Representations, ICLR 2015*.

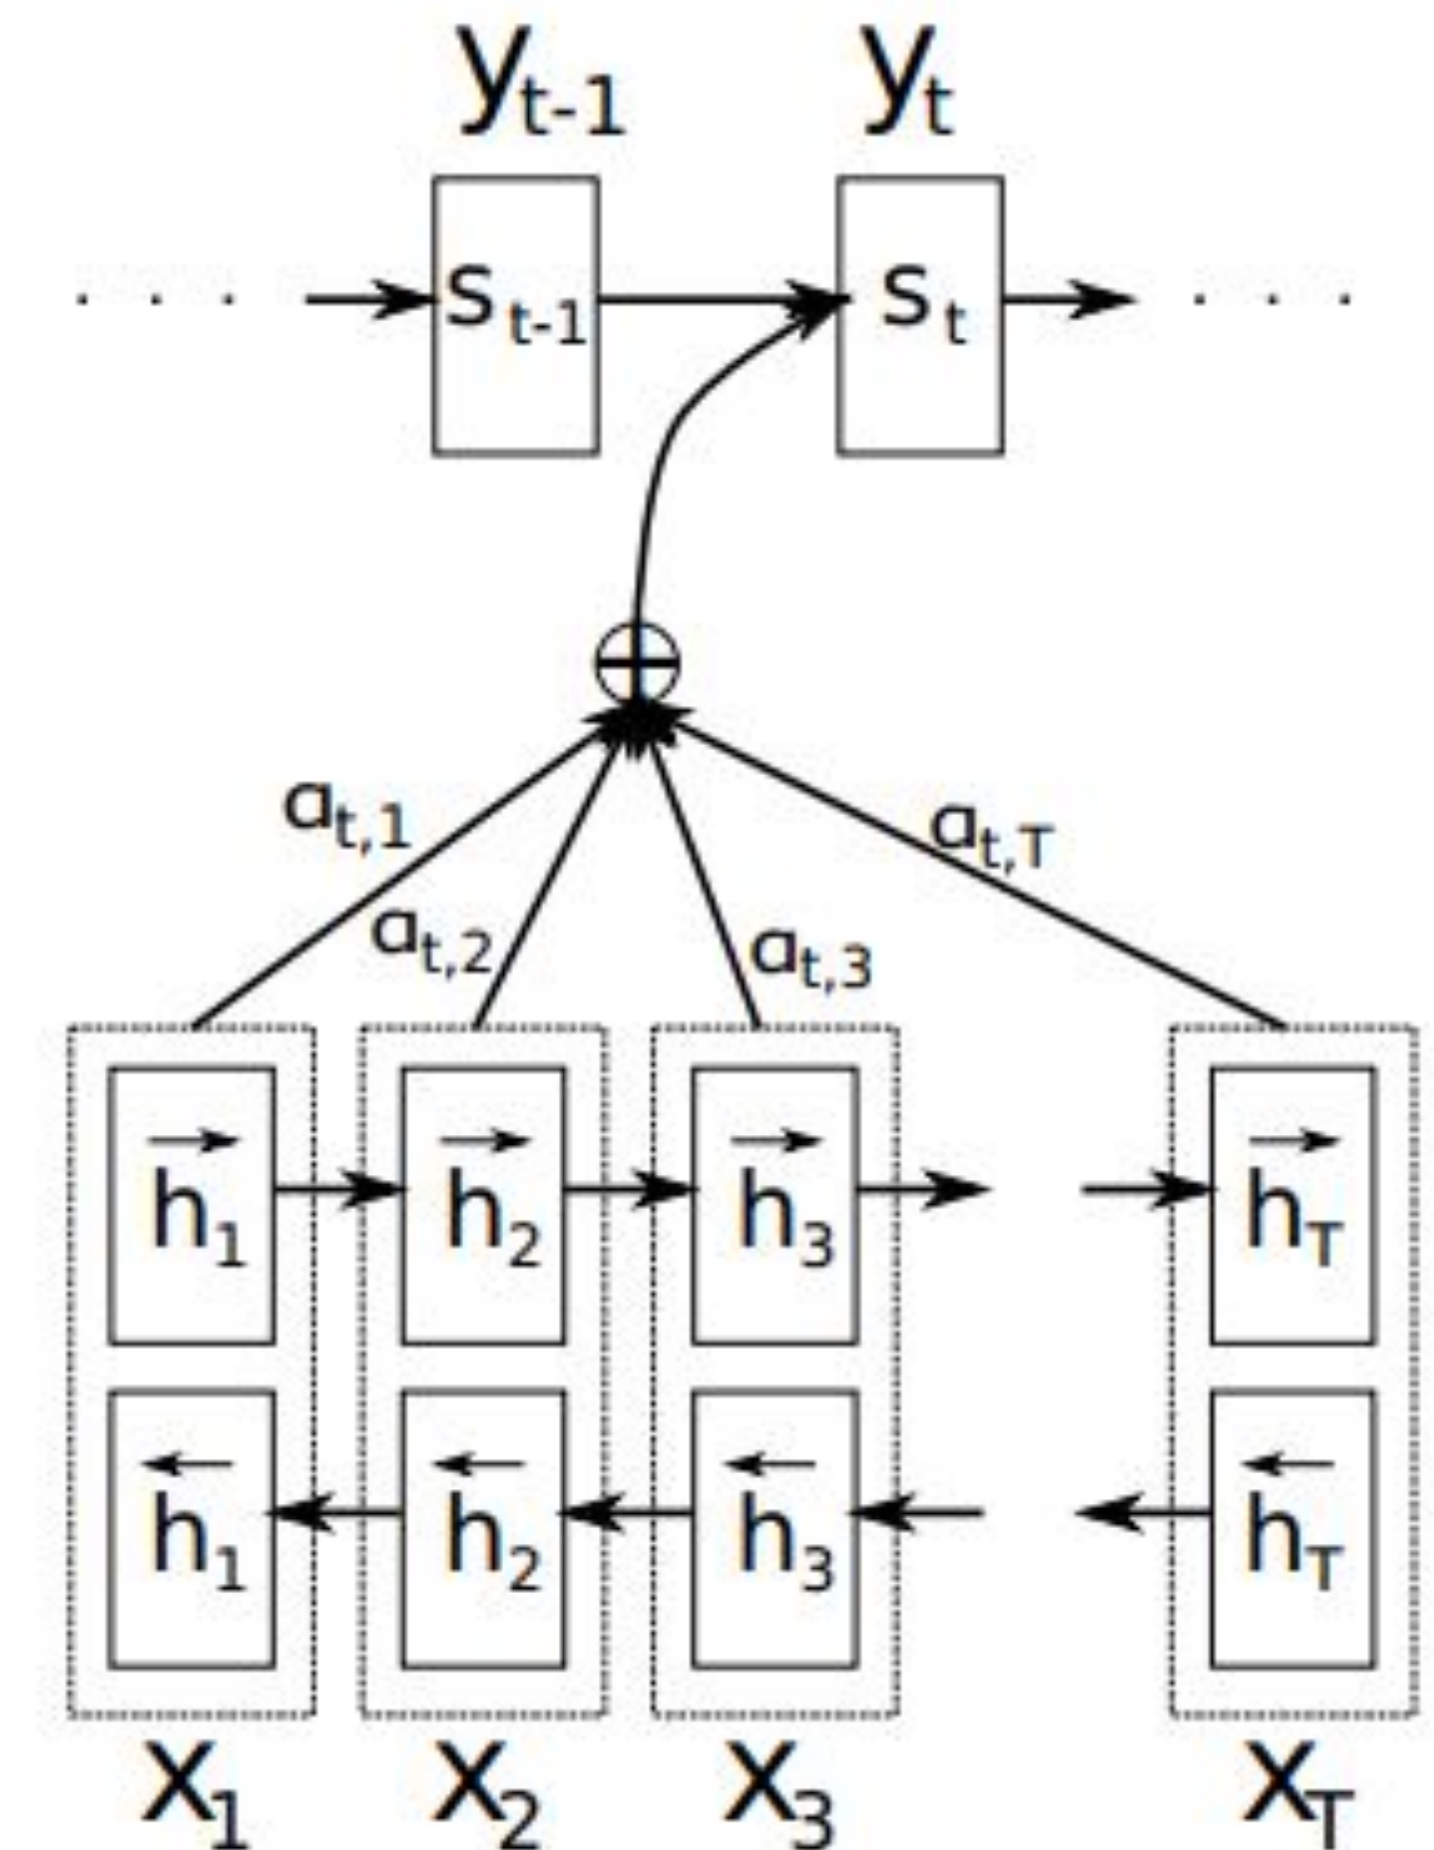
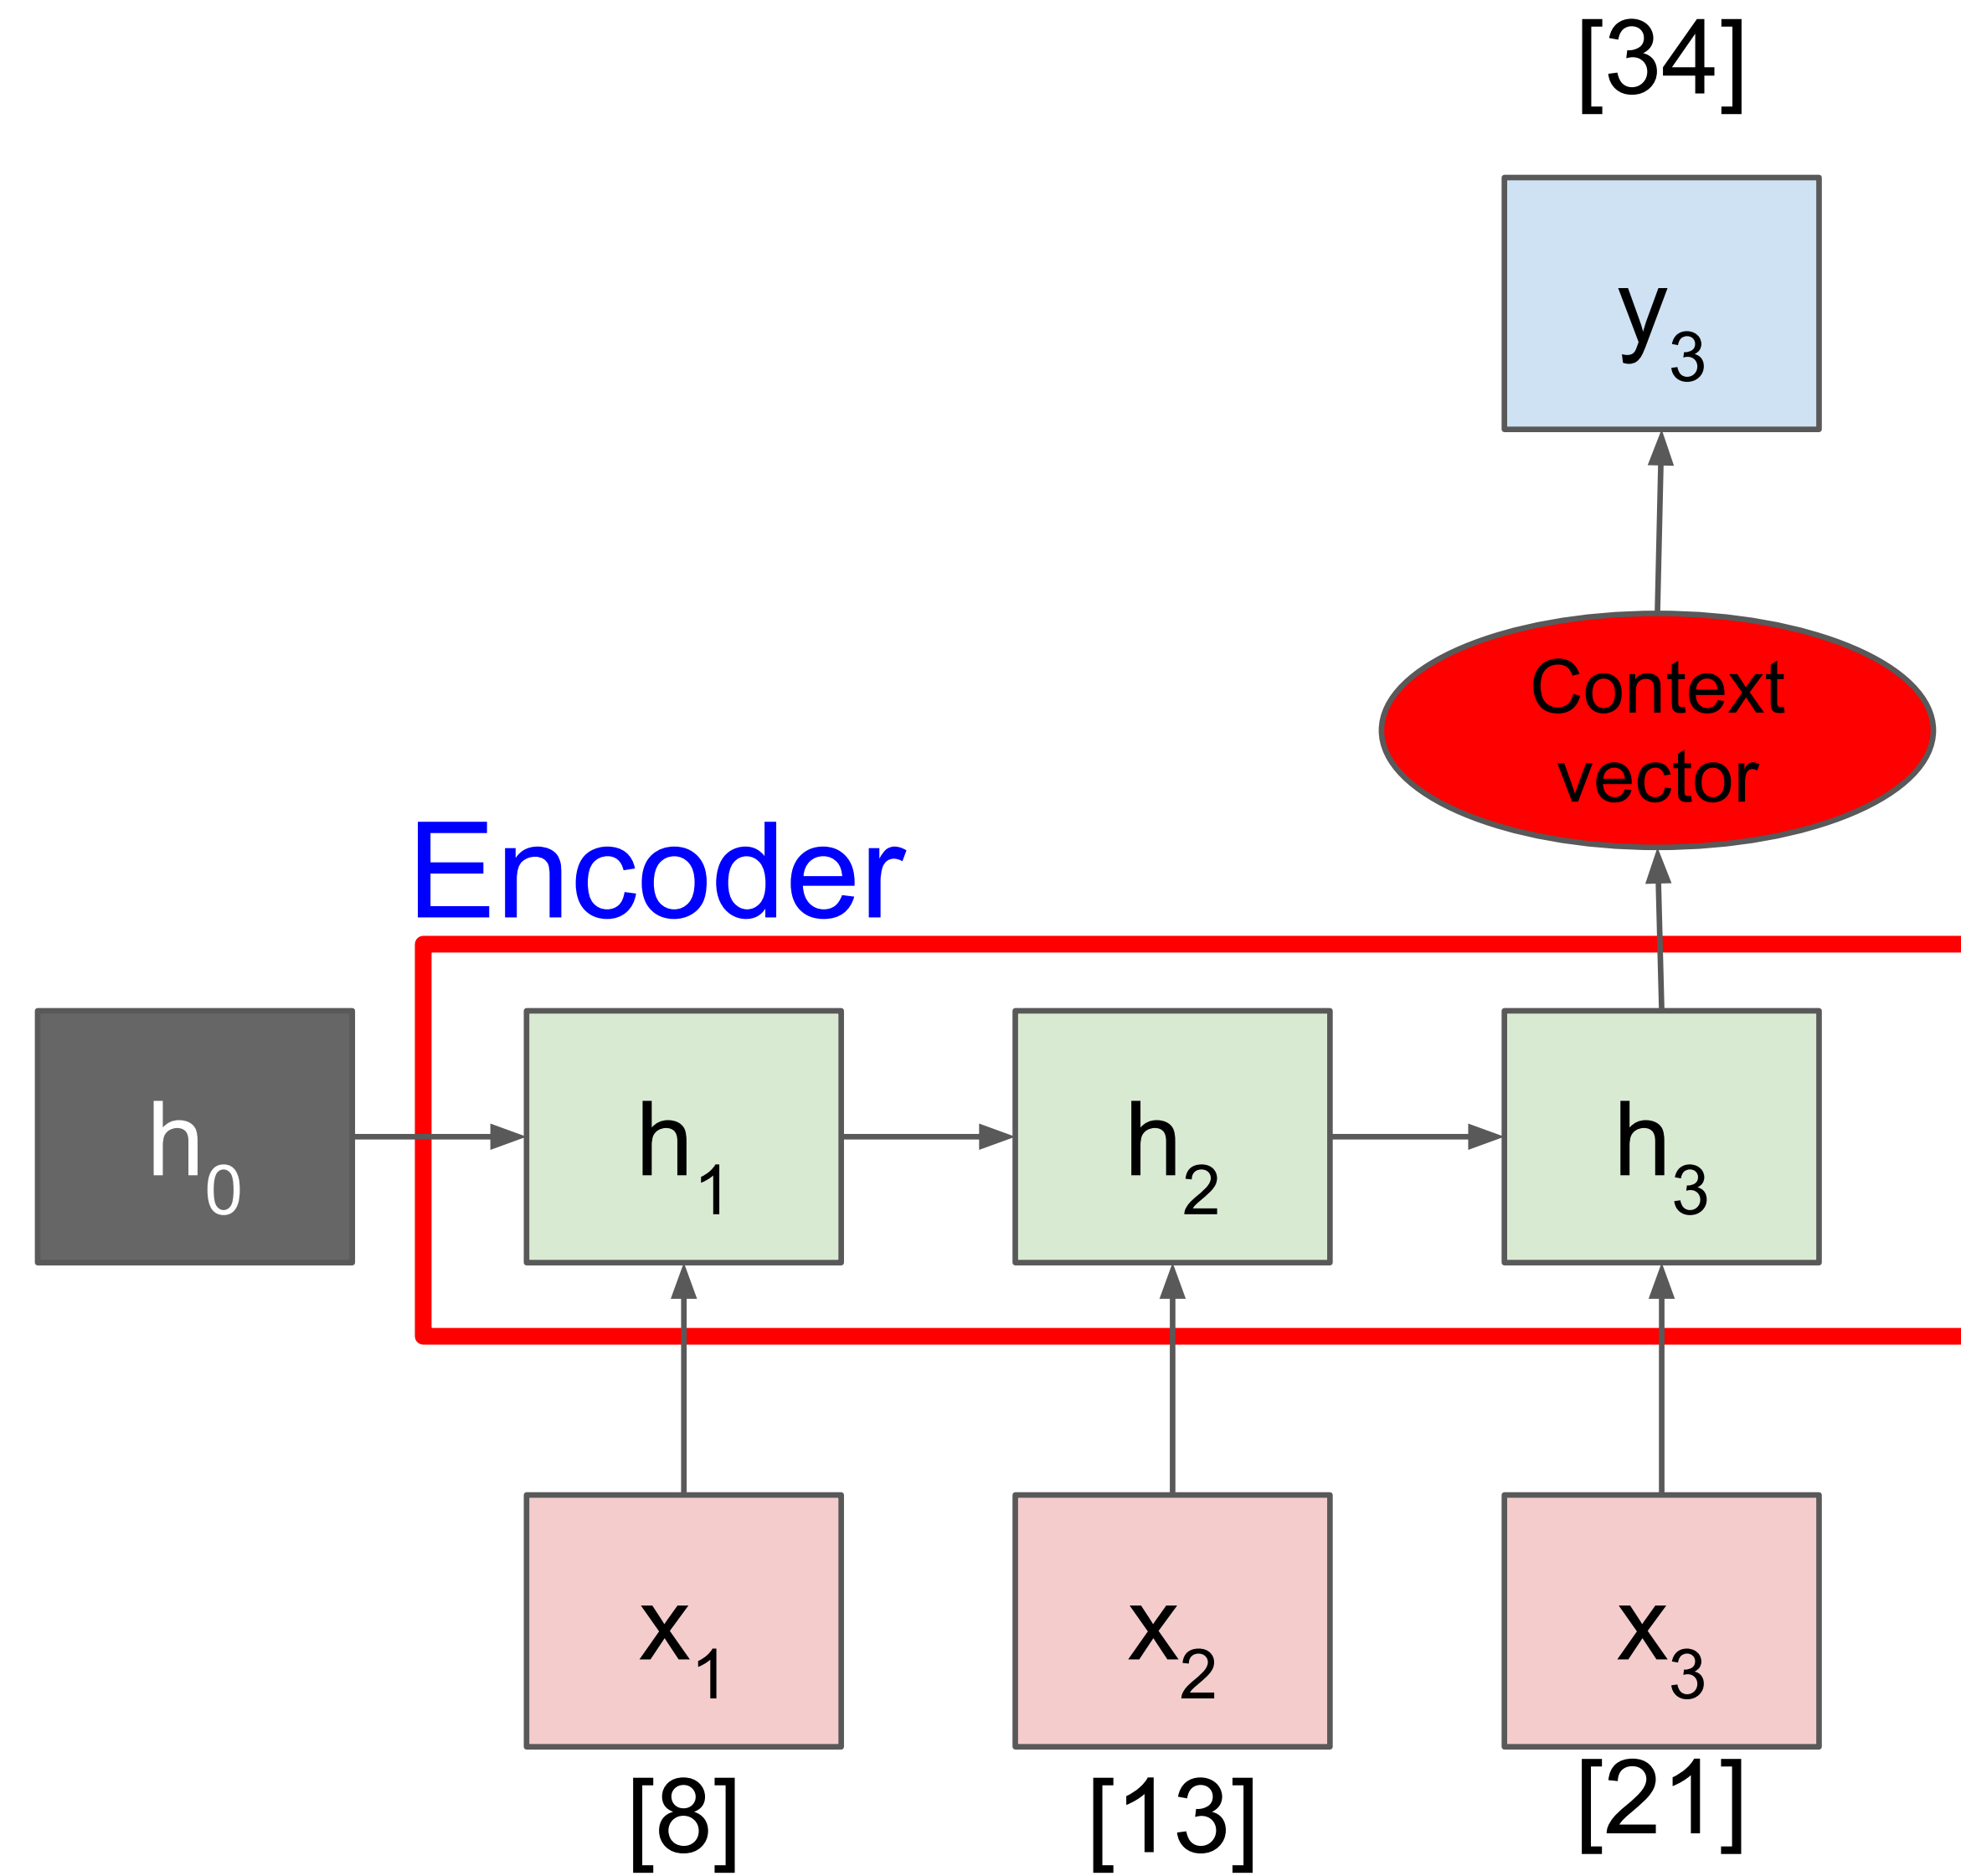
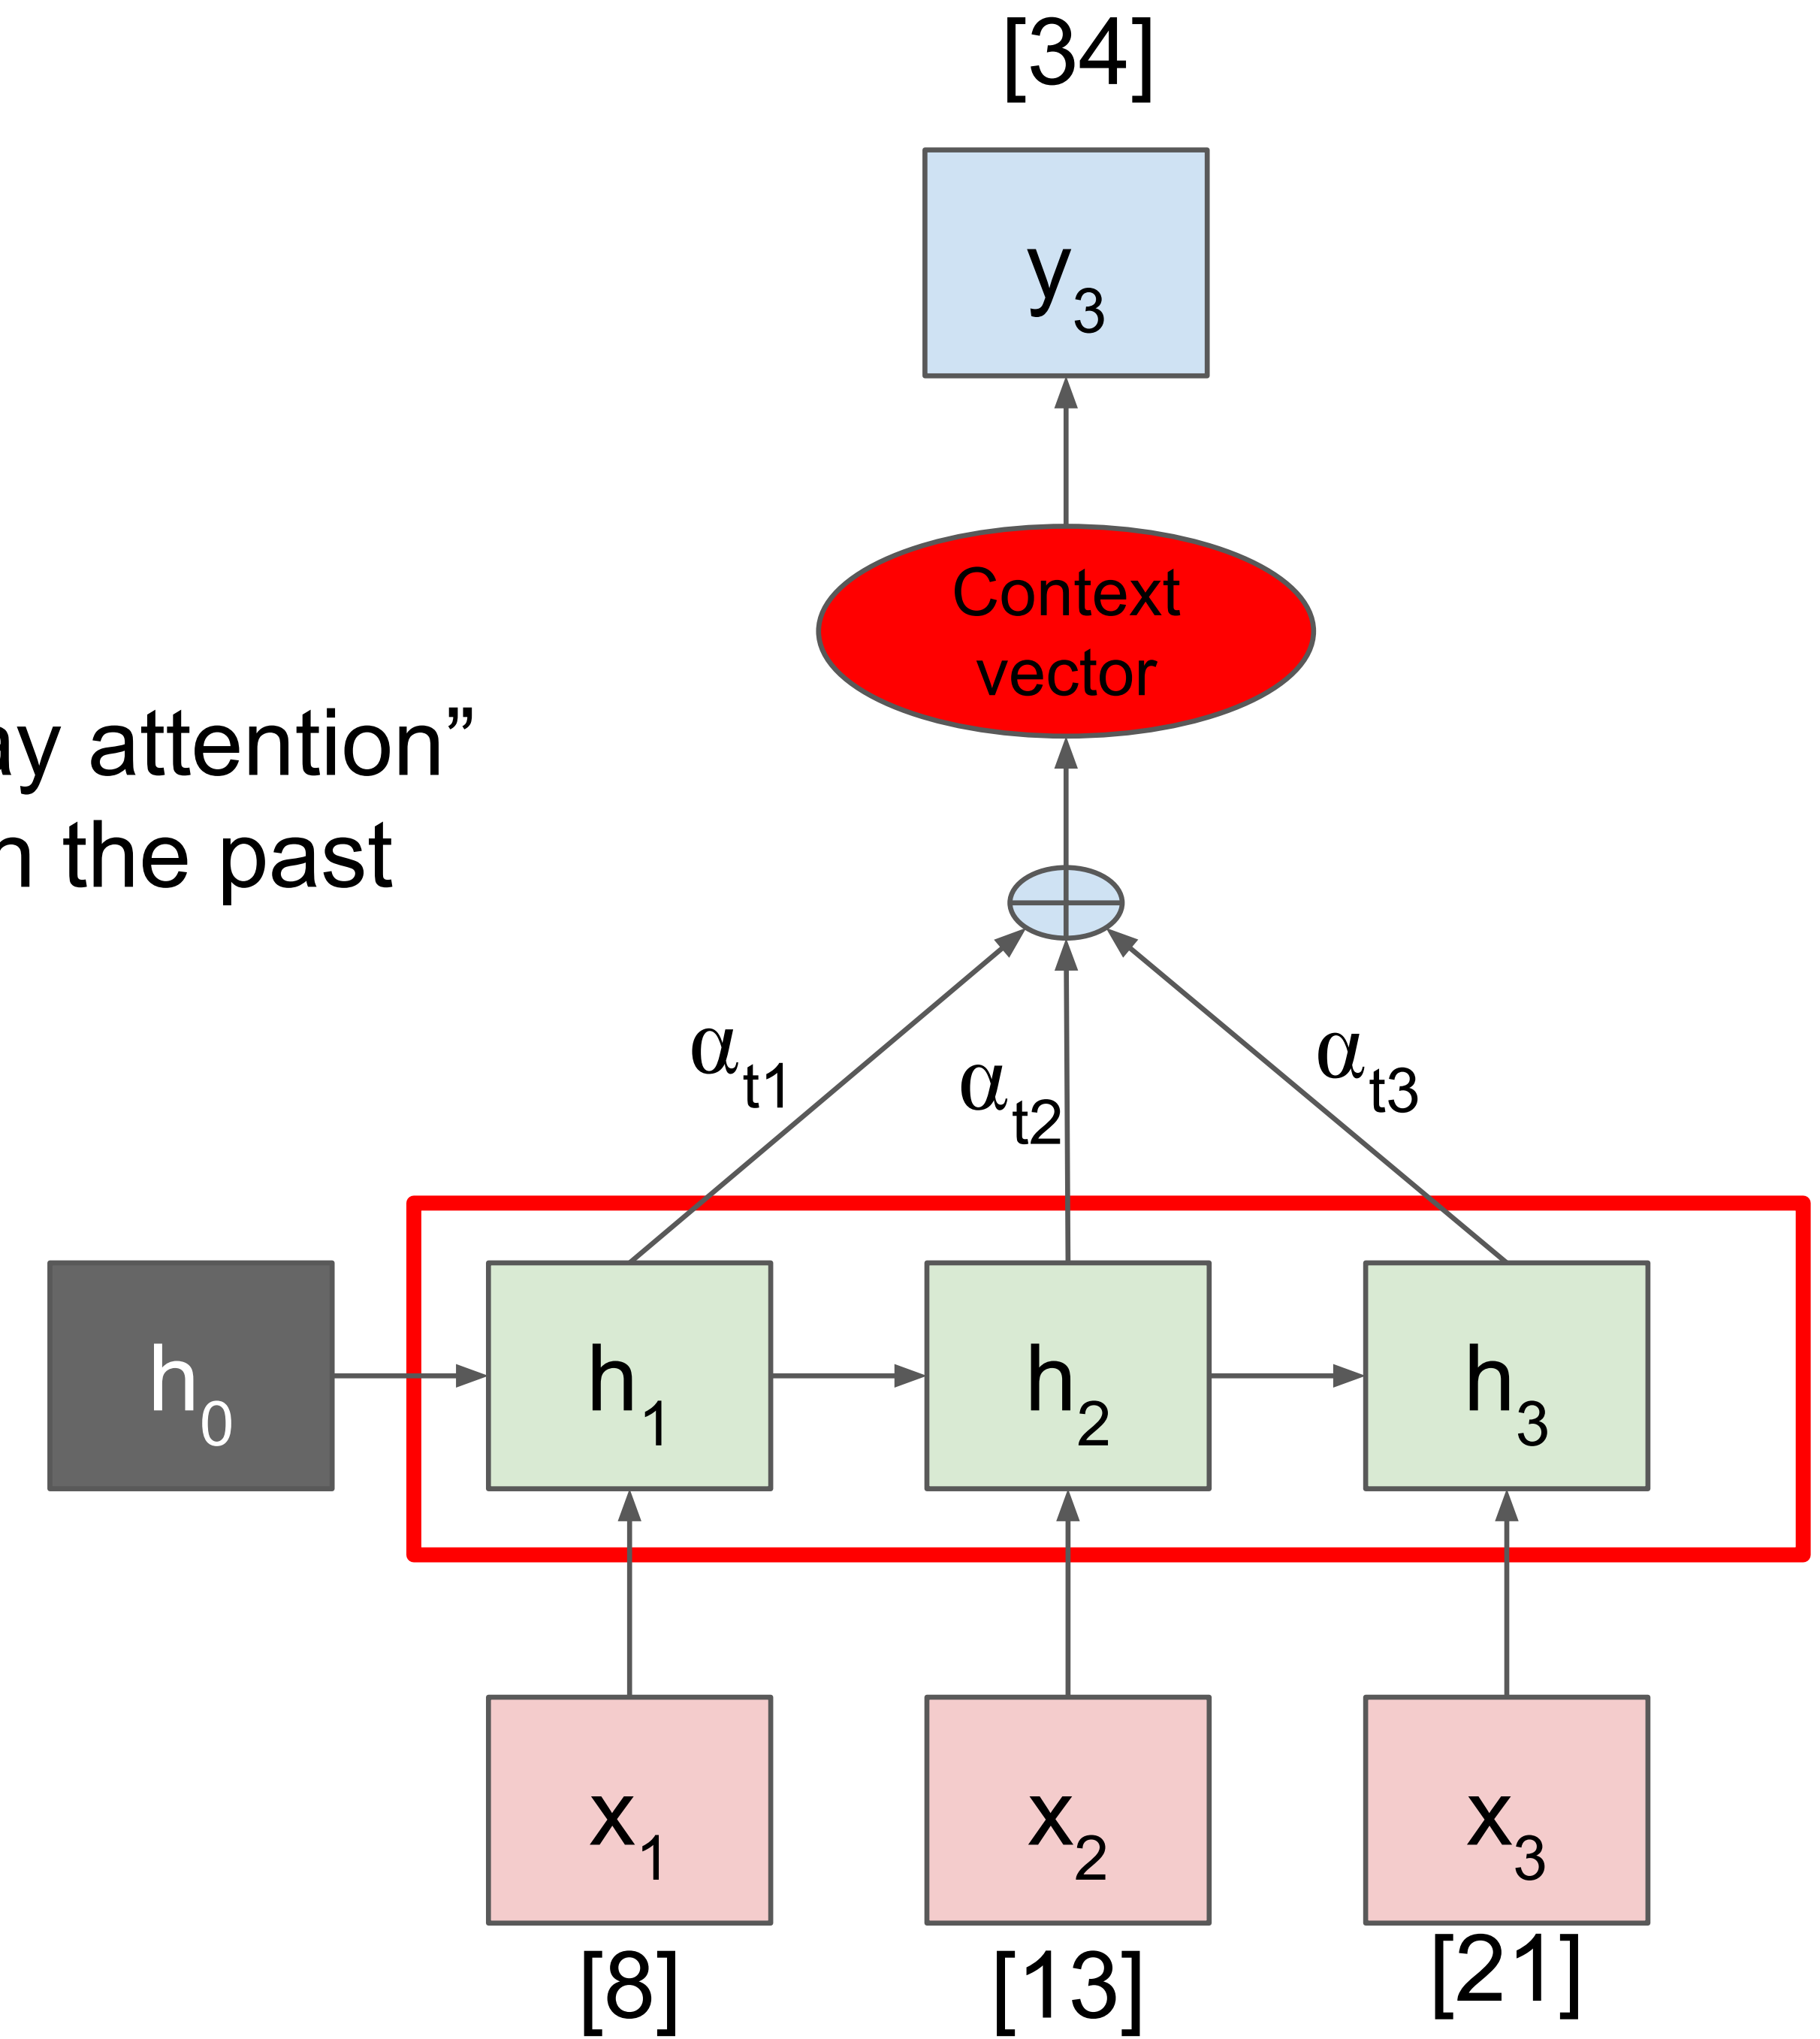



Figure 1: The graphical illustration of the proposed model trying to generate the t -th target word y_t given a source sentence (x_1, x_2, \dots, x_T) .

Many to One with Attention



Assign a weight or “pay attention”
to the specific states in the past



Many to One with Attention

Input: $\mathbf{x} = (x_1, \dots, x_{T_x})$

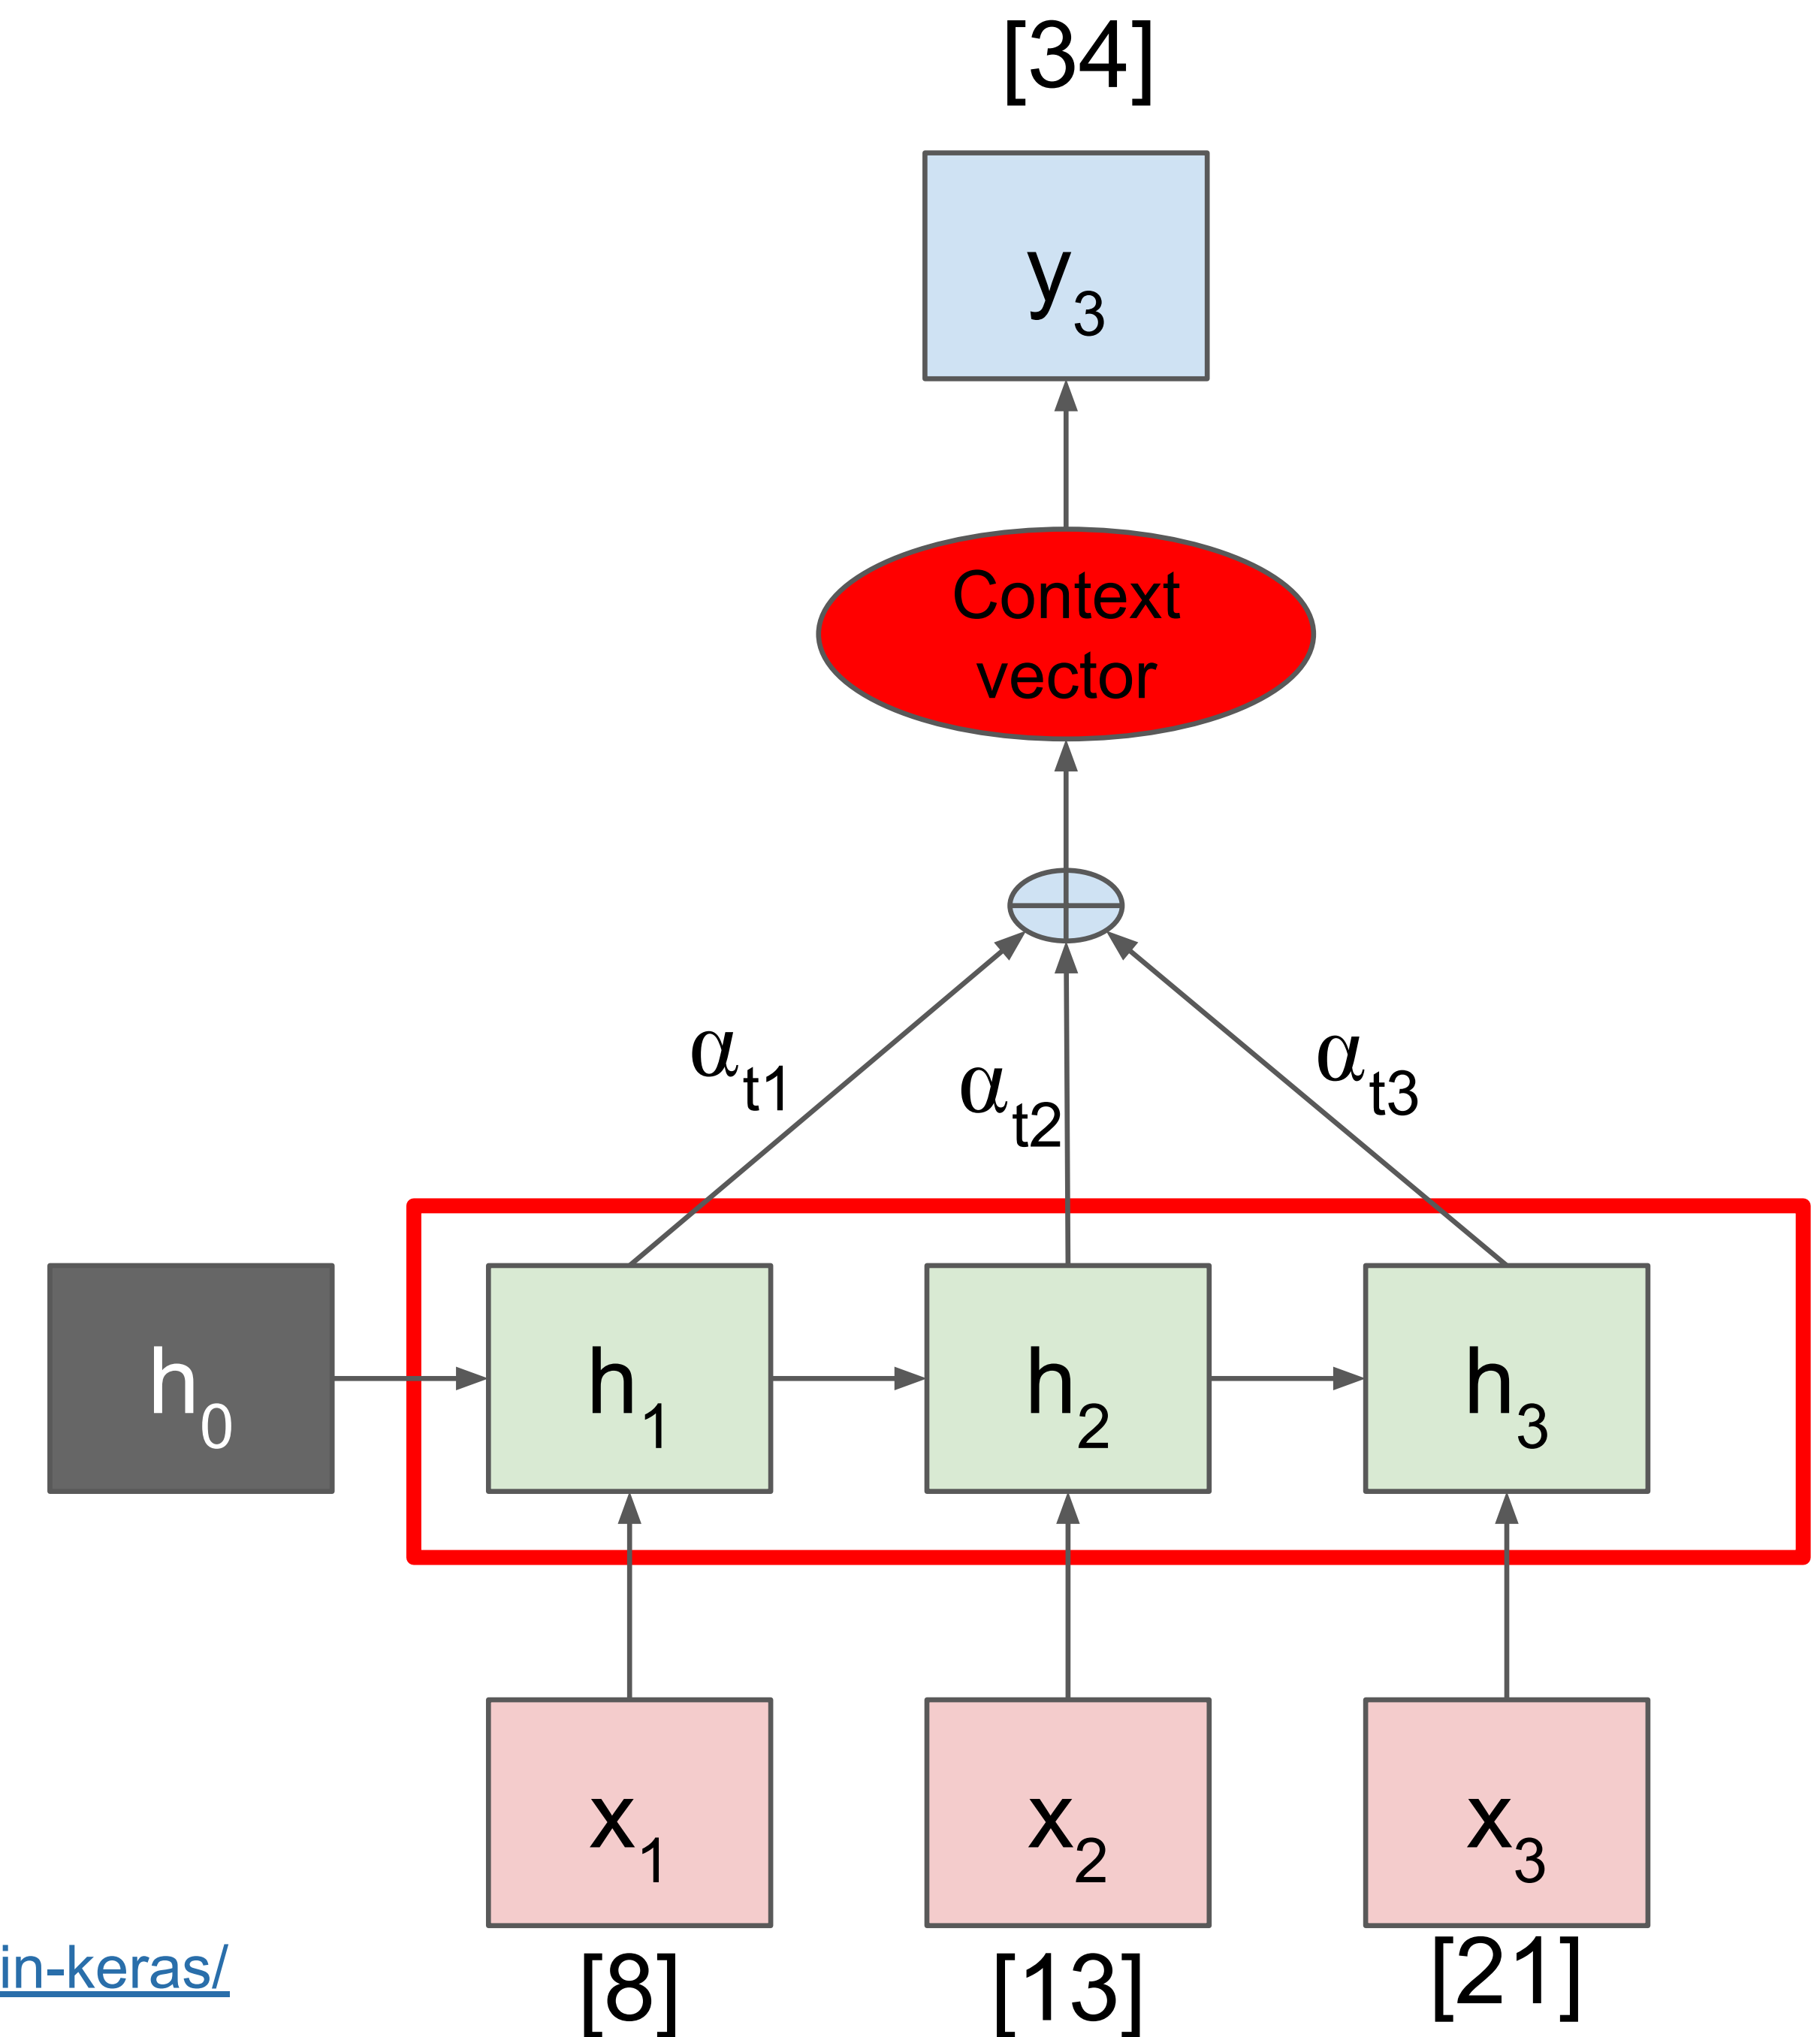
Hidden state at time t: $h_t = f(x_t, h_{t-1})$

Context vector c_i for each target y_i :

$$c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j.$$

For Fibonacci series, the mean square error on the test set is lower with the attention layer.

<https://machinelearningmastery.com/adding-a-custom-attention-layer-to-recurrent-neural-network-in-keras/>




Many to One with Attention Architecture Example

Model: "sequential_49"

Layer (type)	Output Shape	Param #
simple_rnn_52 (SimpleRNN)	(None, 2)	8
dense_62 (Dense)	(None, 1)	3

=====
Total params: 11 (44.00 Byte)
Trainable params: 11 (44.00 Byte)
Non-trainable params: 0 (0.00 Byte)



Model: "sequential_63"

Layer (type)	Output Shape	Param #
simple_rnn_66 (SimpleRNN)	(None, 20, 2)	8
attention_11 (attention)	(None, 2)	22
dense_76 (Dense)	(None, 1)	3

=====
Total params: 33 (132.00 Byte)
Trainable params: 33 (132.00 Byte)
Non-trainable params: 0 (0.00 Byte)

Hidden_units =2

Dense_units = 1

Time_steps = 20

#Params = (1+2+1)*2+(2+1)*1=11

Hidden_units =2

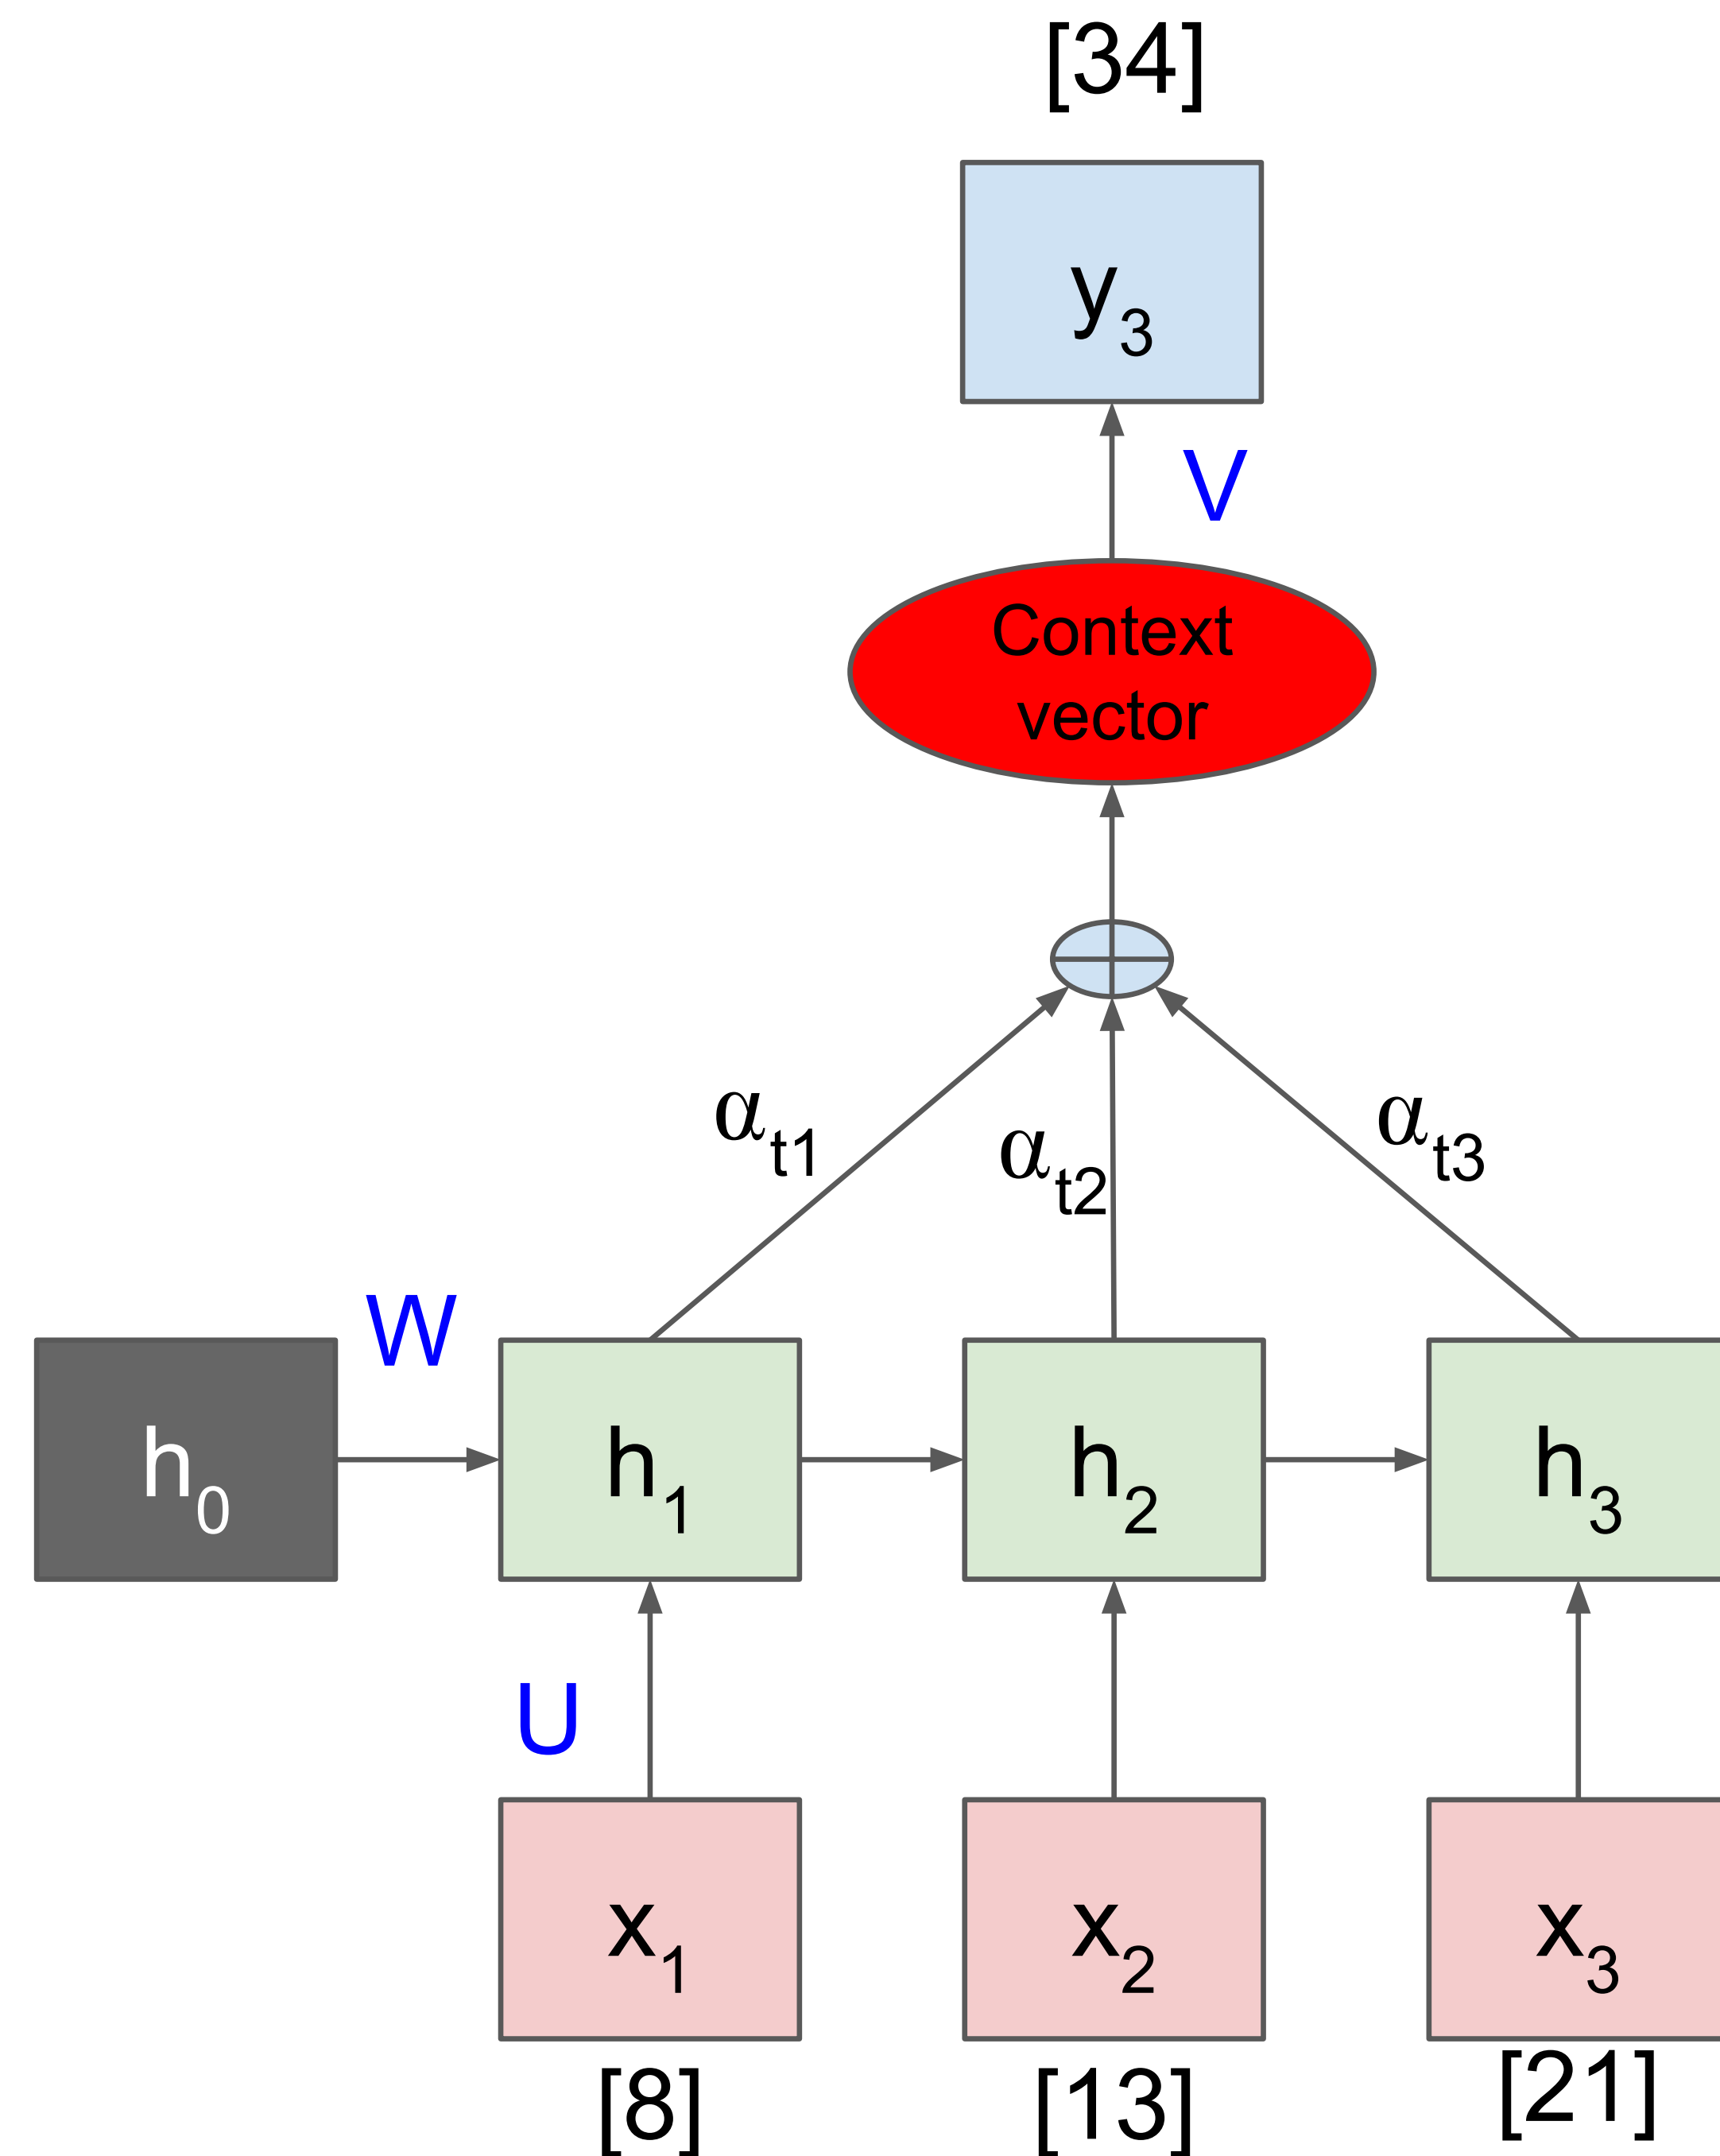
Dense_units = 1

Time_steps = 20

Attention_units=1

- Attention weight: 20 $\alpha_{t1} \dots \alpha_{t_{ts}}$
Shape: (time_steps, attention_units)
- Attention bias: 2
Shape: (hidden_units, attention_units)

Many to One with Attention: Forward Propagation



$$y_t = \tanh(Vh_t + b_{hy})$$

$$c = \alpha_t h_t + b_\alpha$$

$$h_t = \tanh(Ux_t + (Wh_{t-1} + b_{xh}))$$

RNN with Attention: Context Vector

timestep	Atribut 1	alpha_t	ht_neuron1	ht_neuron2
1	0.0000E+00	-0.6759	0.009	0.404
2	2.2664E-251	-0.6923	0.063	0.087
3	4.5327E-251	-0.7468	0.083	0.270
4	9.0654E-251	-0.7456	0.124	0.102
5	1.5865E-250	-0.8025	0.146	0.186
...
19	1.5330E-247	0.7771	0.341	-0.002
20	2.4805E-247	0.7768	0.344	-0.005

$$c = \alpha_t h_t + b_\alpha$$

$$c_1 = \alpha_t h_{t_n1} + b_{\alpha_n1}$$

$$c_2 = \alpha_t h_{t_n2} + b_{\alpha_n2}$$

$$c = [c1, c2]$$

Encoder Decoder without vs with Attention

Encoder:

$$h_t = f(x_t, h_{t-1})$$

$$h_t = f(x_t, h_{t-1})$$

Context Vector:

$$c = q(\{h_1, \dots, h_{T_x}\})$$
$$q(\{h_1, \dots, h_T\}) = h_T$$

$$c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j.$$
$$\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^{T_x} \exp(e_{ik})}$$
$$e_{ij} = a(s_{i-1}, h_j)$$

the process of “paying attention” can be learned during training

Decoder:

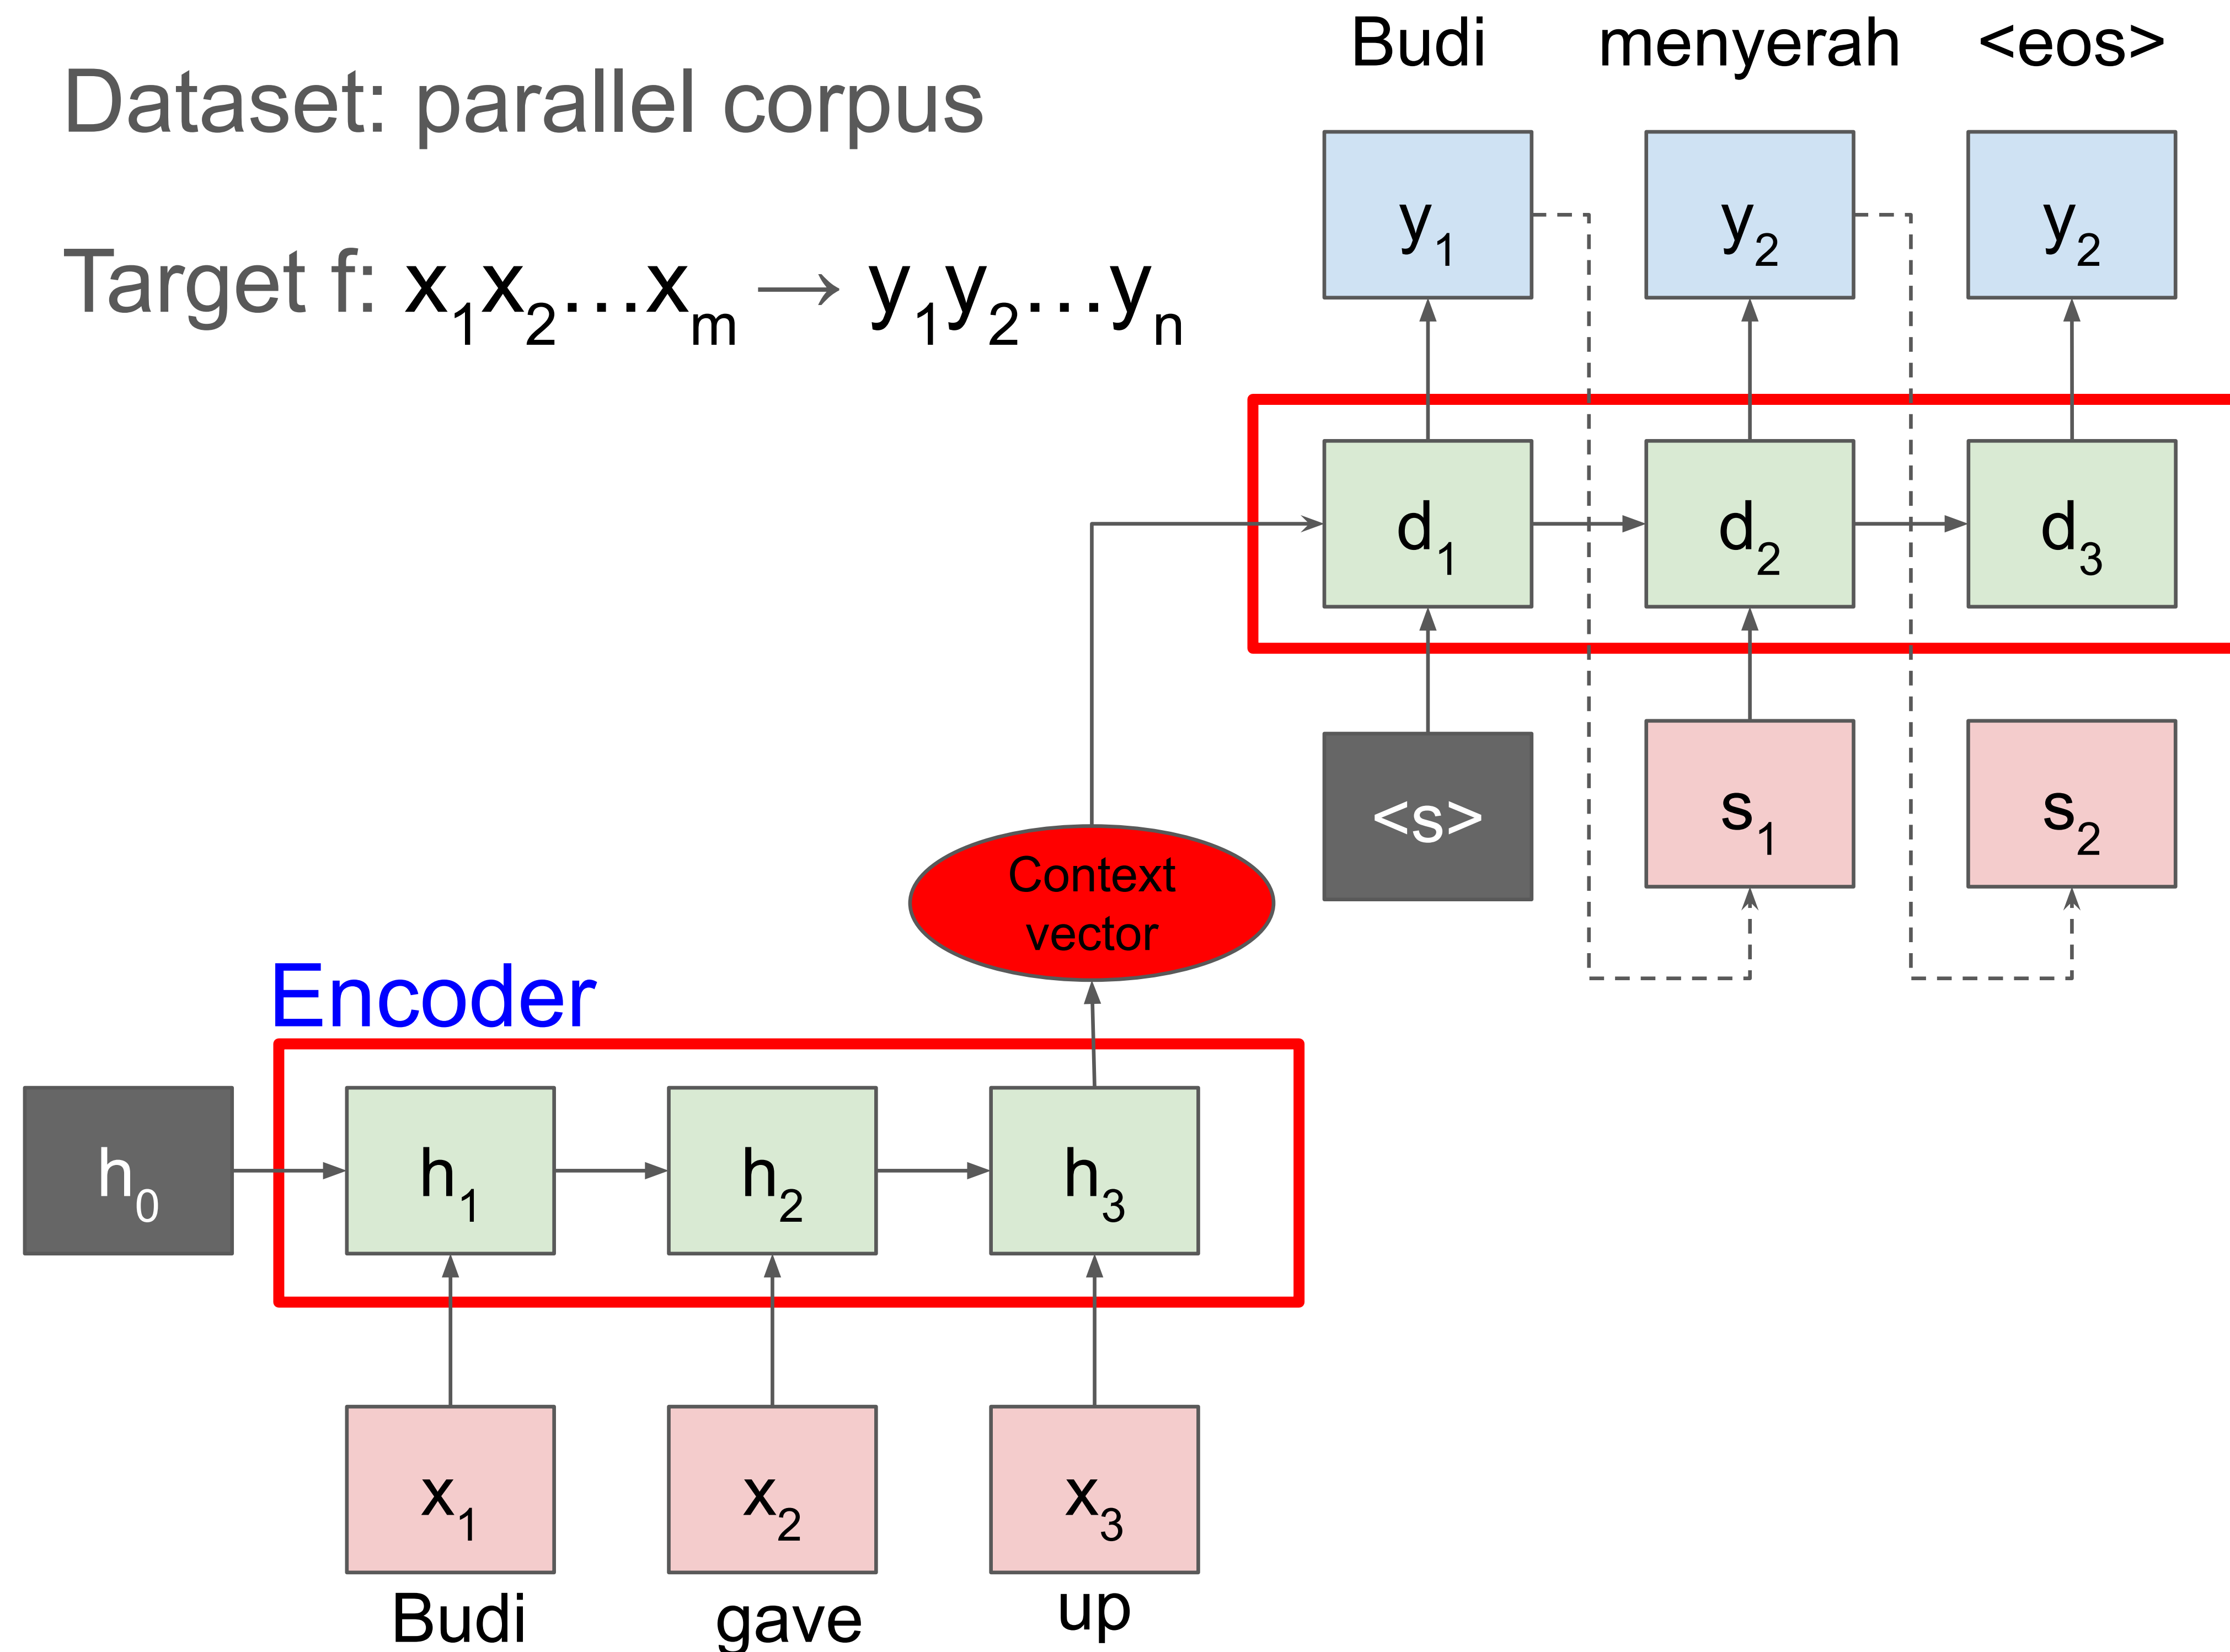
$$p(y_t | \{y_1, \dots, y_{t-1}\}, c) = g(y_{t-1}, s_t, c)$$

$$p(y_i | y_1, \dots, y_{i-1}, \mathbf{x}) = g(y_{i-1}, s_i, c_i)$$
$$s_i = f(s_{i-1}, y_{i-1}, c_i)$$

Encoder - Decoder with Attention

Dataset: parallel corpus

Target f: $x_1 x_2 \dots x_m \rightarrow y_1 y_2 \dots y_n$



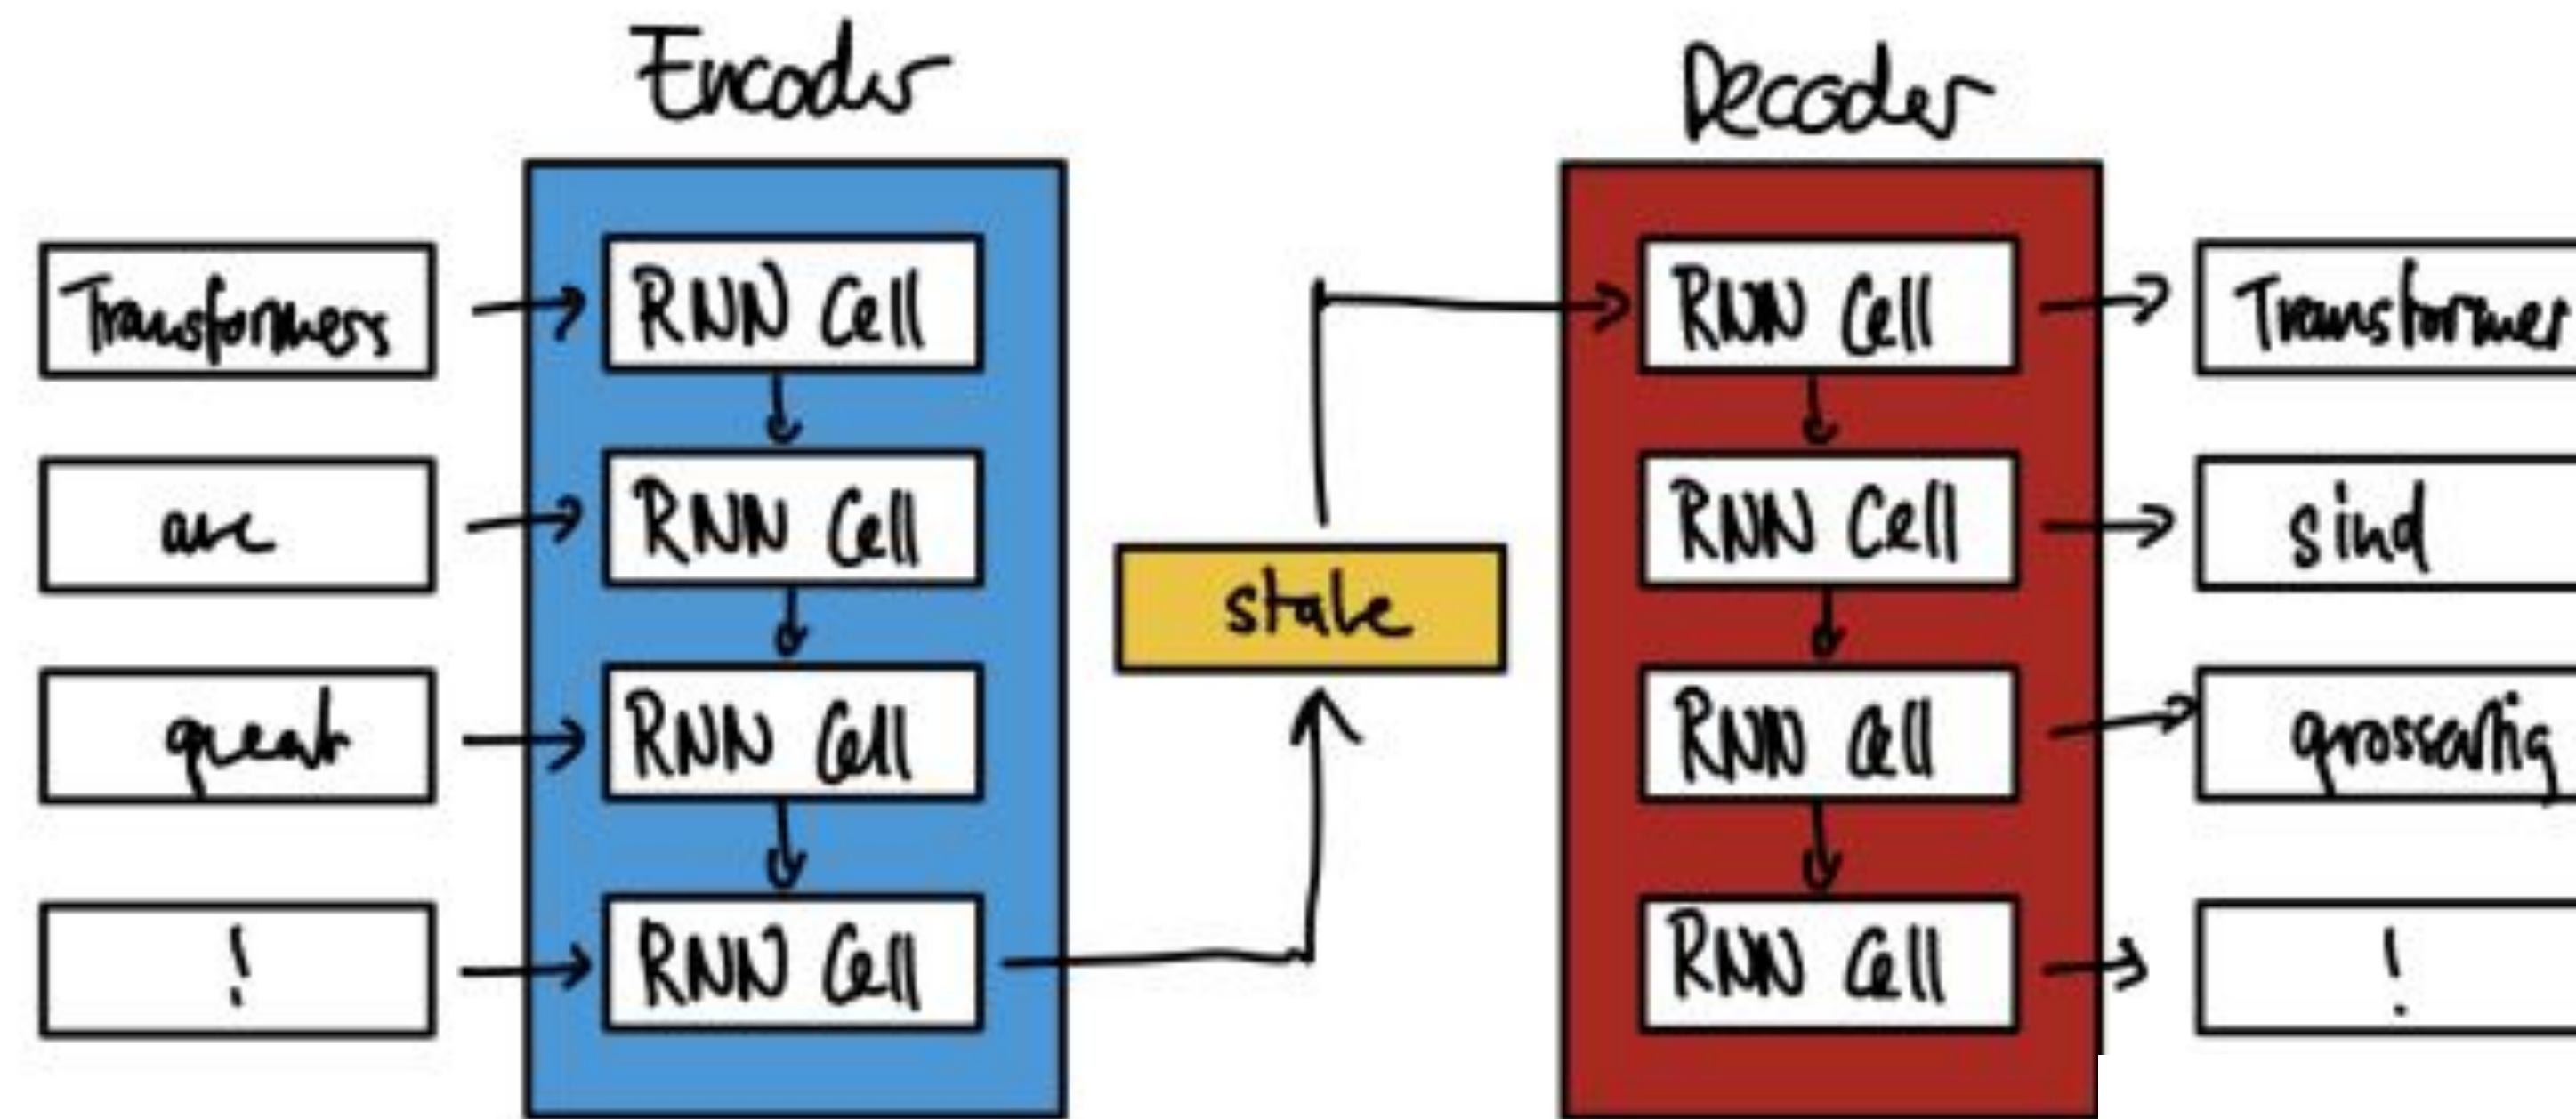
Decoder is trained to predict the **next word** y_t given the **context vector** c and all the previously predicted words $\{y_1, \dots, y_{t-1}\}$.

$$p(y_i | y_1, \dots, y_{i-1}, \mathbf{x}) = g(y_{i-1}, s_i, c_i)$$

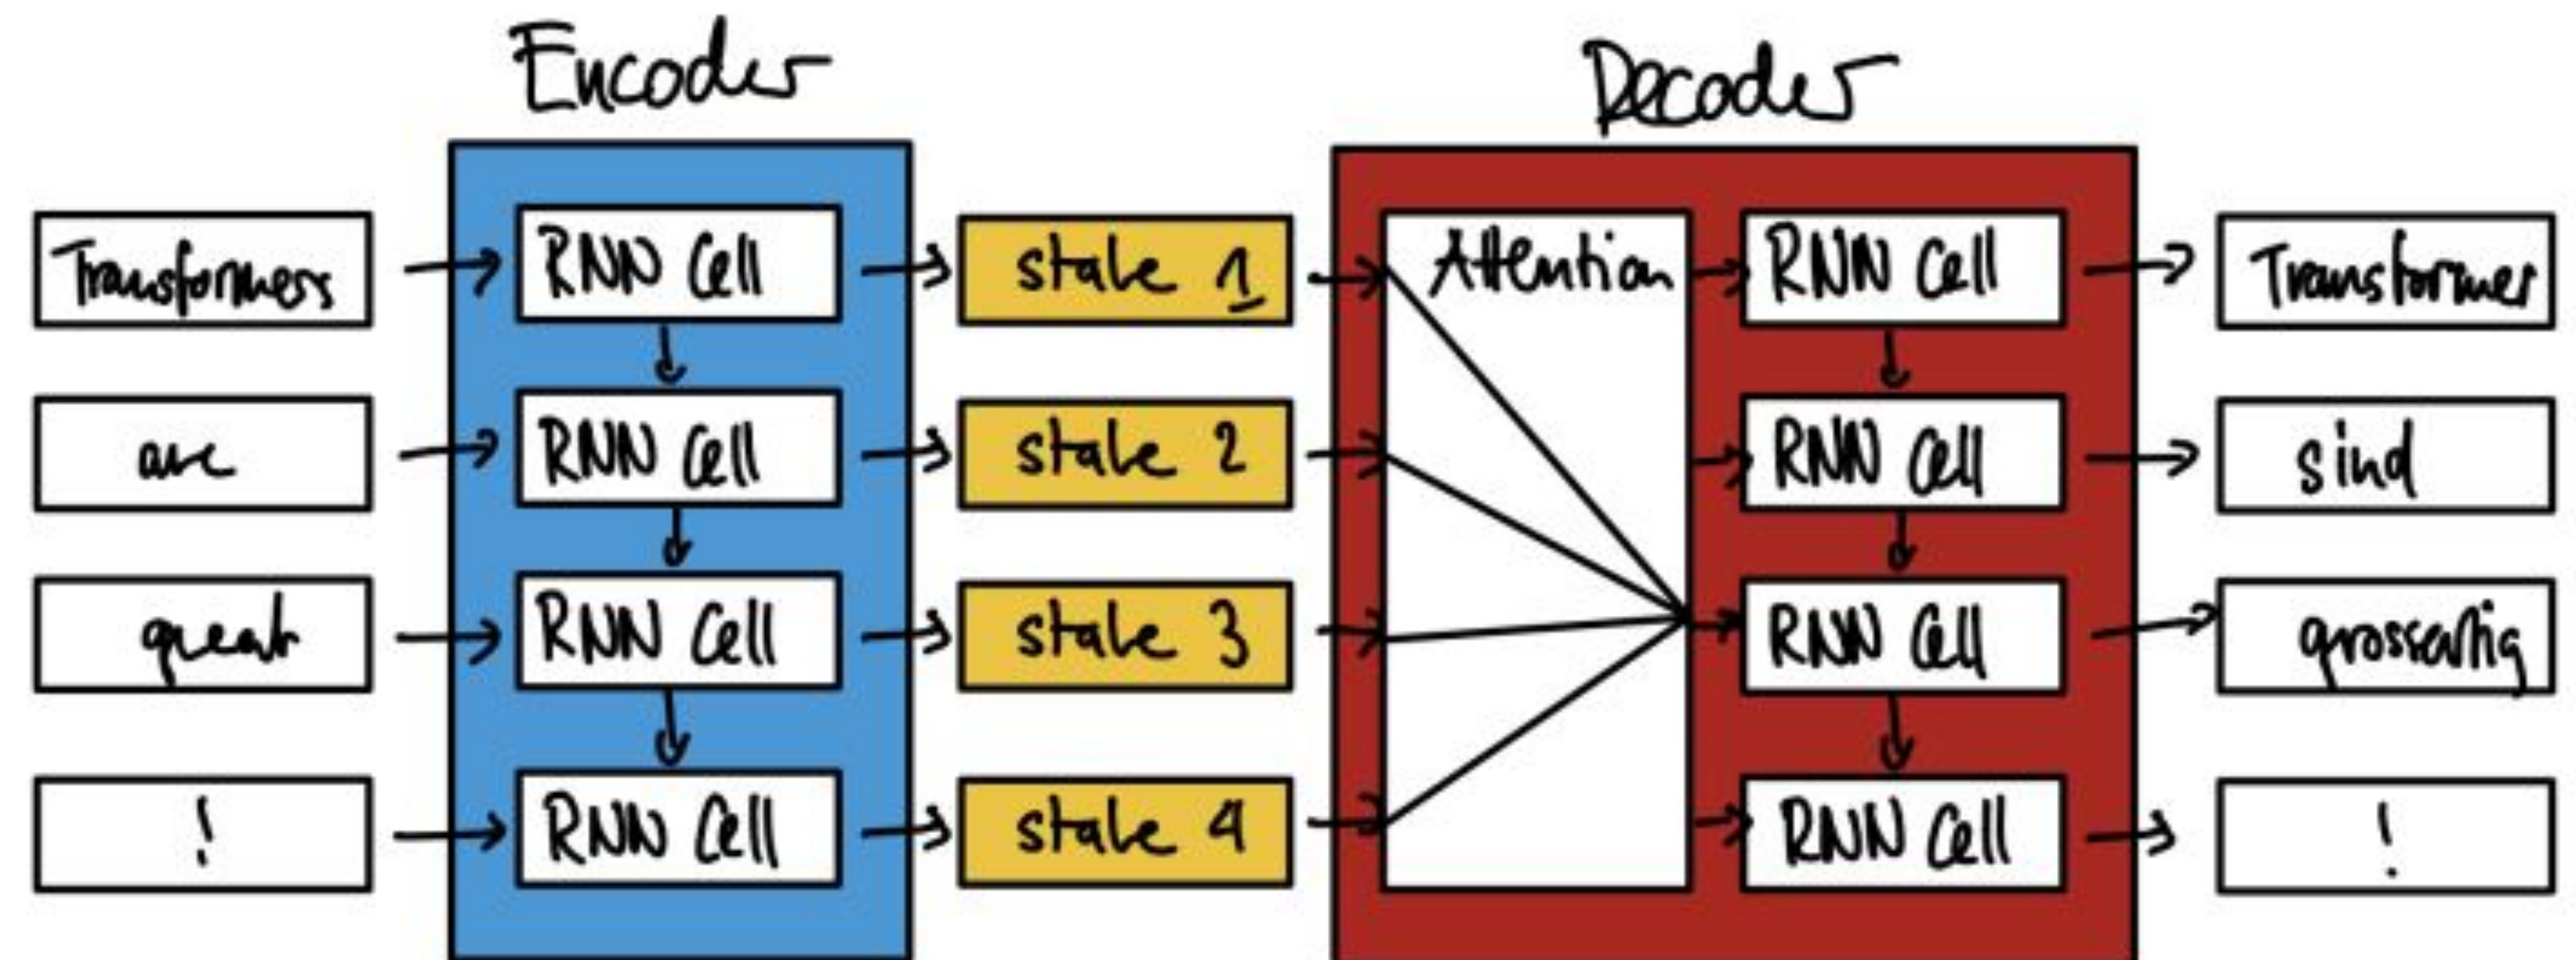
$$p(y_t | \{y_1, \dots, y_{t-1}\}, c) = g(y_{t-1}, s_t, c)$$

$$c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j$$

Vanilla RNN Encoder - Decoder without vs with Attention

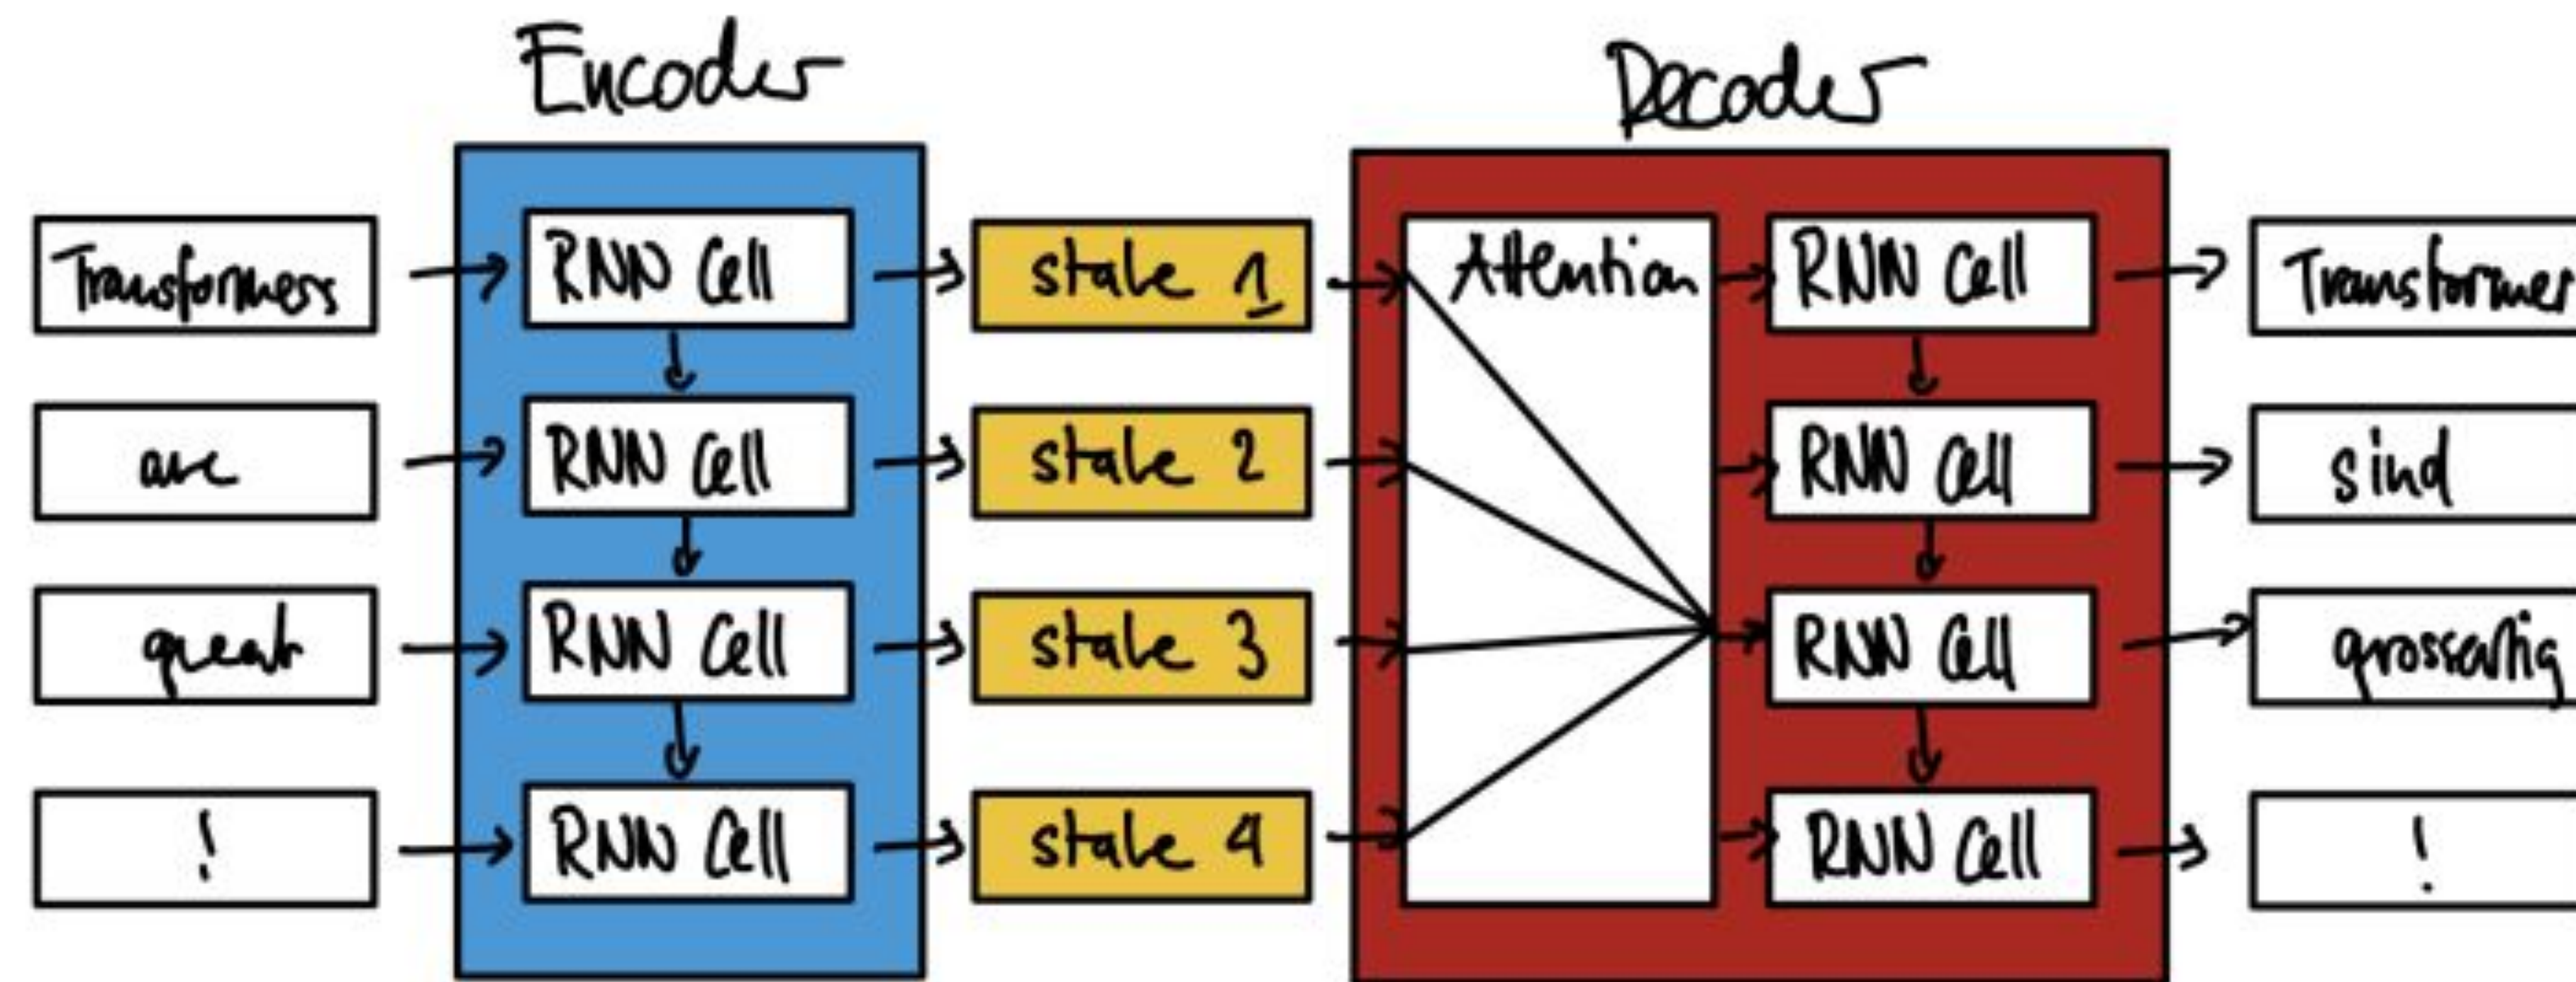


What is the difference between these architectures ?



Tunstall, L., Von Werra, L., & Wolf, T. (2022). *Natural language processing with transformers*. " O'Reilly Media, Inc."

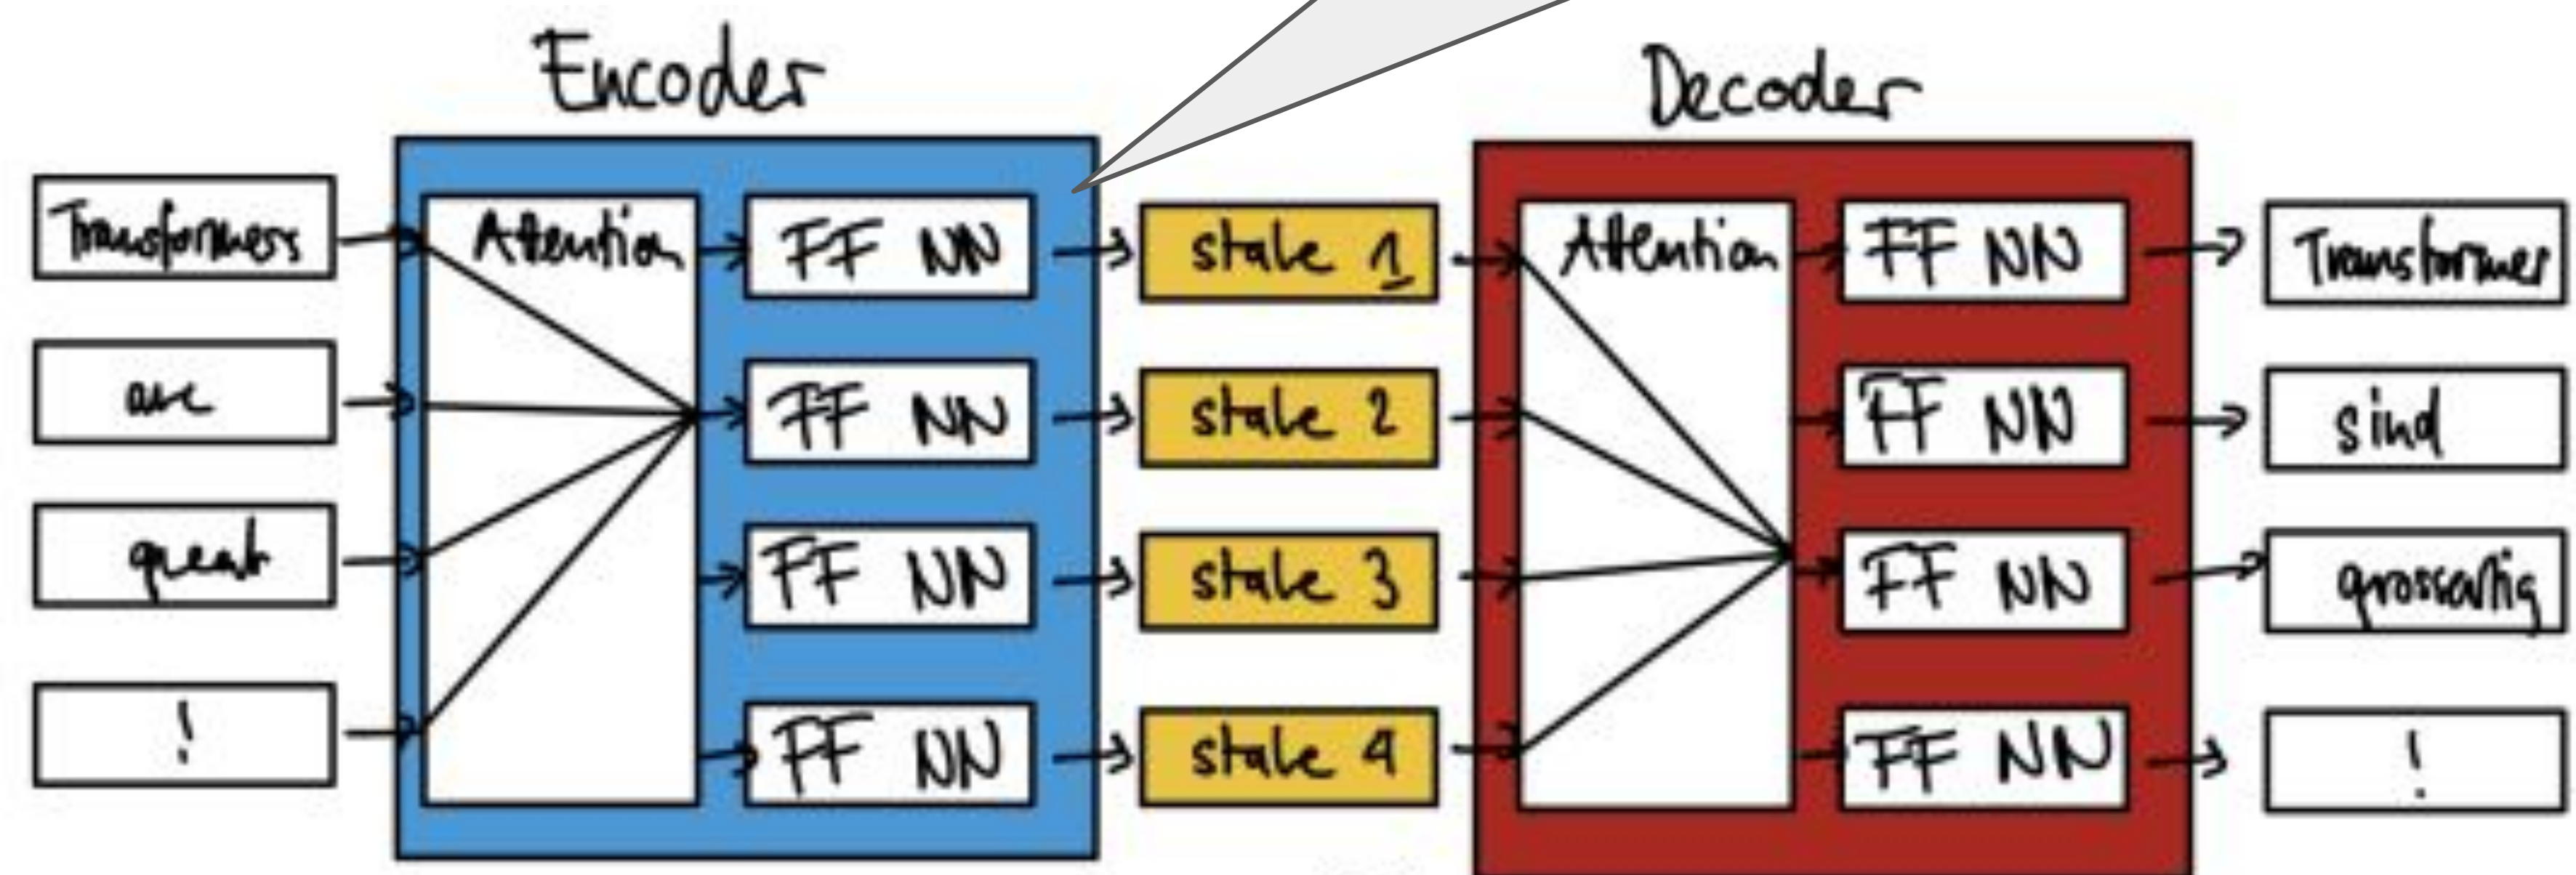
Vanilla RNN Enc-Dec with Attention vs Transformers



The Transformer architecture replaced the recurrent units inside the encoder and decoder entirely with **self-attention** layers and simple **feed-forward** networks.

all the tokens are fed **sequentially** (Vanilla RNN Enc-Dec) vs **in parallel** through the model (transformers)

Tunstall, L., Von Werra, L., & Wolf, T. (2022). *Natural language processing with transformers*. "O'Reilly Media, Inc."



Vanilla RNN Enc-Dec vs Transformers

- Moving from a **sequential** processing to a **fully parallel** processing **unlocked** strong computational **efficiency** gains allowing to train on orders of magnitude **larger corpora** for the same computational cost.
- At the same time, removing the sequential processing bottleneck of information makes the transformer architecture **more efficient** on several task that requires aggregating information over **long time spans**.
- The scaling laws of deep learning models: *larger models trained on more data in many cases yield better results*.
 - scaling models comes at the price of requiring large amounts of training data
- Transformer revolution started: **transfer learning**

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Łukasz Kaiser*
Google Brain
lukaszkaiser@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com

Vaswani, A. (2017). Attention is all you need. *Advances in neural information processing systems*, 30, 1.

... We propose a new simple network architecture, the Transformer, based solely on **attention** mechanisms, dispensing with **recurrence** and **convolutions** entirely. Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring significantly less time to train ...

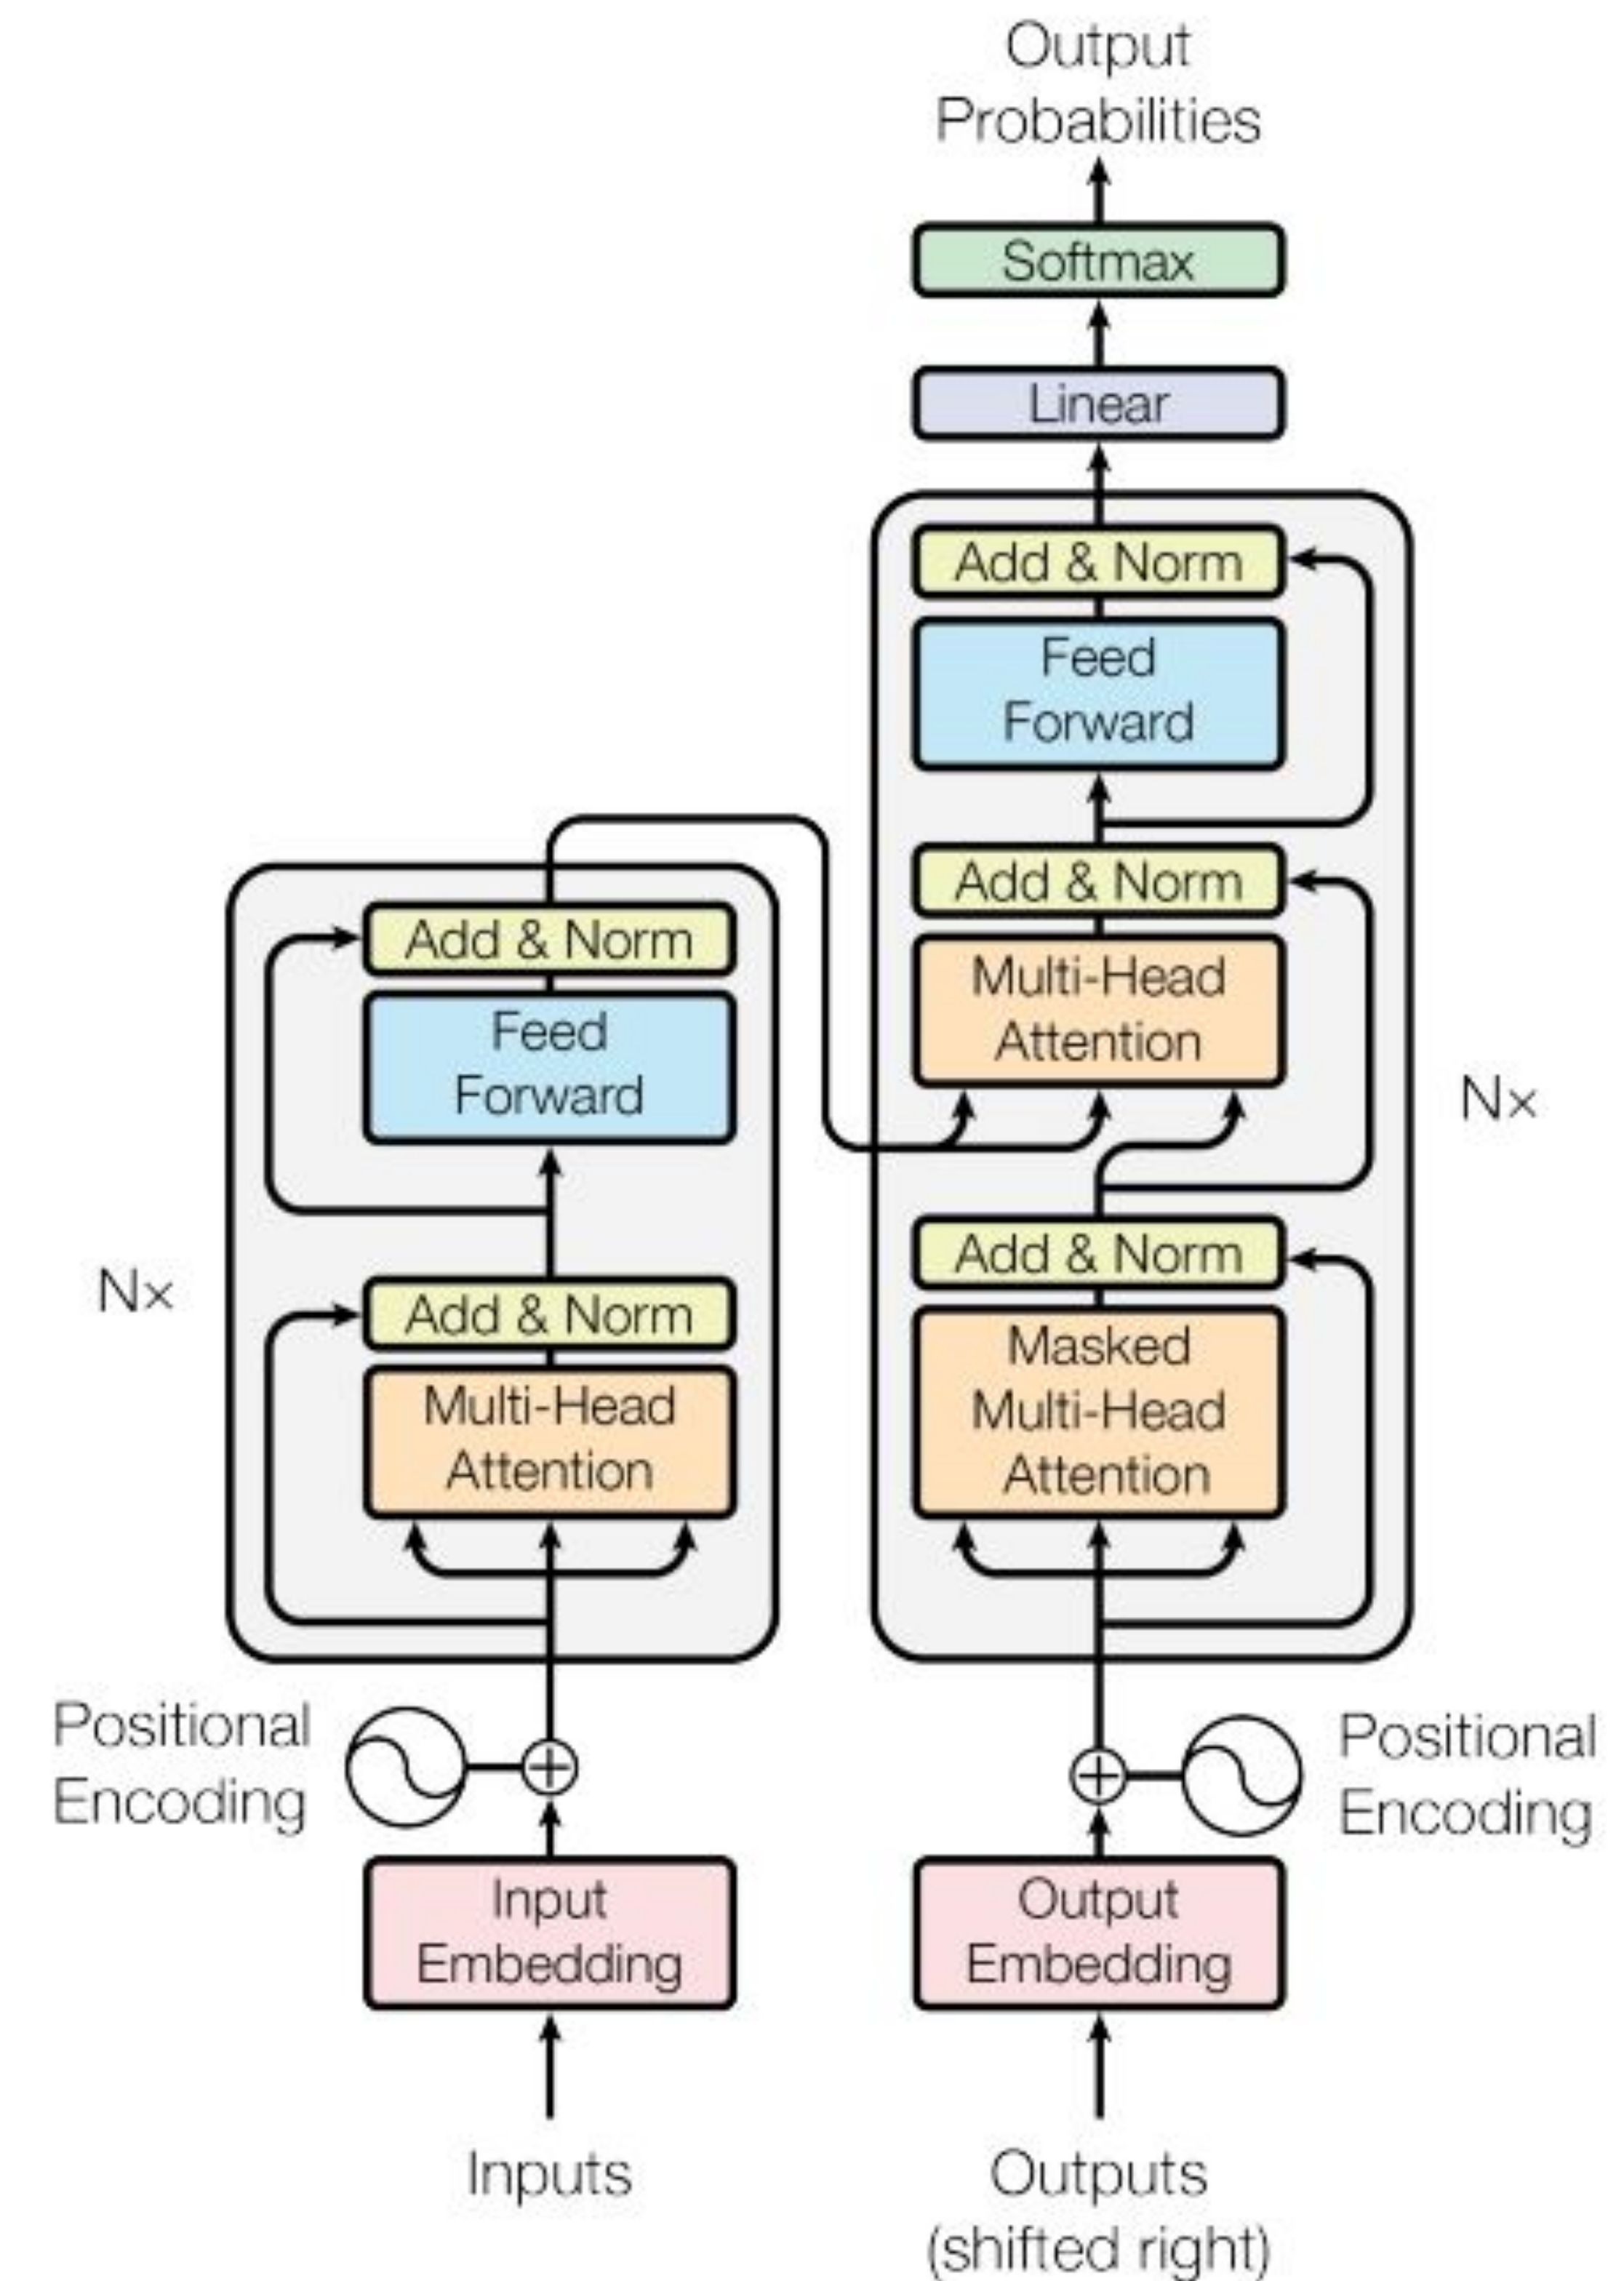


Figure 1: The Transformer - model architecture.

Scientific Breakthrough

Attention for Encoder-Decoder (Seq2Seq)

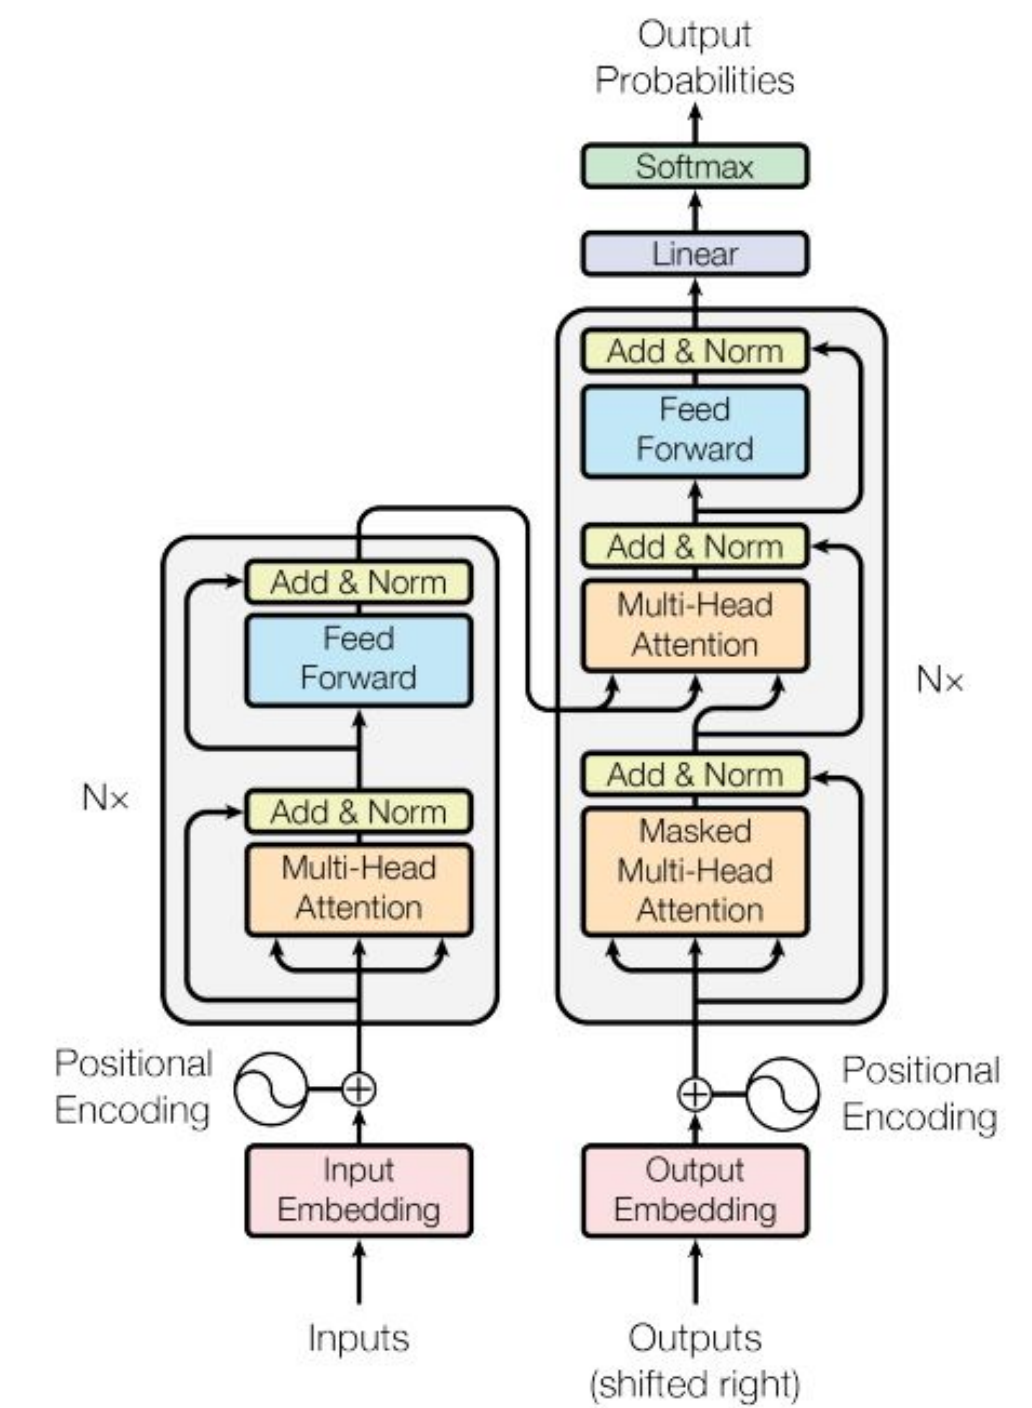

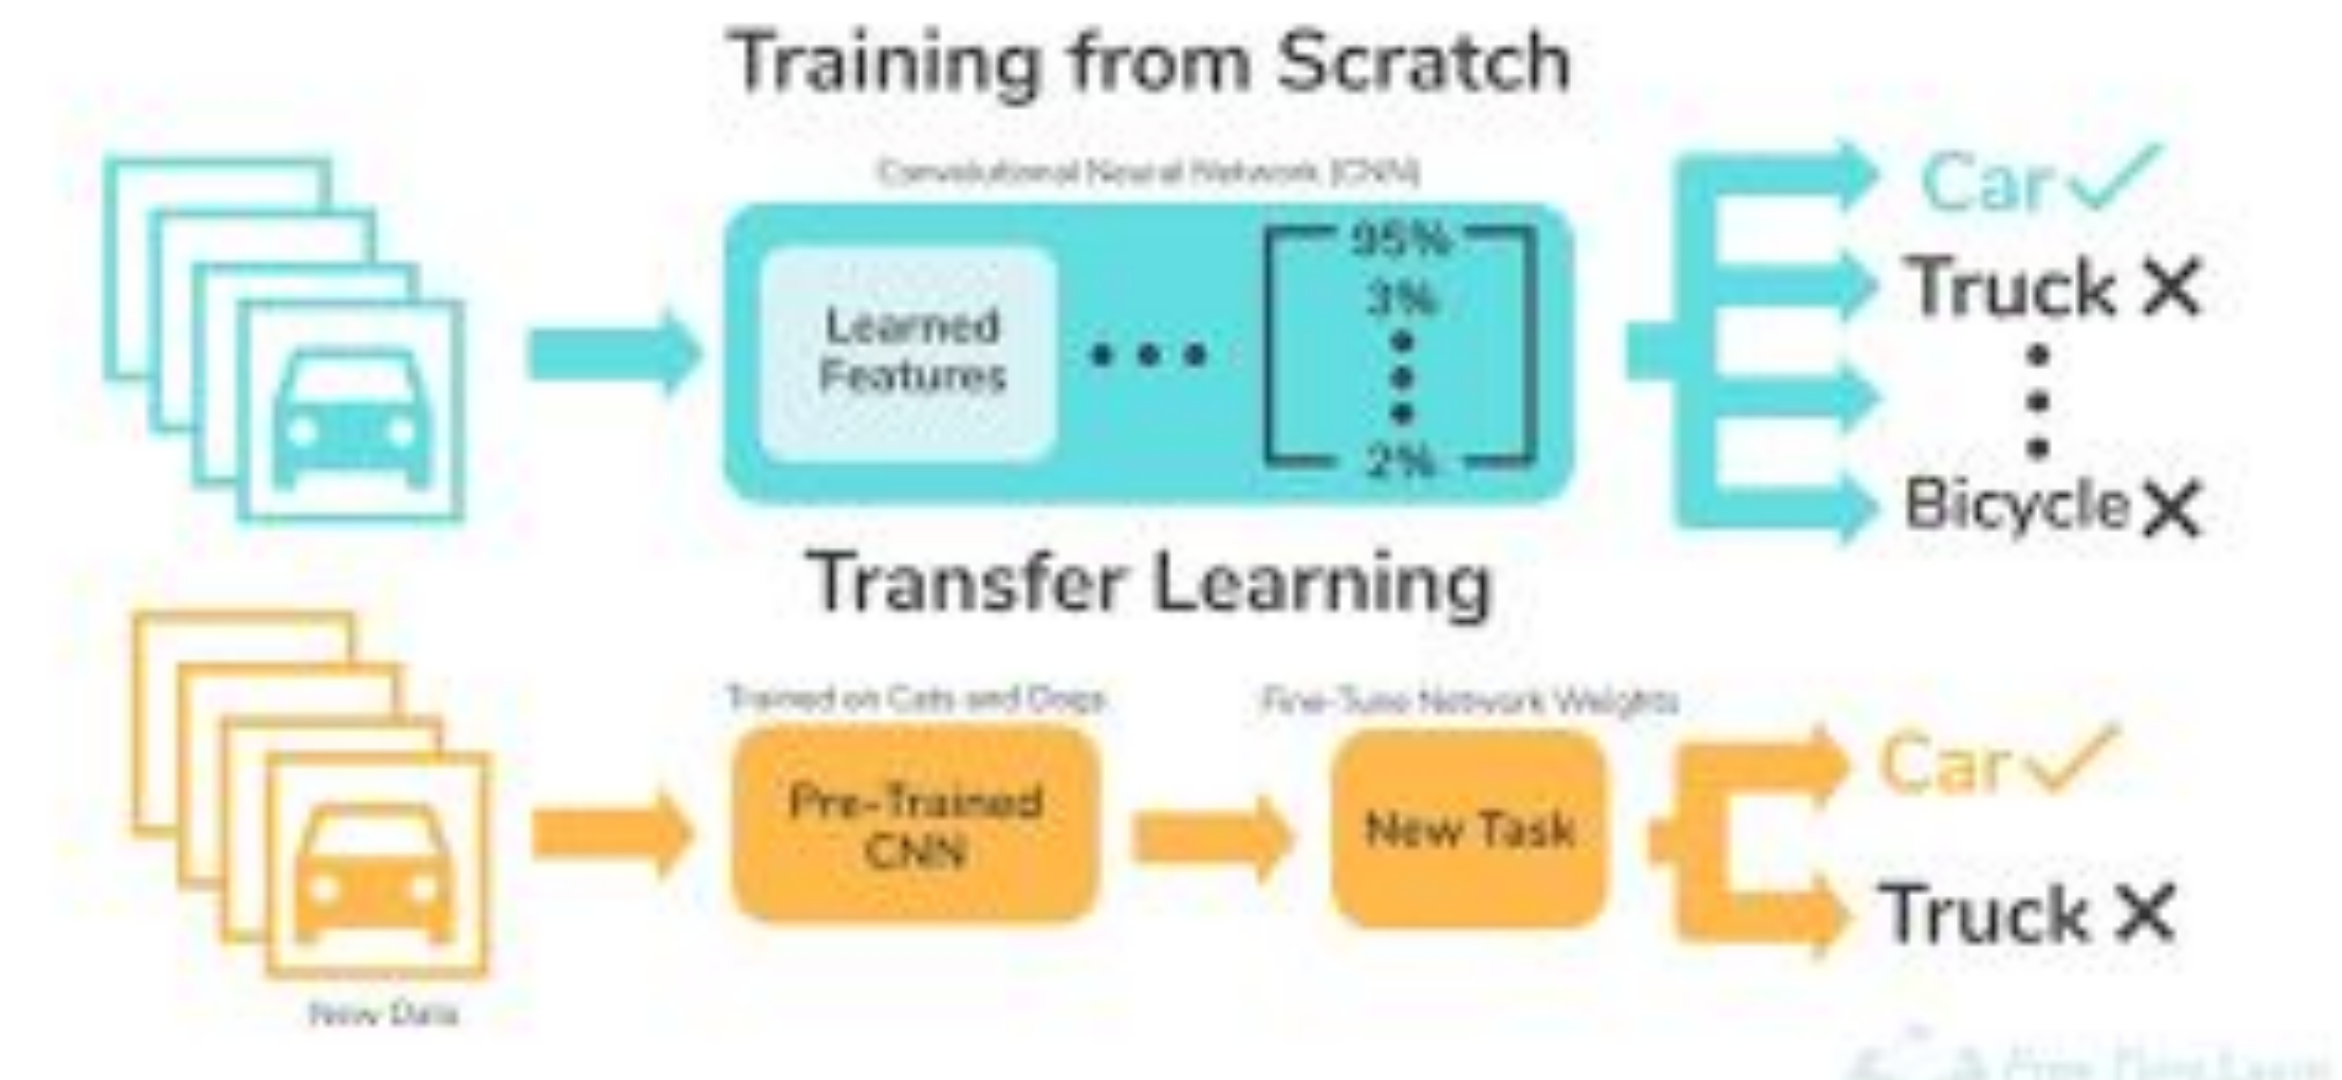


Figure 1: The Transformer - model architecture.

Transfer Learning



Questions ?