

source: NVIDIA GPU Teaching Kit

# IF3230

## Sistem Paralel dan Terdistribusi

### CUDA

Achmad Imam Kistijantoro ([imam@staff.stei.itb.ac.id](mailto:imam@staff.stei.itb.ac.id))

Robithoh Annur ([robithoh@staff.stei.itb.ac.id](mailto:robithoh@staff.stei.itb.ac.id))

Andreas Bara Timur ([bara@staff.stei.itb.ac.id](mailto:bara@staff.stei.itb.ac.id))

Februari 2023  
3/1/2023

# Objective

- To learn convolution, an important method
  - Widely used in audio, image and video processing
  - Foundational to stencil computation used in many science and engineering applications
  - Basic 1D and 2D convolution kernels


# Convolution as a Filter

- Often performed as a filter that transforms signal or pixel values into more desirable values.
  - Some filters smooth out the signal values so that one can see the big-picture trend
  - Others like Gaussian filters can be used to sharpen boundaries and edges of objects in images..

# Convolution – a computational definition

- An array operation where each output data element is a weighted sum of a collection of neighboring input elements
- The weights used in the weighted sum calculation are defined by an input mask array, commonly referred to as the *convolution kernel*
  - We will refer to these mask arrays as convolution masks to avoid confusion.
  - The value pattern of the mask array elements defines the type of filtering done


# 1D Convolution Example




- Commonly used for audio processing
  - Mask size is usually an odd number of elements for symmetry (5 in this example)
- The figure shows calculation of  $P[2]$

$$P[2] = N[0]*M[0] + N[1]*M[1] + N[2]*M[2] + N[3]*M[3] + N[4]*M[4]$$

# Calculation of P[3]



# Convolution Boundary Condition



- Calculation of output elements near the boundaries (beginning and end) of the array need to deal with “ghost” elements
  - Different policies (0, replicates of boundary values, etc.)

# A 1D Convolution Kernel with Boundary Condition Handling

- This kernel forces all elements outside the valid input range to 0

```
__global__ void convolution_1D_basic_kernel(float *N, float *M,  
float *P, int Mask_Width, int Width)  
{  
    int i = blockIdx.x*blockDim.x + threadIdx.x;  
  
    float Pvalue = 0;  
    int N_start_point = i - (Mask_Width/2);  
  
    for (int j = 0; j < Mask_Width; j++) {  
        if (N_start_point + j >= 0 && N_start_point + j < Width) {  
            Pvalue += N[N_start_point + j]*M[j];  
        }  
    }  
  
    P[i] = Pvalue;  
}
```



# A 1D Convolution Kernel with Boundary Condition Handling

- This kernel forces all elements outside the valid input range to 0

```
__global__ void convolution_1D_basic_kernel(float *N, float *M,
                                           float *P, int Mask_Width, int Width)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;

    float Pvalue = 0;
    int N_start_point = i - (Mask_Width/2);

    if (i < Width) {

        for (int j = 0; j < Mask_Width; j++) {
            if (N_start_point + j >= 0 && N_start_point + j < Width) {
                Pvalue += N[N_start_point + j]*M[j];
            }
        }

        P[i] = Pvalue;
    }
}
```

# 2D Convolution

**N**

1	2	3	4	5	6	7
2	3	4	5	6	7	8
3	4	5	6	7	8	9
4	5	6	7	8	5	6
5	6	7	8	5	6	7
6	7	8	9	0	1	2
7	8	9	0	1	2	3

**P**

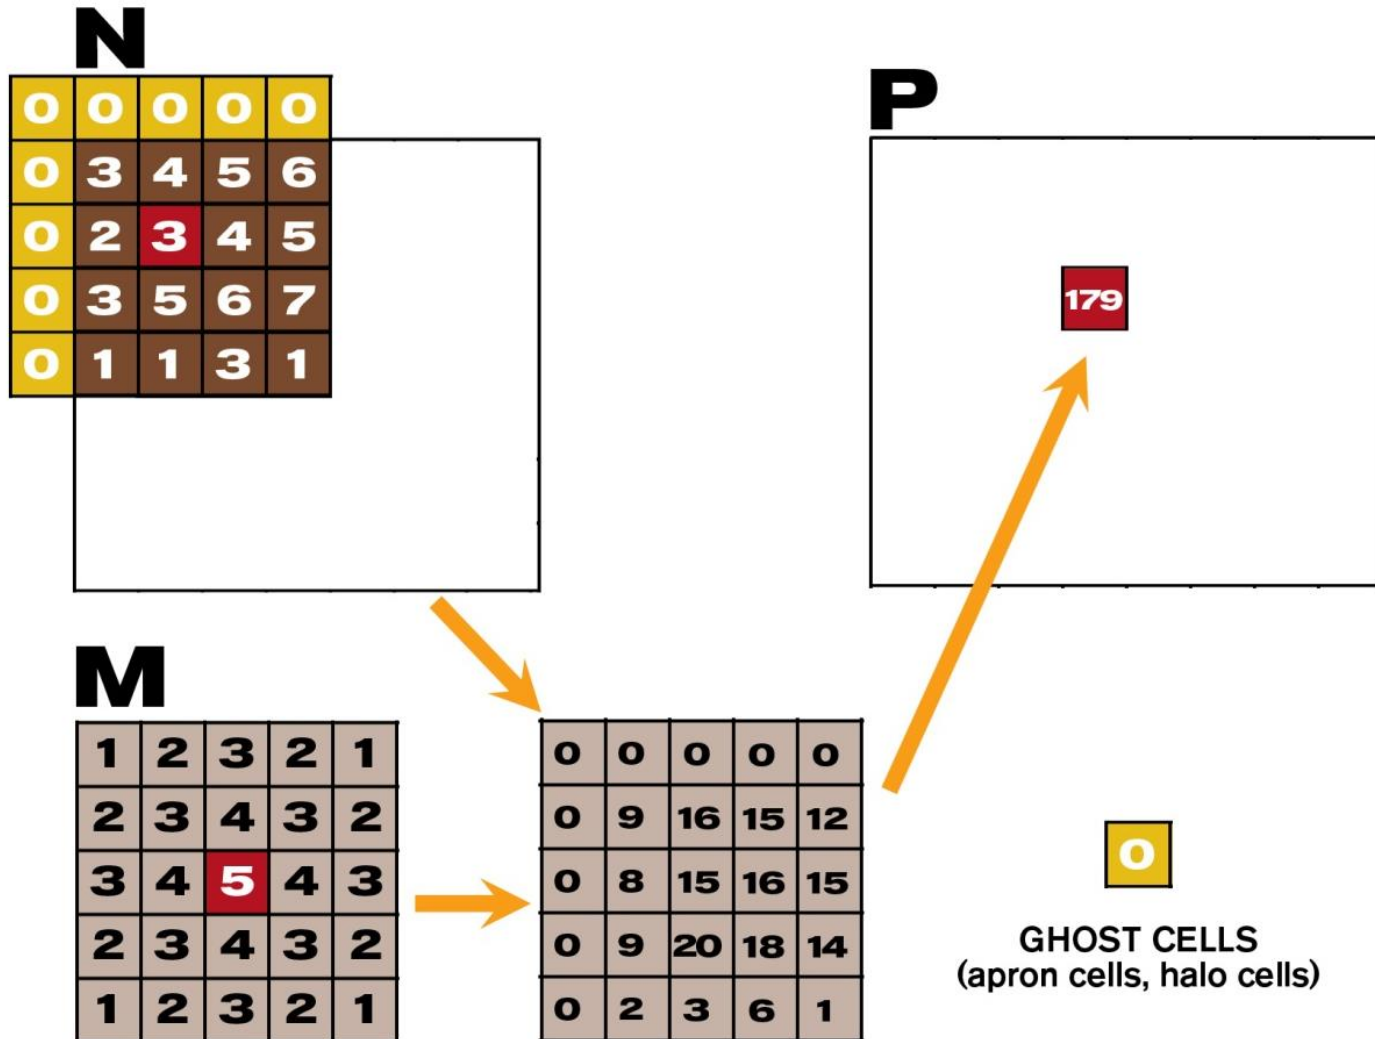
1	2	3	4	5			
2	3	4	5	6			
3	4	321	6	7			
4	5	6	7	8			
5	6	7	8	5			

**M**

1	2	3	2	1
2	3	4	3	2
3	4	5	4	3
2	3	4	3	2
1	2	3	2	1

1	4	9	8	5
4	9	16	15	12
4	16	25	24	21
8	15	24	21	16
5	12	21	16	5

# 2D Convolution – Ghost Cells



\_\_global\_\_

```
void convolution_2D_basic_kernel(unsigned char * in, unsigned char * mask, unsigned char * out,  
                                int maskwidth, int w, int h) {
```

```
    int Col = blockIdx.x * blockDim.x + threadIdx.x;  
    int Row = blockIdx.y * blockDim.y + threadIdx.y;
```

```
    if (Col < w && Row < h) {  
        int pixVal = 0;
```

```
        N_start_col = Col - (maskwidth/2);  
        N_start_row = Row - (maskwidth/2);
```

```
        // Get the of the surrounding box
```

```
        for(int j = 0; j < maskwidth; ++j) {  
            for(int k = 0; k < maskwidth; ++k) {
```

```
                int curRow = N_start_row + j;
```


```
                int curCol = N_start_col + k;
```

```
                // Verify we have a valid image pixel
```

```
                if(curRow > -1 && curRow < h && curCol > -1 && curCol < w)  
                    pixVal += in[curRow * w + curCol] * mask[j*maskwidth+k];  
            }  
        }
```

```
        // Write our new pixel value out
```

```
        out[Row * w + Col] = (unsigned char)(pixVal);
```



\_\_global\_\_

```
void convolution_2D_basic_kernel(unsigned char * in, unsigned char * mask, unsigned char * out,  
    int maskwidth, int w, int h) {  
    int Col = blockIdx.x * blockDim.x + threadIdx.x;  
    int Row = blockIdx.y * blockDim.y + threadIdx.y;
```

```
    if (Col < w && Row < h) {
```

```
        int pixVal = 0;
```

```
        N_start_col = Col - (maskwidth/2);  
        N_start_row = Row - (maskwidth/2);
```

N\_start\_row

```
        // Get the of the surrounding box
```

```
        for(int j = 0; j < maskwidth; ++j) {  
            for(int k = 0; k < maskwidth; ++k) {
```

```
                int curRow = N_Start_row + j;
```

```
                int curCol = N_start_col + k;
```

```
                // Verify we have a valid image pixel
```

```
                if(curRow > -1 && curRow < h && curCol > -1 && curCol < w)  
                    pixVal += in[curRow * w + curCol] * mask[j*maskwidth+k];  
            }  
        }
```

```
        // Write our new pixel value out
```

```
        out[Row * w + Col] = (unsigned char)(pixVal);
```

N\_start\_col

**N**

1	2	3	4	5	6	7
2	3	4	5	6	7	8
3	4	5	6	7	8	9
4	5	6	7	8	5	6
5	6	7	8	5	6	7
6	7	8	9	0	1	2
7	8	9	0	1	2	3

**M**

1	2	3	2	1
2	3	4	3	2
3	4	5	4	3
2	3	4	3	2
1	2	3	2	1

1	4	9	8	5
4	9	16	15	12
4	16	25	24	21
8	15	24	21	16
5	12	21	16	5

\_\_global\_\_

```
void convolution_2D_basic_kernel(unsigned char * in, unsigned char * mask, unsigned char * out,  
                                int maskwidth, int w, int h) {  
    int Col = blockIdx.x * blockDim.x + threadIdx.x;  
    int Row = blockIdx.y * blockDim.y + threadIdx.y;
```

```
    if (Col < w && Row < h) {  
        int pixVal = 0;
```

```
        N_start_col = Col - (maskwidth/2);  
        N_start_row = Row - (maskwidth/2);
```

```
        // Get the of the surrounding box
```

```
        for(int j = 0; j < maskwidth; ++j) {  
            for(int k = 0; k < maskwidth; ++k) {  
  
                int curRow = N_Start_row + j;  
                int curCol = N_start_col + k;  
                // Verify we have a valid image pixel  
                if(curRow > -1 && curRow < h && curCol > -1 && curCol < w) {  
                    pixVal += in[curRow * w + curCol] * mask[j*maskwidth+k];  
                }  
            }  
        }  
    }
```

```
        // Write our new pixel value out
```

```
        out[Row * w + Col] = (unsigned char)(pixVal);
```




```
    }
```

# Objective

- To learn about tiled convolution algorithms
  - Some intricate aspects of tiling algorithms
  - Output tiles versus input tiles

# Tiling Opportunity Convolution


- Calculation of adjacent output elements involve shared input elements
  - E.g.,  $N[2]$  is used in calculation of  $P[0]$ ,  $P[1]$ ,  $P[2]$ .  $P[3]$  and  $P[5]$  assuming a 1D convolution Mask\_Width of width 5
- We can load all the input elements required by all threads in a block into the shared memory to reduce global memory accesses






# Input Data Needs

- Assume that we want to have each block to calculate  $T$  output elements
  - $T + \text{Mask\_Width} - 1$  input elements are needed to calculate  $T$  output elements
  - $T + \text{Mask\_Width} - 1$  is usually not a multiple of  $T$ , except for small  $T$  values
  - $T$  is usually significantly larger than  $\text{Mask\_Width}$



# Definition – output tile




Each thread block calculates an output tile

Each output tile width is O\_TILE\_WIDTH

For each thread,

O\_TILE\_WIDTH is 4 in this example

# Definition - Input Tiles

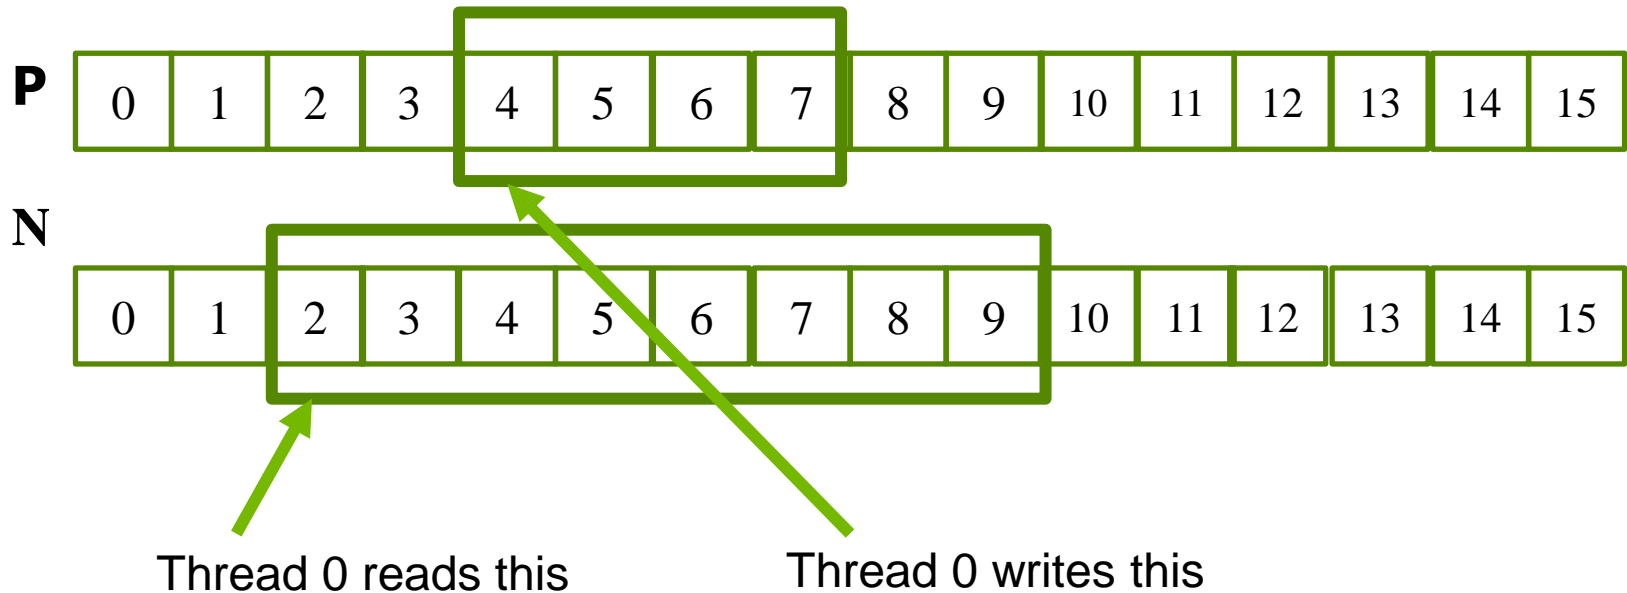


**Each input tile has all values needed to calculate the corresponding output tile.**

# Two Design Options

- Design 1: The size of each thread block matches the size of an output tile
  - All threads participate in calculating output elements
  - `blockDim.x` would be 4 in our example
  - Some threads need to load more than one input element into the shared memory
- Design 2: The size of each thread block matches the size of an input tile
  - Some threads will not participate in calculating output elements
  - `blockDim.x` would be 8 in our example
  - Each thread loads one input element into the shared memory
- We will present Design 2 and leave Design 1 as an exercise.

# Thread to Input and Output Data Mapping



For each thread,  
 $\text{Index}_i = \text{index}_o - n$

were  $n$  is  $\text{Mask\_Width} / 2$   
 $n$  is 2 in this example

# All Threads Participate in Loading Input Tiles

```
float output = 0.0f;

if((index_i >= 0) && (index_i < Width)) {
    Ns[tx] = N[index_i];
}
else{
    Ns[tx] = 0.0f;
}
```

# Some threads do not participate in calculating output

```
if (threadIdx.x < O_TILE_WIDTH) {  
    output = 0.0f;  
    for(j = 0; j < Mask_Width; j++) {  
        output += M[j] * Ns[j+threadIdx.x];  
    }  
    P[index_o] = output;  
}
```

- $\text{index\_o} = \text{blockIdx.x} * \text{O\_TILE\_WIDTH} + \text{threadIdx.x}$
- Only Threads 0 through  $\text{O\_TILE\_WIDTH}-1$  participate in calculation of output.

# Setting Block Size

```
#define O_TILE_WIDTH 1020
#define BLOCK_WIDTH (O_TILE_WIDTH + 4)

dim3 dimBlock(BLOCK_WIDTH, 1, 1);

dim3 dimGrid((Width-1)/O_TILE_WIDTH+1, 1, 1)
```

The Mask\_Width is 5 in this example

In general, block width should be

output tile width + (mask width-1)



# Shared Memory Data Reuse

**N\_ds**

Mask\_Width is 5



Element 2 is used by thread 4 (1X)

Element 3 is used by threads 4, 5 (2X)

Element 4 is used by threads 4, 5, 6 (3X)

Element 5 is used by threads 4, 5, 6, 7 (4X)

Element 6 is used by threads 4, 5, 6, 7 (4X)


Element 7 is used by threads 5, 6, 7 (3X)

Element 8 is used by threads 6, 7 (2X)

Element 9 is used by thread 7 (1X)

# Ghost Cells

# N



# Objective

- To learn to write a 2D convolution kernel
  - 2D Image data types and API functions
  - Using constant caching
  - Input tiles vs. output tiles in 2D
  - Thread to data index mapping
  - Handling boundary conditions

# 2D Image Matrix with Automated Padding


- It is sometimes desirable to pad each row of a 2D matrix to multiples of DRAM bursts
  - So each row starts at the DRAM burst boundary
  - Effectively adding columns
  - This is usually done automatically by matrix allocation function
  - Pitch can be different for different hardware
- Example: a 3X3 matrix padded into a 3X4 matrix

Height is 3

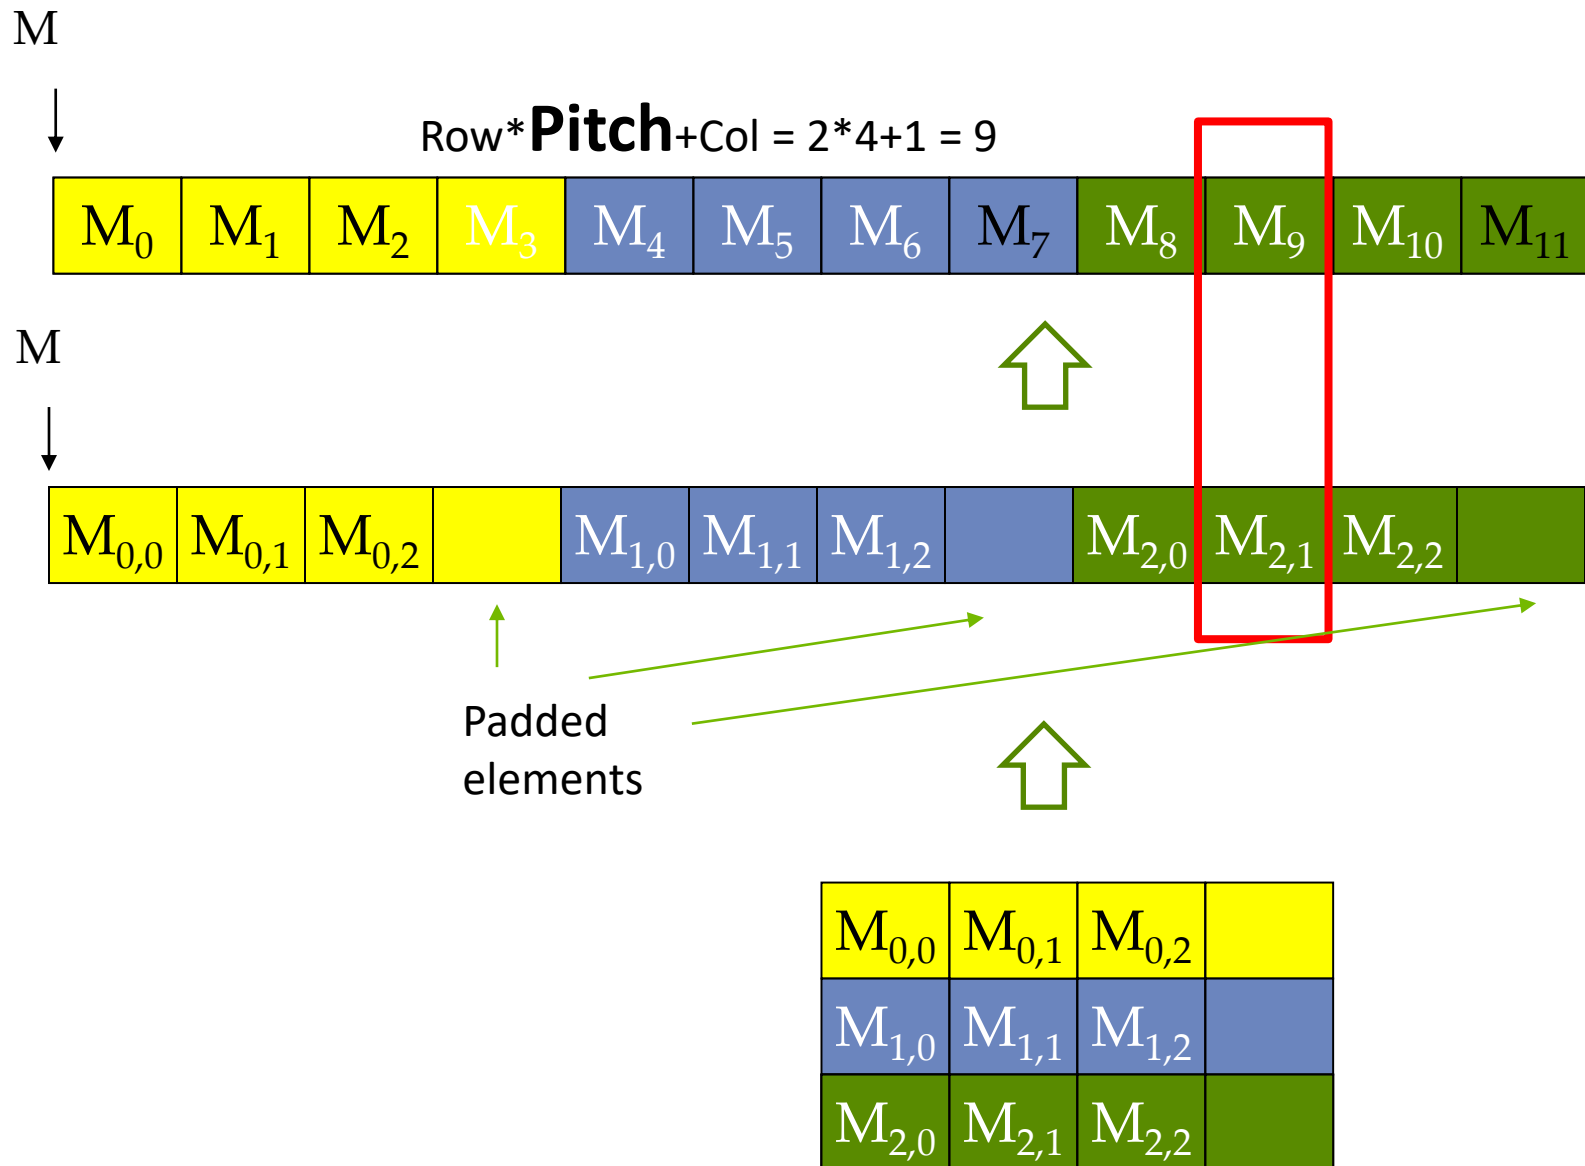
Width is 3

Channels is 1 (See MP Description)

Pitch is 4



# Row-Major Layout with Pitch



# Image Matrix Type in this Course

```
// Image Matrix Structure declaration
//
typedef struct {
    int width;
    int height;
    int pitch;
    int channels;
    float* data;
} * wbImage_t;
```

This type will only be used in the host code of the MP.

# wbImage\_t API Function for Your Lab

```
wbImage_t  wbImage_new(int height, int  
width, int channels)  
wbImage_t  wbImport(char * File);
```

```
void wbImage_delete(wbImage_t img)
```

```
int wbImage_getWidth(wbImage_t img)  
int wbImage_getHeight(wbImage_t img)  
int wbImage_getChannels(wbImage_t img)  
int wbImage_getPitch(wbImage_t img)
```

```
float *wbImage_getData(wbImage_t img)
```

For simplicity, the pitch of all matrices are set to be width \* channels (no padding) for our labs.

The use of all API functions has been done in the provided host code.

# Setting Block Size

```
#define O_TILE_WIDTH 12
#define BLOCK_WIDTH (O_TILE_WIDTH + 4)

dim3 dimBlock(BLOCK_WIDTH, BLOCK_WIDTH);
dim3 dimGrid((wbImage_getWidth(N)-1)/O_TILE_WIDTH+1,
             (wbImage_getHeight(N)-1)/O_TILE_WIDTH+1, 1)
```

In general, BLOCK\_WIDTH should be  
 $O\_TILE\_WIDTH + (MASK\_WIDTH - 1)$



# Using constant memory and caching for Mask


- Mask is used by all threads but not modified in the convolution kernel
  - All threads in a warp access the same locations at each point in time
- CUDA devices provide constant memory whose contents are aggressively cached
  - Cached values are broadcast to all threads in a warp
  - Effectively magnifies memory bandwidth without consuming shared memory
- Use of `const __restrict__` qualifiers for the mask parameter informs the compiler that it is eligible for constant caching, for example:

```
__global__ void convolution_2D_kernel(float *P,  
    float *N, height, width, channels,  
    const float __restrict__ *M) {
```

# Shifting from output coordinates to input coordinate

```
int tx = threadIdx.x;  
int ty = threadIdx.y;  
int row_o = blockIdx.y*_TILE_WIDTH + ty;  
int col_o = blockIdx.x*_TILE_WIDTH + tx;
```

```
int row_i = row_o - 2;  
int col_i = col_o - 2;
```



# Taking Care of Boundaries (1 channel example)

```
if((row_i >= 0) && (row_i < height) &&
    (col_i >= 0) && (col_i < width)) {
    Ns[ty][tx] = data[row_i * width + col_i];
} else{
    Ns[ty][tx] = 0.0f;
}
```



Use of width here is OK since pitch is set to width for this MP.

## Some threads do not participate in calculating output. (1 channel example)

```
float output = 0.0f;
if (ty < O_TILE_WIDTH && tx < O_TILE_WIDTH) {
    for (i = 0; i < MASK_WIDTH; i++) {
        for (j = 0; j < MASK_WIDTH; j++) {
            output += M[i][j] * Ns[i+ty][j+tx];
        }
    }
}
```

## Some threads do not write output (1 channel example)

```
if(row_o < height && col_o < width)
    data[row_o*width + col_o] = output;
```

You need to write the kernel for a 3-channel (RGB) image.  
See more details in the Lab MP Description.

# Objective

- To learn to analyze the cost and benefit of tiled parallel convolution algorithms
  - More complex reuse pattern than matrix multiplication
  - Less uniform access patterns

# An 8-element Convolution Tile

**N\_ds**



**P**


Mask\_Width is 5



For Mask\_Width=5, we load  $8+5-1=12$  elements  
(12 memory loads)

# Each output P element uses 5 N elements

**N\_ds**



P[8] uses N[6], N[7], N[8], N[9], N[10]

P[9] uses N[7], N[8], N[9], N[10], N[11]

P[10] use N[8], N[9], N[10], N[11], N[12]

...

P[14] uses N[12], N[13], N[14], N[15], N[16]

P[15] uses N[13], N[14], N[15], N[16], N[17]



# A simple way to calculate tiling benefit

- $(8+5-1)=12$  elements loaded
- $8*5$  global memory accesses replaced by shared memory accesses
- This gives a bandwidth reduction of  $40/12=3.3$

# In General, for 1D TILED CONVOLUTION


- $O\_TILE\_WIDTH + MASK\_WIDTH - 1$  elements loaded for each input tile
- $O\_TILE\_WIDTH * MASK\_WIDTH$  global memory accesses replaced by shared memory accesses
- This gives a reduction factor of

$$(O\_TILE\_WIDTH * MASK\_WIDTH) / (O\_TILE\_WIDTH + MASK\_WIDTH - 1)$$

This ignores ghost elements in edge tiles.

# Another Way to Look at Reuse

**N\_ds**



N[6] is used by P[8] (1X)  
N[7] is used by P[8], P[9] (2X)  
N[8] is used by P[8], P[9], P[10] (3X)  
N[9] is used by P[8], P[9], P[10], P[11] (4X)  
N[10] is used by P[8], P[9], P[10], P[11], P[12] (5X)  
... (5X)  
N[14] is used by P[12], P[13], P[14], P[15] (4X)  
N[15] is used by P[13], P[14], P[15] (3X)

# Another Way to Look at Reuse

The total number of global memory accesses  
(to the  $(8+5-1)=12$  N elements) replaced by shared  
memory accesses is:

$$\begin{aligned} &1 + 2 + 3 + 4 + 5 * (8-5+1) + 4 + 3 + 2 + 1 \\ &= 10 + 20 + 10 \\ &= 40 \end{aligned}$$

So the reduction is:

$$40/12 = 3.3$$

# In General, for 1D

- The total number of global memory accesses to the input tile can be calculated as

$$\begin{aligned} & 1 + 2 + \dots + \text{MASK\_WIDTH} - 1 + \text{MASK\_WIDTH} * (\text{O\_TILE\_WIDTH} - \\ & \quad \text{MASK\_WIDTH} + 1) + \text{MASK\_WIDTH} - 1 + \dots + 2 + 1 \\ = & \text{MASK\_WIDTH} * (\text{MASK\_WIDTH} - 1) + \text{MASK\_WIDTH} * \\ & \quad (\text{O\_TILE\_WIDTH} - \text{MASK\_WIDTH} + 1) \\ = & \text{MASK\_WIDTH} * \text{O\_TILE\_WIDTH} \end{aligned}$$

For a total of  $\text{O\_TILE\_WIDTH} + \text{MASK\_WIDTH} - 1$  input tile elements

# Examples of Bandwidth Reduction for 1D

The reduction ratio is:

$$\text{MASK\_WIDTH} * (\text{O\_TILE\_WIDTH}) / (\text{O\_TILE\_WIDTH} + \text{MASK\_WIDTH} - 1)$$

O_TILE_WIDTH	16	32	64	128	256
MASK_WIDTH= 5	4.0	4.4	4.7	4.9	4.9
MASK_WIDTH = 9	6.0	7.2	8.0	8.5	8.7

# For 2D Convolution Tiles

- $(O\_TILE\_WIDTH + MASK\_WIDTH - 1)^2$  input elements need to be loaded into shared memory
- The calculation of each output element needs to access  $MASK\_WIDTH^2$  input elements
- $O\_TILE\_WIDTH^2 * MASK\_WIDTH^2$  global memory accesses are converted into shared memory accesses
- The reduction ratio is

$$O\_TILE\_WIDTH^2 * MASK\_WIDTH^2 / (O\_TILE\_WIDTH + MASK\_WIDTH - 1)^2$$

# Bandwidth Reduction for 2D

The reduction ratio is:

$$\frac{O\_TILE\_WIDTH^2 * MASK\_WIDTH^2}{(O\_TILE\_WIDTH + MASK\_WIDTH - 1)^2}$$

O_TILE_WIDTH	8	16	32	64
MASK_WIDTH = 5	11.1	16	19.7	22.1
MASK_WIDTH = 9	20.3	36	51.8	64

Tile size has significant effect on of the memory bandwidth reduction ratio.

This often argues for larger shared memory size.