

# IF3230 – Sistem Terdistribusi

## Clock Synchronization

Achmad Imam Kistijantoro ([imam@informatika.org](mailto:imam@informatika.org))

Judhi Santoso ([judhi@informatika.org](mailto:judhi@informatika.org))

Anggrahita Bayu Sasmita ([bayu.anggrahita@informatika.org](mailto:bayu.anggrahita@informatika.org))

# Clock Synchronization

---

- ▶ physical clock
- ▶ logical clock
- ▶ vector clock



# Physical clock

---

- ▶ Koordinasi antar proses yang berjalan konkuren sering memerlukan order (keterurutan) antar event
- ▶ Misal:
  - ▶ untuk menentukan urutan update terhadap data yang terreplikasi
  - ▶ Menentukan urutan pesan yang akan diproses/ditampilkan
  - ▶ Scheduler, timeout, failure detectors, performance measurements, cache validity
- ▶ Clock pada komputer berbasis quartz crystal clock, untuk yang standar dapat memiliki akurasi 6 ppm (sekitar  $\frac{1}{2}$  detik/hari)
- ▶ Clock yang bagus dapat mencapai akurasi 1 detik dalam 10 tahun, namun sensitif thd perubahan suhu, dan freq dapat berubah sesuai dengan usia quartz crystal



# Physical clock

---

- ▶ Atomic clock:
- ▶ Waktu referensi didefinisikan sebagai 9,192,631,770 periode radiasi yang berkorespondensi dengan 2 hyperfine level dari cesium-133
- ▶ Akurasi 1 detik dalam 6 juta tahun
- ▶ Standar NIST sejak 1960
- ▶ Pewaktuan standar berbasis atomic clock: UTC (Coordinated Universal Time)



# Leap second

---

- ▶ UTC menggunakan atomic clock, yang tidak persis sama dengan GMT (solar time) yang menggunakan rotasi bumi dan matahari => perputaran rotasi bumi tidak selalu konstan
- ▶ Kadang perlu dilakukan koreksi detik (leap second), dan dilakukan pada 30 juni dan 31 desember setiap tahun
  - ▶ Dimajukan 1 detik
  - ▶ Tetap
  - ▶ Mundur 1 detik



# Leap second - problem

---

- ▶ Penanganan software/komputer terhadap leap second?
  - ▶ Diabaikan
- ▶ OS dan system terdistribusi sering bergantung pada timing dengan akurasi  $< 1$  s
- ▶ 30 Juni 2012: bug pada linux mengakibatkan livelock pada leap second, menyebabkan banyak layanan Internet yang down

<https://www.wired.com/2012/07/leap-second-glitch-explained/>



The screenshot shows the top portion of a Wired article. The Wired logo is on the left, followed by navigation links: BACKCHANNEL, BUSINESS, CULTURE, GEAR, IDEAS, SCIENCE, and SECURITY. On the right are links for SIGN IN and SUBSCRIBE, along with a search icon. Below the navigation bar, the article's byline reads 'ROBERT MCMILLAN CADE METZ BUSINESS JUL 2, 2012 7:54 PM'. The main headline is 'The Inside Story of the Extra Second That Crashed the Web'. The first line of the article text is visible: 'The "leap second" crash -- which hit several web operations on Saturday evening -- can be traced to a single glitch in the Linux operating system. Here's the inside story on what happened.'

Wired

BACKCHANNEL BUSINESS CULTURE GEAR IDEAS SCIENCE SECURITY

SIGN IN SUBSCRIBE

ROBERT MCMILLAN CADE METZ BUSINESS JUL 2, 2012 7:54 PM

## The Inside Story of the Extra Second That Crashed the Web

The "leap second" crash -- which hit several web operations on Saturday evening -- can be traced to a single glitch in the Linux operating system. Here's the inside story on what happened.

# Physical clock

---


- ▶ Problem: 2 komputer tidak pernah memiliki physical clock yang sinkron
- ▶ Setiap quartz crystal memiliki frekuensi yang sedikit berbeda
  - ▶ Antar clock memiliki gap yang membesar dengan rate tertentu, yang disebut sebagai **time drift**
  - ▶ Selisih waktu antar 2 clock disebut sebagai **time skew**



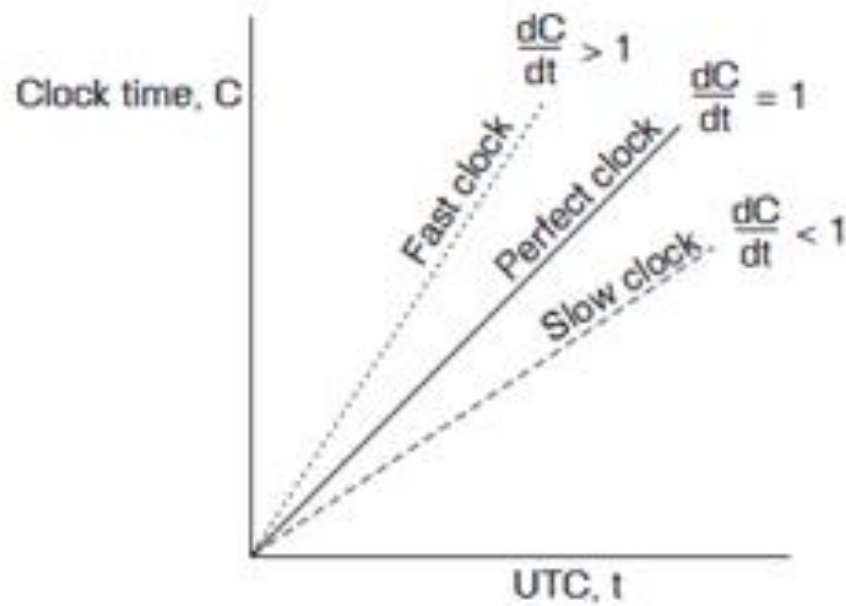
# Quartz clock error: drift

---

- ▶ Dipengaruhi lingkungan, e.g. suhu
- ▶ 1 ppm = 1 microsecond/second = 86 ms/day = 32 s /year
- ▶ Typical computer: 50 ppm







In practice:  $1 - \rho \leq \frac{dC}{dt} \leq 1 + \rho$ .

# Penanganan drift

---

- ▶ Bagaimana mencocokkan waktu yang mengalami drift
- ▶ Clock sebaiknya tidak di-set mundur
  - ▶ Mengacaukan order message dan lingkungan pengembangan software




# Penanganan drift


---

- ▶ **Dengan koreksi gradual**

- ▶ Jika terlalu cepat, buat clock berjalan lebih lambat hingga sinkron
- ▶ Jika terlalu lambat, buat clock berjalan lebih cepat hingga sinkron
- ▶ Pada komputer, hal ini dapat dilakukan pada level OS, yaitu dengan mengubah frekuensi saat pembangkitan interrupt clock
- ▶ Pada UNIX-based, disediakan oleh fungsi adjtime (lihat man pada Linux)







# Mendapatkan waktu akurat

---

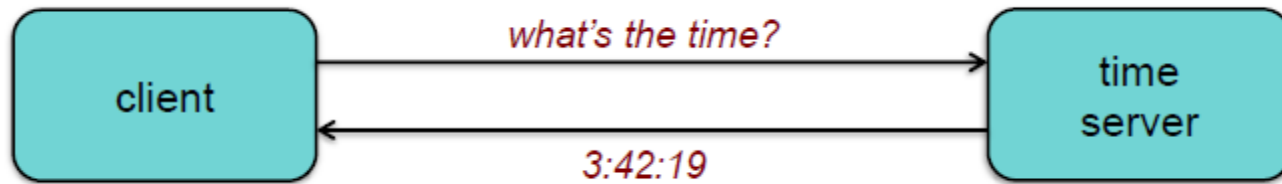
- ▶ Menggunakan GPS receiver yang terhubung ke komputer
  - ▶ Dapat memberikan waktu akurat hingga selisih 1 ms terhadap UTC
  - ▶ Tidak dapat digunakan in-door
- ▶ Di US, dapat menggunakan WWV radio receiver
  - ▶ Akurasi 3 - 10 ms, tergantung lokasi
- ▶ Menggunakan GOES (Geostationary Operational Environment Satellites)
  - ▶ Akurasi 0.1 ms
- ▶ Solusi di atas tidak praktis, sehingga umumnya sinkronisasi dilakukan dengan mencocokkan waktu dengan komputer lain yang lebih akurat => time server



# Sinkronisasi

---

- ▶ Cara sederhana
  - ▶ Meminta waktu melalui jaringan
  - ▶ Set waktu sesuai dengan jawaban




- ▶ Belum mempertimbangkan network delay

# Algoritma Cristian

---


- ▶ **Kompensasi delay**
  - ▶  $T_0$  : request dikirim
  - ▶  $T_1$  : reply diterima
- ▶ **Asumsi network delay simetrik**





# Algoritma Cristian


---



►  $T_{new} = T_{server} + (T_1 - T_0)/2$

# Algoritma Cristian

- ▶ Jika waktu pengiriman pesan minimum diketahui, dapat dihitung batasan akurasi



- ▶ Akurasi:  $(T_1 - T_0)/2 - T_{min}$

# Algoritma Berkeley

---

- ▶ Gusella & Zatti, 1989
- ▶ Asumsi: tidak ada mesin yang memiliki sumber waktu akurat
- ▶ Menghitung rata2 waktu dari semua komputer
- ▶ Sinkronisasi semua komputer dengan waktu rata2



# Algoritma Berkeley


---

- ▶ Komputer menjalankan daemon yang mengimplementasikan protokol
- ▶ 1 mesin berfungsi sebagai master, lainnya slave
- ▶ Master poll setiap komputer periodik, menanyakan waktu
  - ▶ Dapat menggunakan algoritma cristian untuk kompensasi latency
- ▶ Saat reply diterima master, hitung waktu rata2
- ▶ Master mengirimkan informasi offset ke semua komputer
- ▶ Mesin yang memiliki beda waktu besar diabaikan

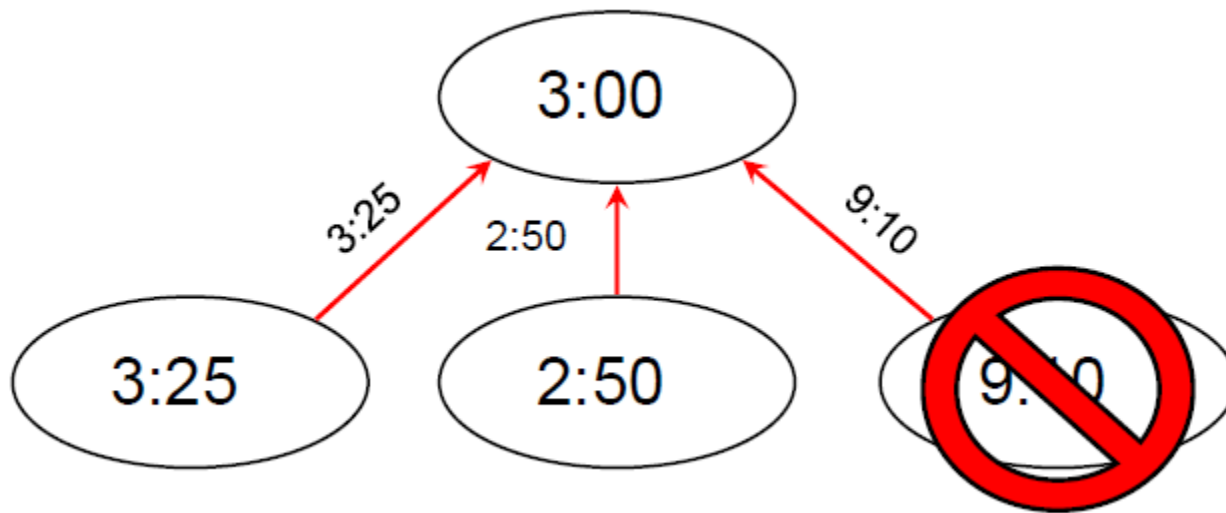



# Berkeley Algorithm

---



- 
- ▶ Waktu rata2: 3:25, 2:50, 3:00 = 3:05





# Network Time Protocol

---

- ▶ 1991, 1992, Internet Standard v3: RFC 1305
- ▶ June 2010
  - ▶ Internet Standard v4: RFC 5905-5908
  - ▶ IPv6 support
  - ▶ Dynamic server discovery





# NTP Goals

---

- Enable clients across Internet to be **accurately** synchronized to UTC despite message delays
  - Use statistical techniques to filter data and gauge quality of results
- Provide **reliable** service
  - Survive lengthy losses of connectivity
  - Redundant paths
  - Redundant servers
- Provide **scalable** service
  - Enable clients to **synchronize frequently**
  - Offset effects of clock drift
- Provide **protection** against interference
  - Authenticate source of data




# NTP Servers

---

## Arranged in strata

- 1<sup>st</sup> stratum: machines connected directly to accurate time source
- 2<sup>nd</sup> stratum: machines synchronized from 1<sup>st</sup> stratum machines
- ...



Synchronization Subnet

# NTP Synchronization Modes

---

## Multicast mode

- for high speed LANS
- Lower accuracy but efficient

## Procedure call mode

- Similar to Cristian's algorithm

## Symmetric mode


- Intended for master servers
- Peer servers can synchronize with each other to provide mutual backup
  - Pair of servers retain data to improve synchronization over time

All messages delivered unreliably with UDP



# Simple NTP

---



- ▶  $d = (T_4 - T_1) - (T_3 - T_2)$
- ▶ Offset:  $((T_2 - T_1) + (T_3 - T_4)) / 2$

# Logical Clock

---



# the happened before relationship

---

## Problem

We first need to introduce a notion of ordering before we can order anything.

## The happened-before relation

- If  $a$  and  $b$  are two events in the same process, and  $a$  comes before  $b$ , then  $a \rightarrow b$ .
- If  $a$  is the sending of a message, and  $b$  is the receipt of that message, then  $a \rightarrow b$ .
- If  $a \rightarrow b$  and  $b \rightarrow c$ , then  $a \rightarrow c$ .

## Note

This introduces a **partial ordering of events** in a system with concurrently operating processes.



# logical clock

---

## Problem

How do we maintain a global view on the system's behavior that is consistent with the happened-before relation?

## Solution

Attach a timestamp  $C(e)$  to each event  $e$ , satisfying the following properties:

- P1 If  $a$  and  $b$  are two events in the same process, and  $a \rightarrow b$ , then we demand that  $C(a) < C(b)$ .
- P2 If  $a$  corresponds to sending a message  $m$ , and  $b$  to the receipt of that message, then also  $C(a) < C(b)$ .

## Problem

How to attach a timestamp to an event when there's no global clock  $\Rightarrow$  maintain a **consistent** set of logical clocks, one per process.



# logical clock

---

## Solution

Each process  $P_i$  maintains a **local** counter  $C_i$  and adjusts this counter according to the following rules:

- 1: For any two **successive events** that take place within  $P_i$ ,  $C_i$  is incremented by 1.
- 2: Each time a message  $m$  is **sent** by process  $P_i$ , the message receives a timestamp  $ts(m) = C_i$ .
- 3: Whenever a message  $m$  is **received** by a process  $P_j$ ,  $P_j$  adjusts its local counter  $C_j$  to  **$\max\{C_j, ts(m)\}$** ; then executes step 1 before passing  $m$  to the application.

## Notes


- Property **P1** is satisfied by (1); Property **P2** by (2) and (3).
- It can still occur that two events happen at the same time. Avoid this by **breaking ties through process IDs**.



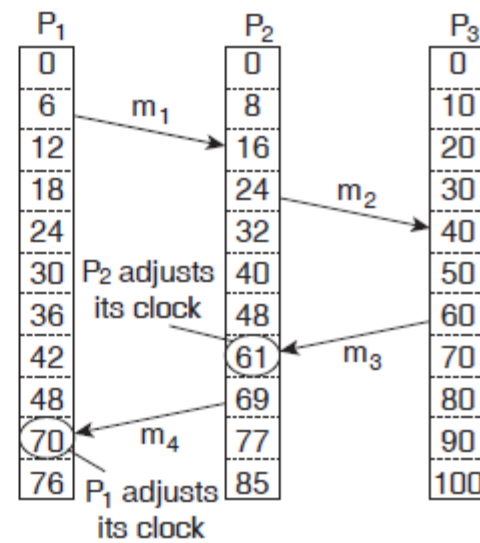


# logical clock

---



(a)

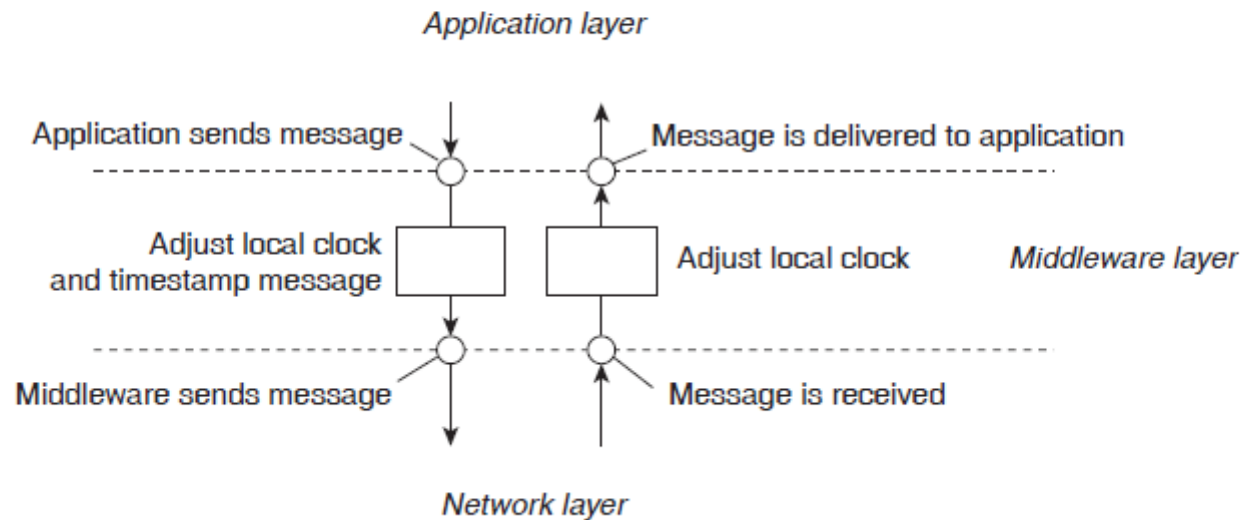


(b)

# logical clock - example

## Note

Adjustments take place in the middleware layer




# example – totally ordered multicast

## Problem

We sometimes need to guarantee that concurrent updates on a replicated database are seen in the same order everywhere:

- $P_1$  adds \$100 to an account (initial value: \$1000)
- $P_2$  increments account by 1%
- There are two replicas



## Result

In absence of proper synchronization:

replica #1  $\leftarrow$  \$1111, while replica #2  $\leftarrow$  \$1110.

# example – totally ordered multicast

---

## Solution

- Process  $P_i$  sends timestamped message  $msg_i$  to all others. The message itself is put in a local queue  $queue_i$ .
- Any incoming message at  $P_j$  is queued in  $queue_j$ , according to its timestamp, and acknowledged to every other process.

$P_j$  passes a message  $msg_i$  to its application if:

- (1)  $msg_i$  is at the head of  $queue_j$
- (2) for each process  $P_k$ , there is a message  $msg_k$  in  $queue_j$  with a larger timestamp.

## Note


We are assuming that communication is reliable and FIFO ordered.



# vector clock

## Observation

Lamport's clocks do not guarantee that if  $C(a) < C(b)$  that  $a$  causally preceded  $b$



## Observation

Event  $a$ :  $m_1$  is received at  $T = 16$ ;  
Event  $b$ :  $m_2$  is sent at  $T = 20$ .

## Note

We cannot conclude that  $a$  causally precedes  $b$ .

# vector clock

---

## Solution


- Each process  $P_i$  has an array  $VC_i[1..n]$ , where  $VC_i[j]$  denotes the number of events that process  $P_i$  knows have taken place at process  $P_j$ .
- When  $P_i$  sends a message  $m$ , it adds 1 to  $VC_i[i]$ , and sends  $VC_i$  along with  $m$  as vector timestamp  $vt(m)$ . Result: upon arrival, recipient knows  $P_i$ 's timestamp.
- When a process  $P_j$  delivers a message  $m$  that it received from  $P_i$  with vector timestamp  $ts(m)$ , it
  - (1) updates each  $VC_j[k]$  to  $\max\{VC_j[k], ts(m)[k]\}$
  - (2) increments  $VC_j[j]$  by 1.

## Question


What does  $VC_i[j] = k$  mean in terms of messages sent and received?



# Contoh



# Contoh





# Sumber

---

- ▶ Paul Krzyzanowski, Clock Synchronization, Lectures on Distributed Systems, Rutgers University 2013
- ▶ Andrew S. Tanenbaum & Marten v. Steen, Distributed Systems Principles and Paradigms, 2<sup>nd</sup> edition, Chapter 6, Prentice Hall, 2007

