

IF3230 – Sistem Terdistribusi

Distributed File Systems – Andrew File Systems

Achmad Imam Kistijantoro (imam@informatika.org)

Judhi Santoso (judhi@informatika.org)

Anggrahita Bayu Sasmita (bayu.anggrahita@informatika.org)

Topik

- Generic Distributed File System
- NFS Network File System
- **AFS Andrew File System**
- Google File System
- Hadoop Distributed File System (HDFS)



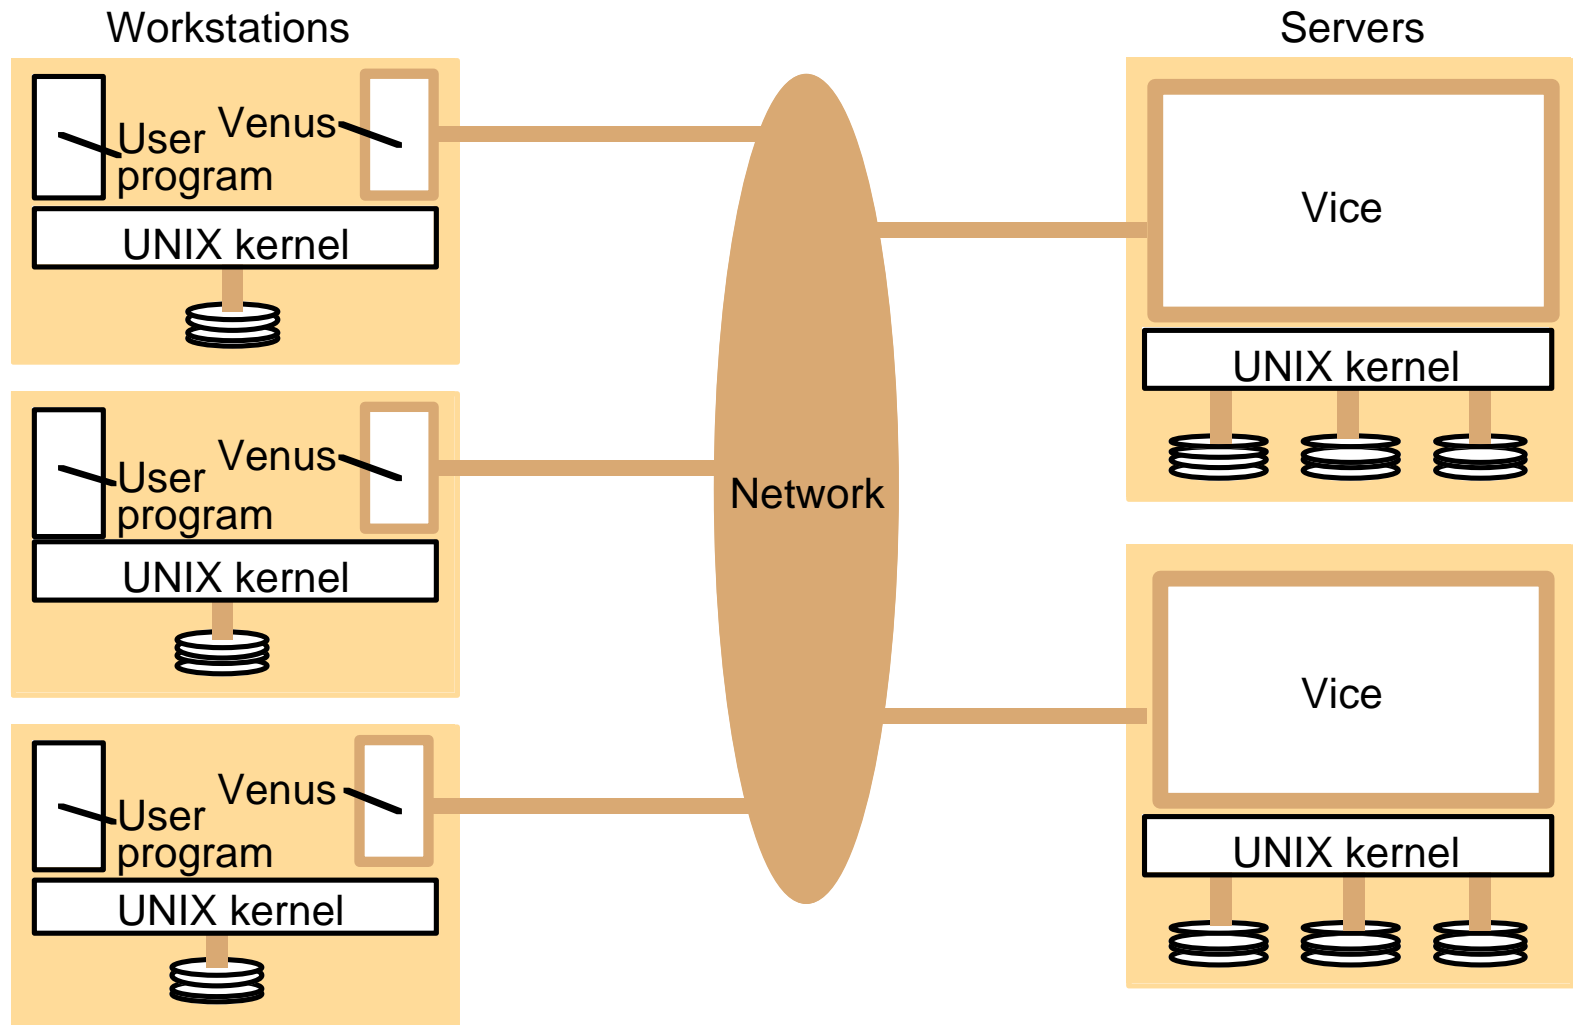
Andrew File System

- ▶ berbeda dengan NFS, AFS lebih mengutamakan scalability.
- ▶ scalability dicapai dengan melakukan caching keseluruhan file pada client nodes.
- ▶ Sebuah client cache dapat menyimpan beberapa ratus file yang most recently used pada komputer tersebut
- ▶ Cache bersifat permanent.
- ▶ Saat client membuka file, cache diperiksa dan digunakan jika sudah pernah di cache sebelumnya.

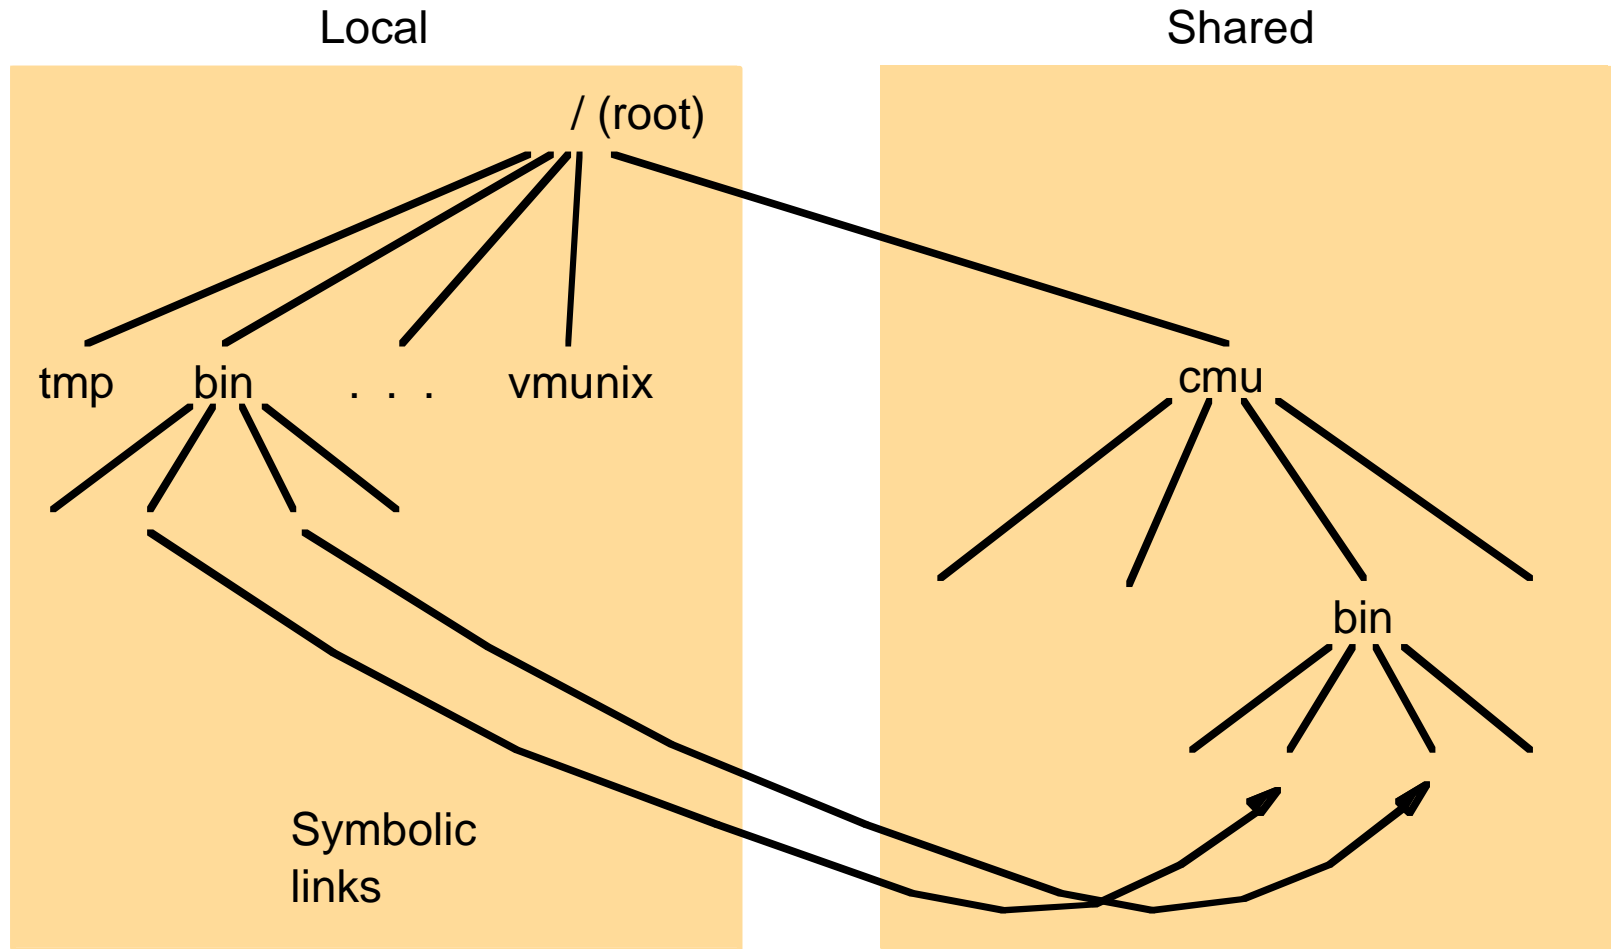
Andrew File System - Typical Scenario

- ▶ Saat code client membuka file, client cache dibaca terlebih dahulu. Jika tidak ada, lokasi server ditentukan dan file dibaca dari server tersebut.
- ▶ Copy file disimpan pada sisi client dan dibuka.
- ▶ reads dan writes berikutnya dilakukan pada copy client.
- ▶ Saat client menutup file – jika file berubah, file tersebut dikirim balik ke server. Copy file tetap disimpan pada client.
- ▶ Contoh: library dan command files. File yang hanya digunakan oleh single user => mayoritas akses file pada Unix
- ▶ **Principle: Make the common case fast.**


Distribution of processes in the Andrew File System



File name space seen by clients of AFS



System call interception in AFS



Implementation of file system calls in AFS

<i>User process</i>	<i>UNIX kernel</i>	<i>Venus</i>	<i>Net</i>	<i>Vice</i>
<i>open(FileName, mode)</i>	<p>If <i>FileName</i> refers to a file in shared file space, pass the request to Venus.</p> <p>Open the local file and return the file descriptor to the application.</p>	<p>Check list of files in local cache. If not present or there is no valid <i>callback promise</i>, send a request for the file to the Vice server that is custodian of the volume containing the file.</p> <p>Place the copy of the file in the local file system, enter its local name in the local cache list and return the local name to UNIX.</p>		<p>Transfer a copy of the file and a <i>callback promise</i> to the workstation. Log the callback promise.</p>
<i>read(FileDescriptor, Buffer, length)</i>	Perform a normal UNIX read operation on the local copy.			
<i>write(FileDescriptor, Buffer, length)</i>	Perform a normal UNIX write operation on the local copy.			
<i>close(FileDescriptor)</i>	Close the local copy and notify Venus that the file has been closed.	<p>If the local copy has been changed, send a copy to the Vice server that is the custodian of the file.</p>		<p>Replace the file contents and send a <i>callback</i> to all other clients holding <i>callback promises</i> on the file.</p>

The main components of the Vice service interface

<i>Fetch(fid) -> attr, data</i>	Returns the attributes (status) and, optionally, the contents of file identified by the <i>fid</i> and records a callback promise on it.
<i>Store(fid, attr, data)</i>	Updates the attributes and (optionally) the contents of a specified file.
<i>Create() -> fid</i>	Creates a new file and records a callback promise on it.
<i>Remove(fid)</i>	Deletes the specified file.
<i>SetLock(fid, mode)</i>	Sets a lock on the specified file or directory. The mode of the lock may be shared or exclusive. Locks that are not removed expire after 30 minutes.
<i>ReleaseLock(fid)</i>	Unlocks the specified file or directory.
<i>RemoveCallback(fid)</i>	Informs server that a Venus process has flushed a file from its cache.
<i>BreakCallback(fid)</i>	This call is made by a Vice server to a Venus process. It cancels the callback promise on the relevant file.
